



Feature Extraction – COMP4423 Computer Vision

Xiaoyong Wei (魏驍勇)

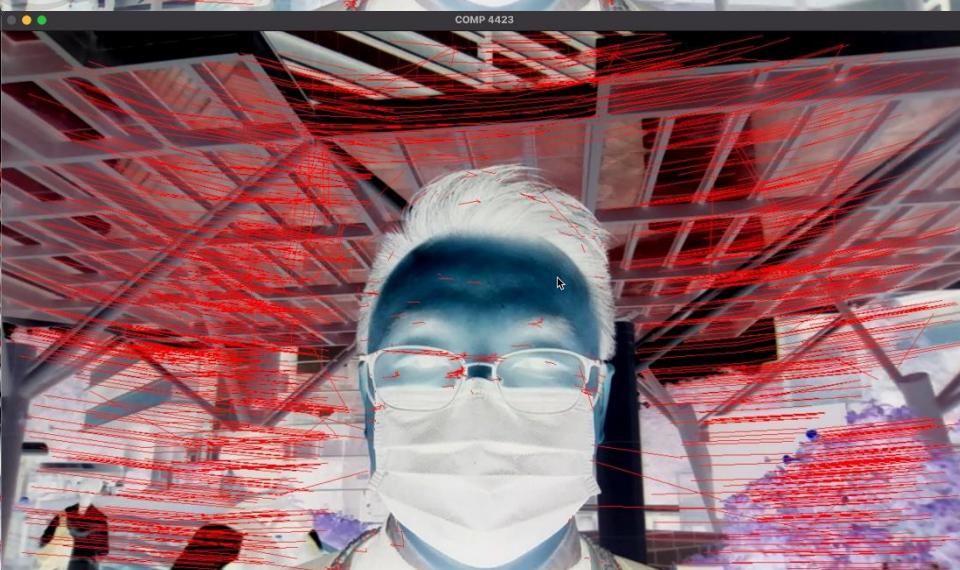
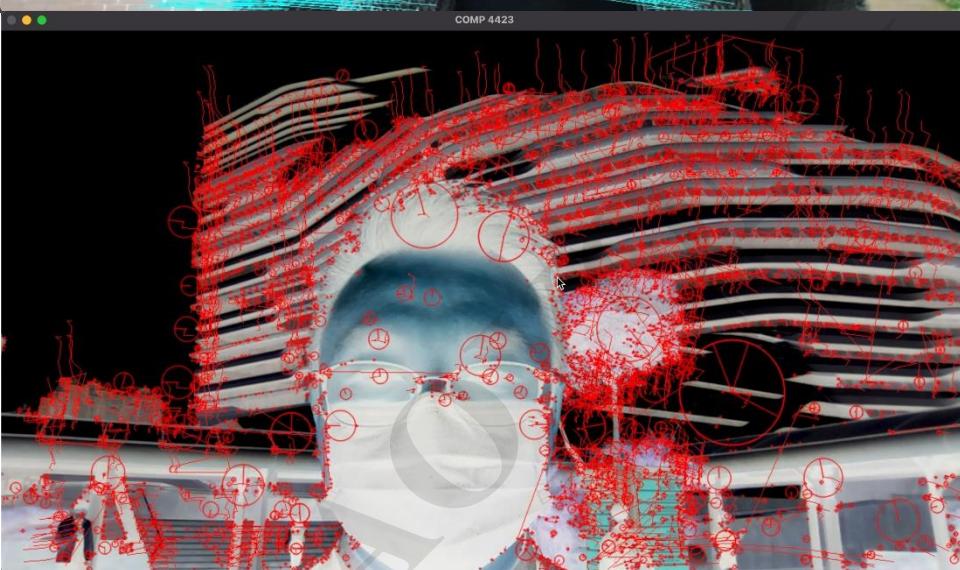
x1wei@polyu.edu.hk



A New Toy



A New Toy



Outline

- > Features and Quantization
- > Metrics (Distance and Similarity)
- > Global and Local Features (Color Histograms, LBP, SIFT)

How do you group them?



It's easy!



Yellow Hair, Blue Eyes,
White Skin



Black Hair, Brown Eyes,
Yellow Skin



Black Hair, Brown Eyes,
Yellow



Yellow Hair, Brown Eyes,
White Skin



Black Hair, Brown Eyes,
Black Skin

It's easy!

	Hair Color	Skin Color
	Yellow	White
	Black	Yellow
	Black	Yellow
	Yellow	White
	Black	Black

This is how human see it.

Is there a way for computers to
“see” it?

Quantization

	Hair Color	Skin Color
	Yellow (2)	White (1)
	Black (3)	Yellow (2)
	Black (3)	Yellow (2)
	Yellow (2)	White (1)
	Black (3)	Black (3)

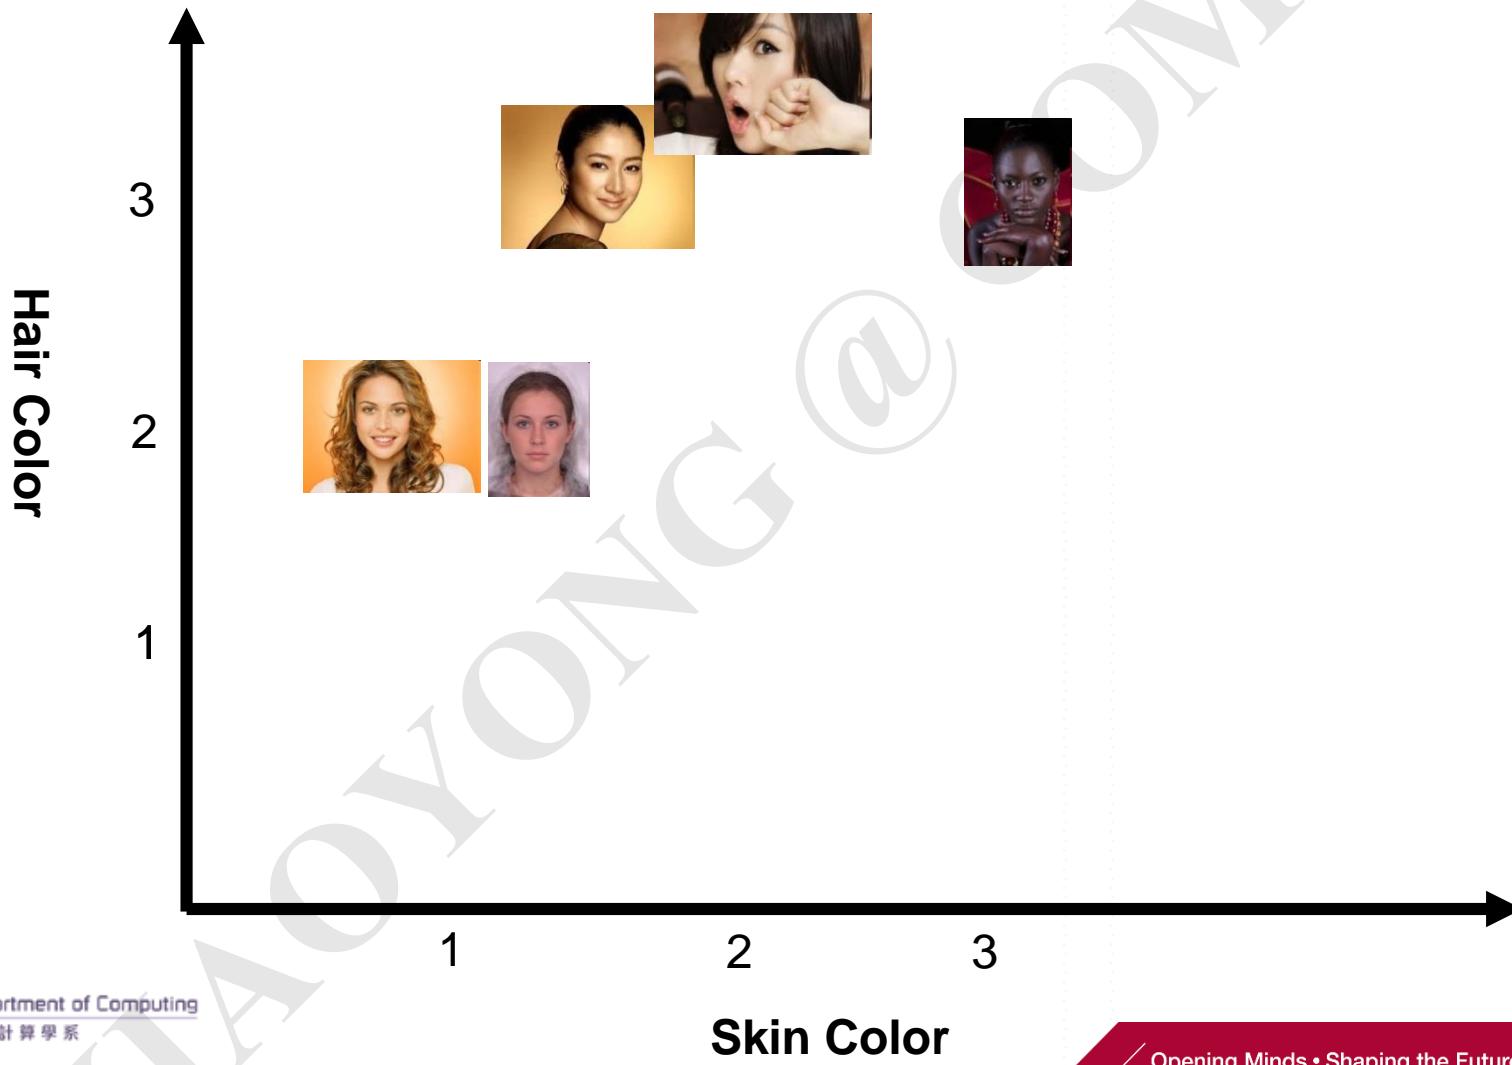
Great! Seems that now we
can build grouping rules
easily using if-else?

Right and wrong!

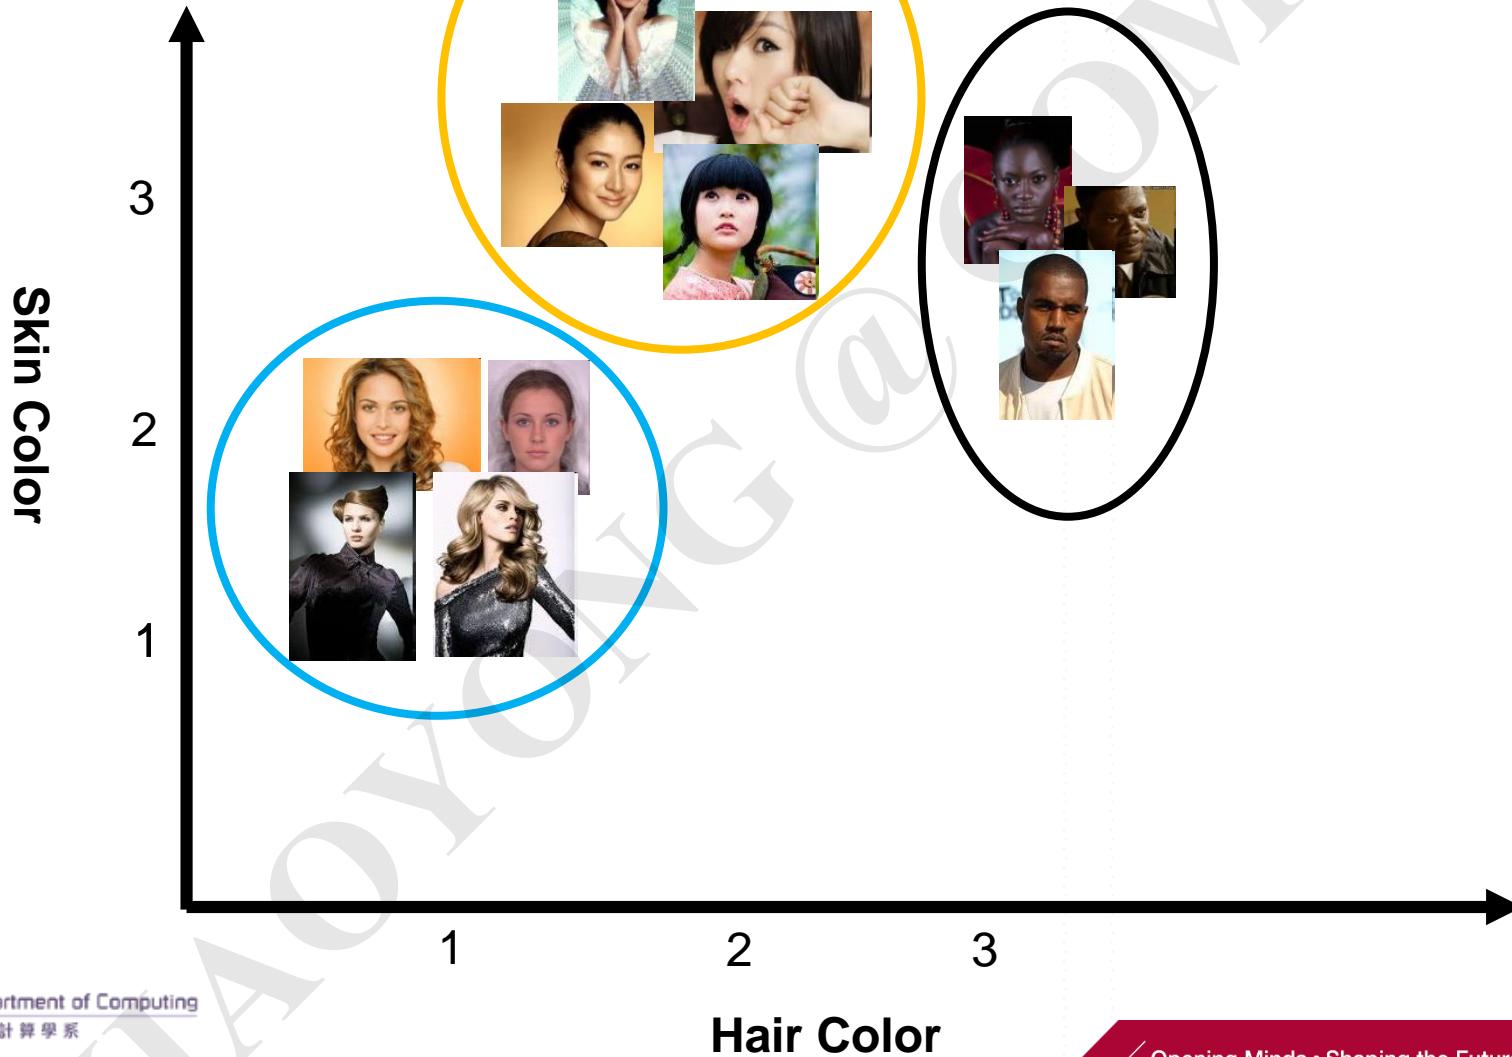
Feature Vectors - Quantization

	Hair Color	Skin Color
	Yellow (2.135)	White (1.05)
	Black (3.324)	Yellow (2.008)
	Black (2.945)	Yellow (2.355)
	Yellow (2.035)	White (1.035)
	Black (2.994)	Black (3.208)

Feature Space



Feature Space



Can we teach computers to draw circles for grouping?

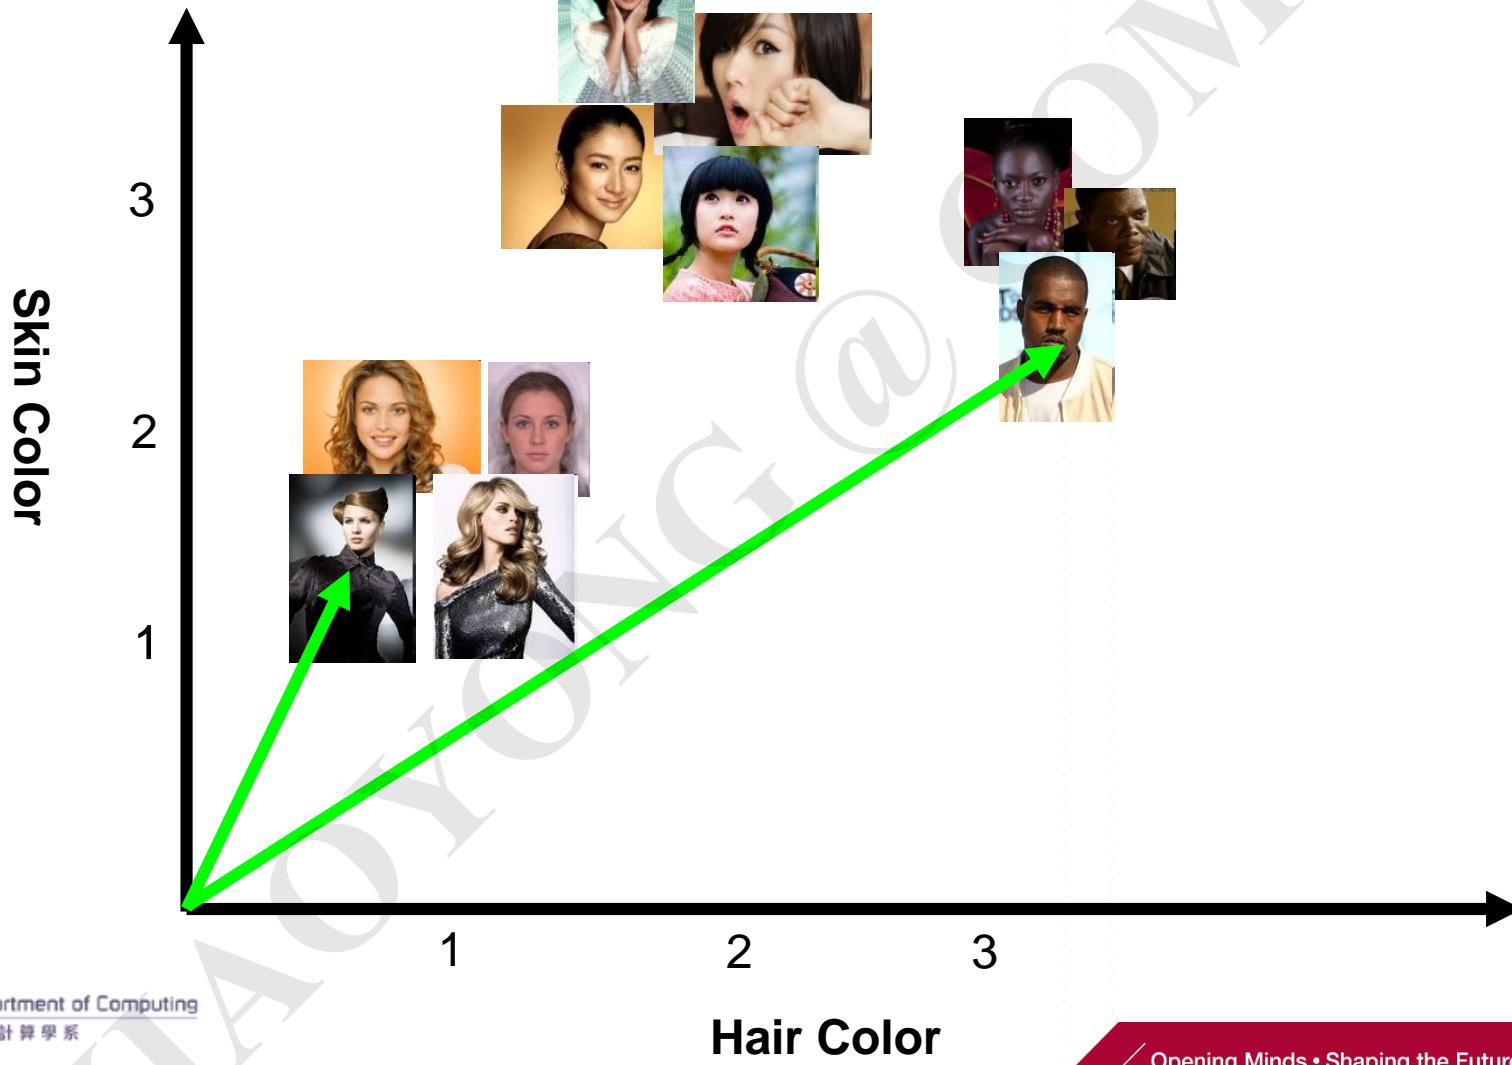
GOOD GUESS



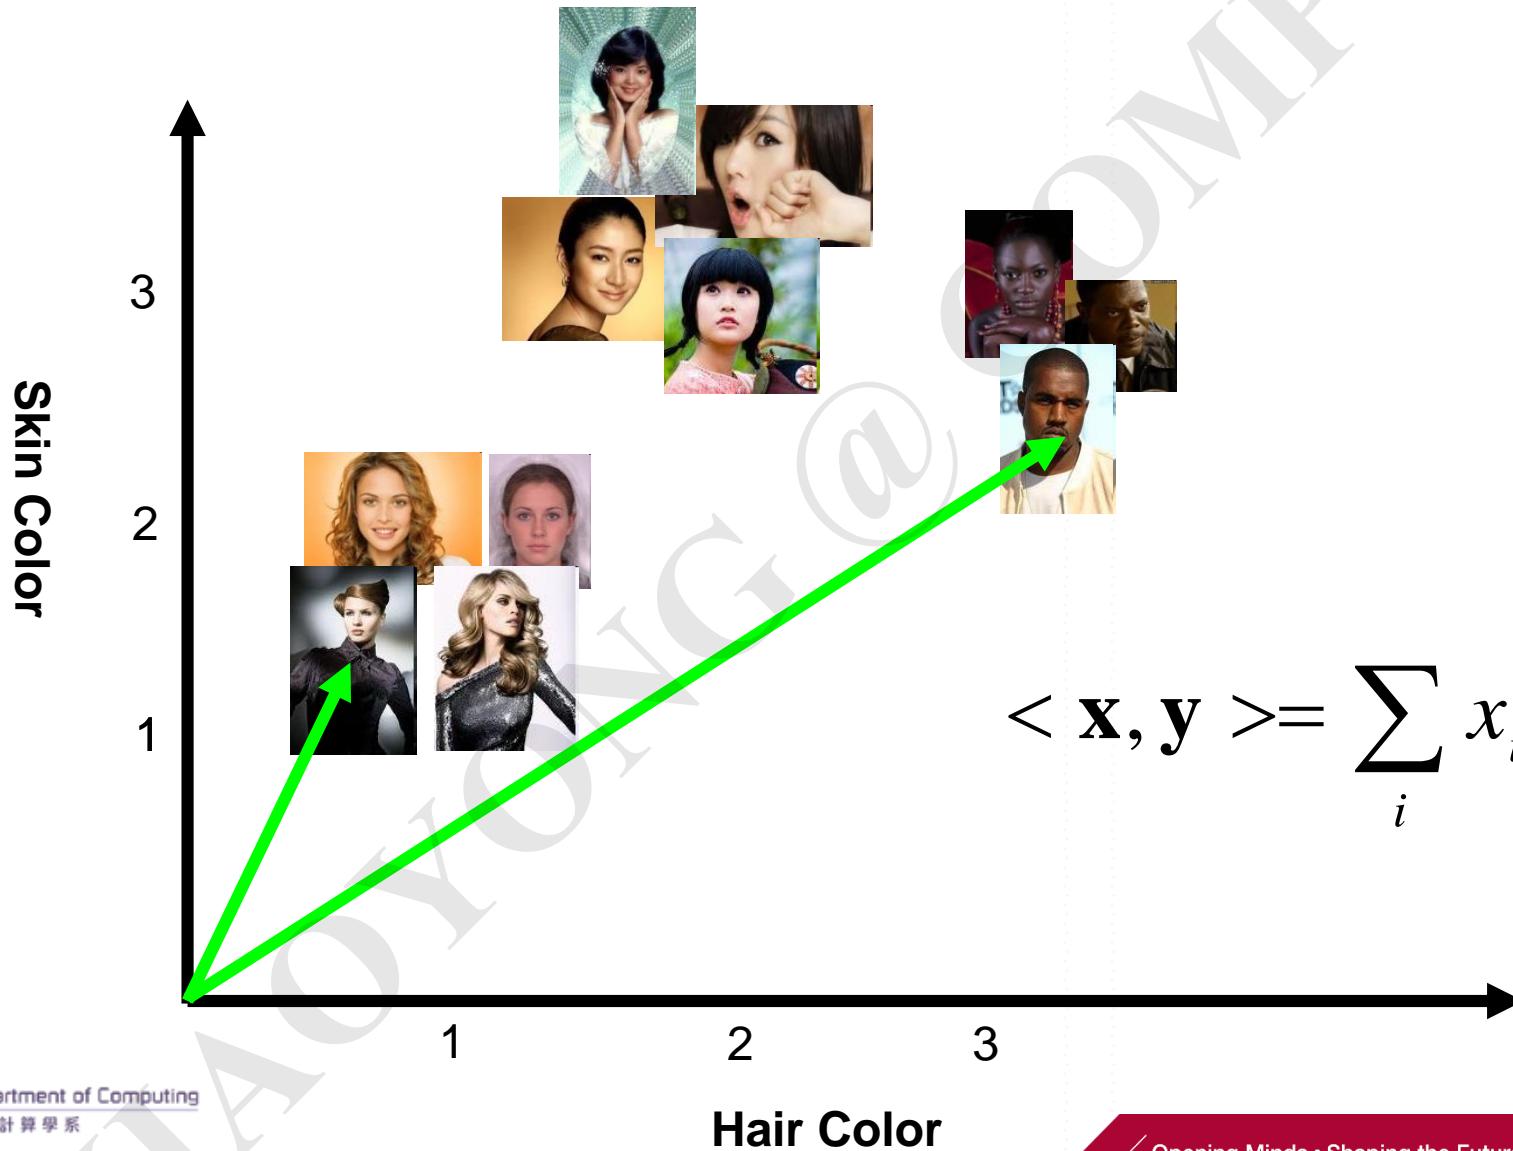
This is what we teach in this course!

But let's start from the most straightforward way by teaching computers to tell the similarity/difference.

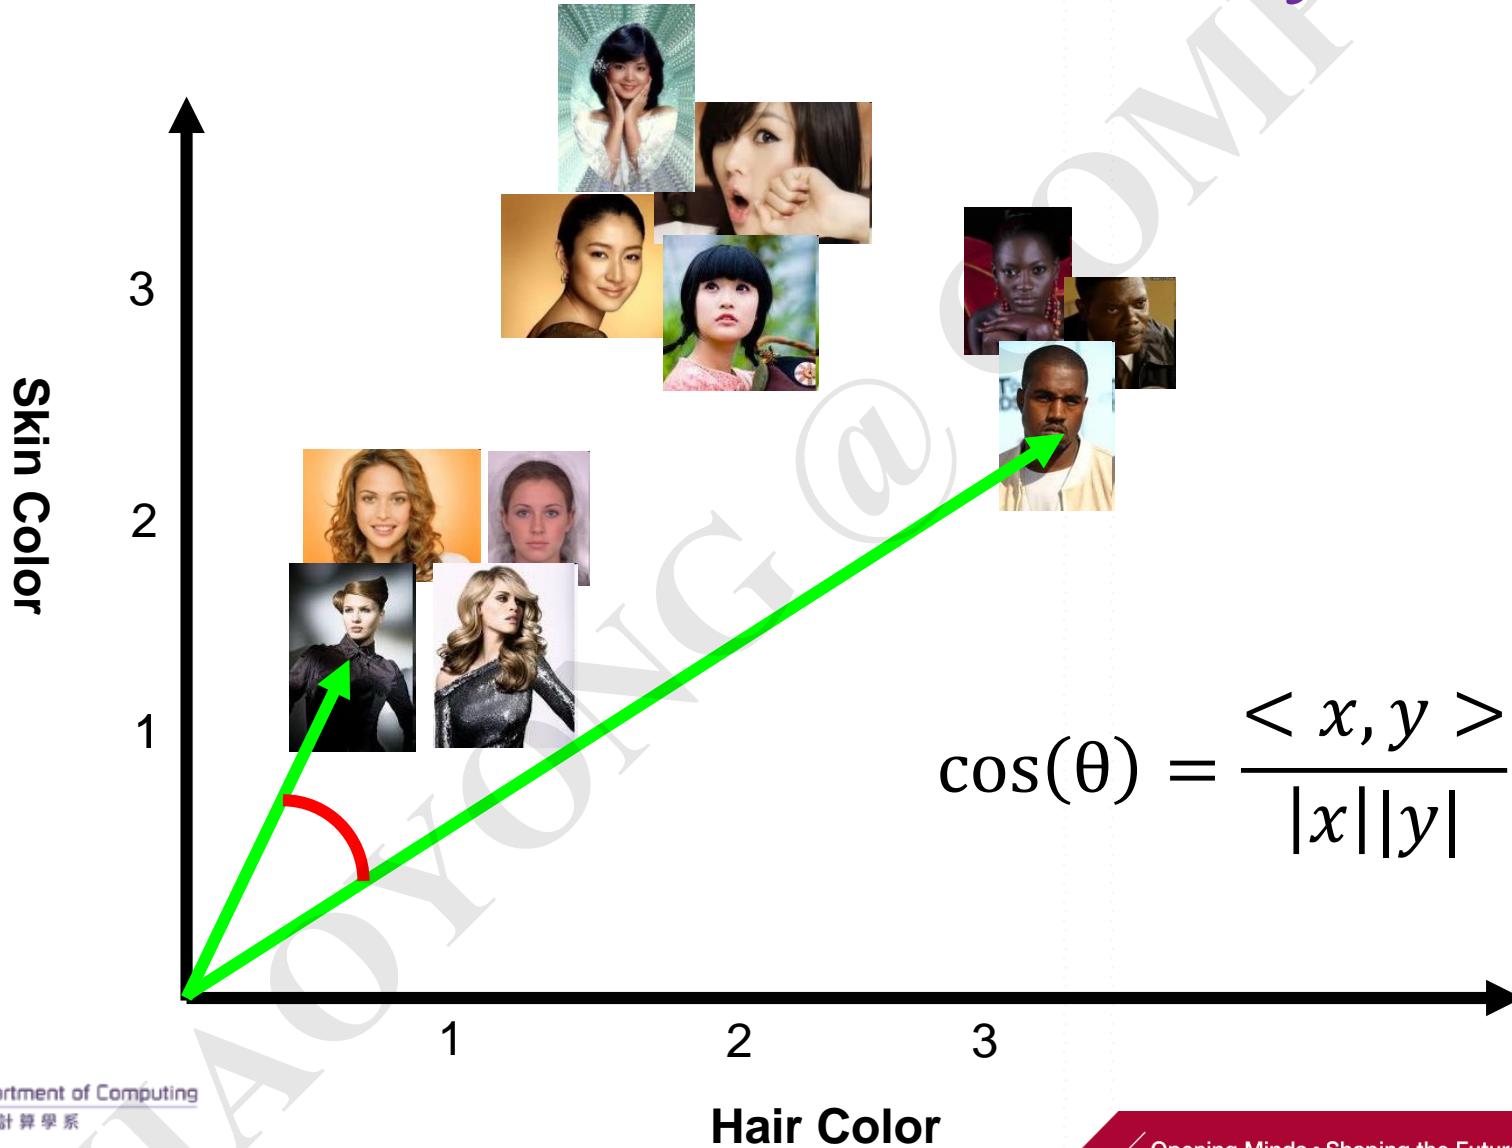
Metrics



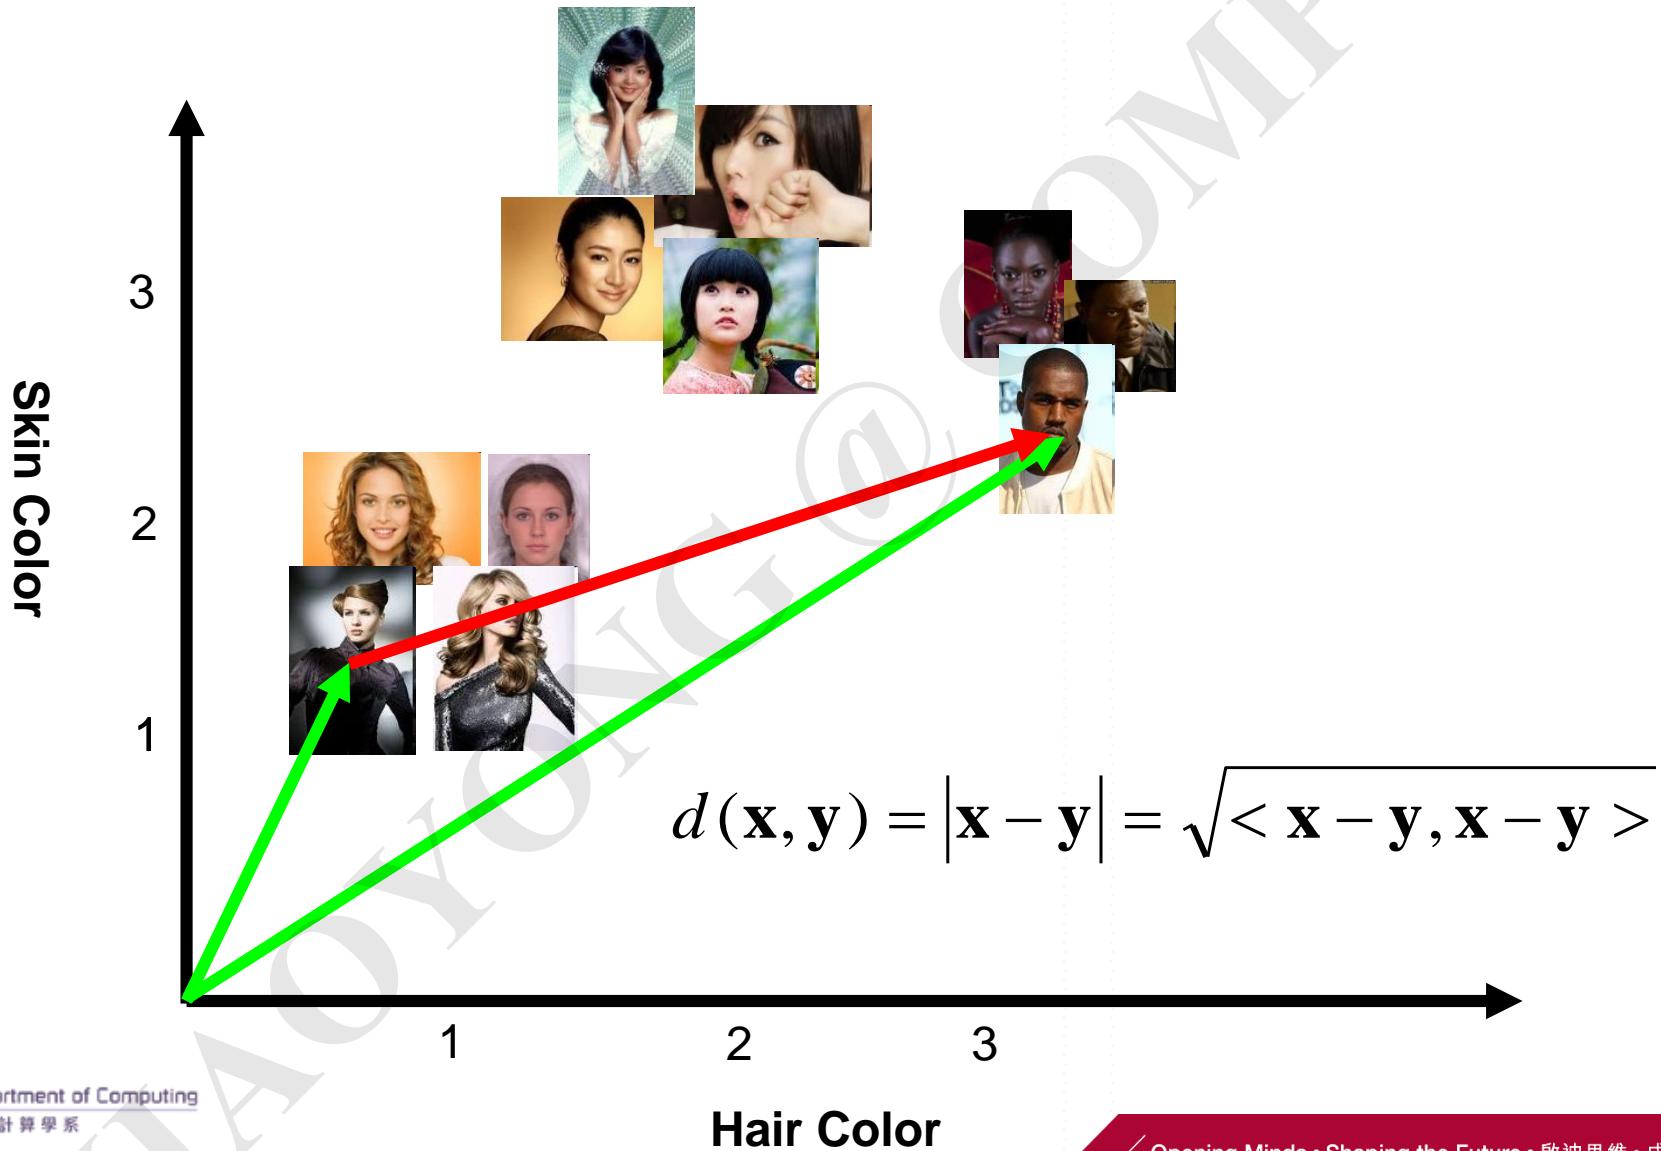
Metrics – Inner Product



Metrics – Cosine Similarity



Metrics – Euclidian Distance



I see!

We can tell if two images are from the same group by measuring the similarity/distance?

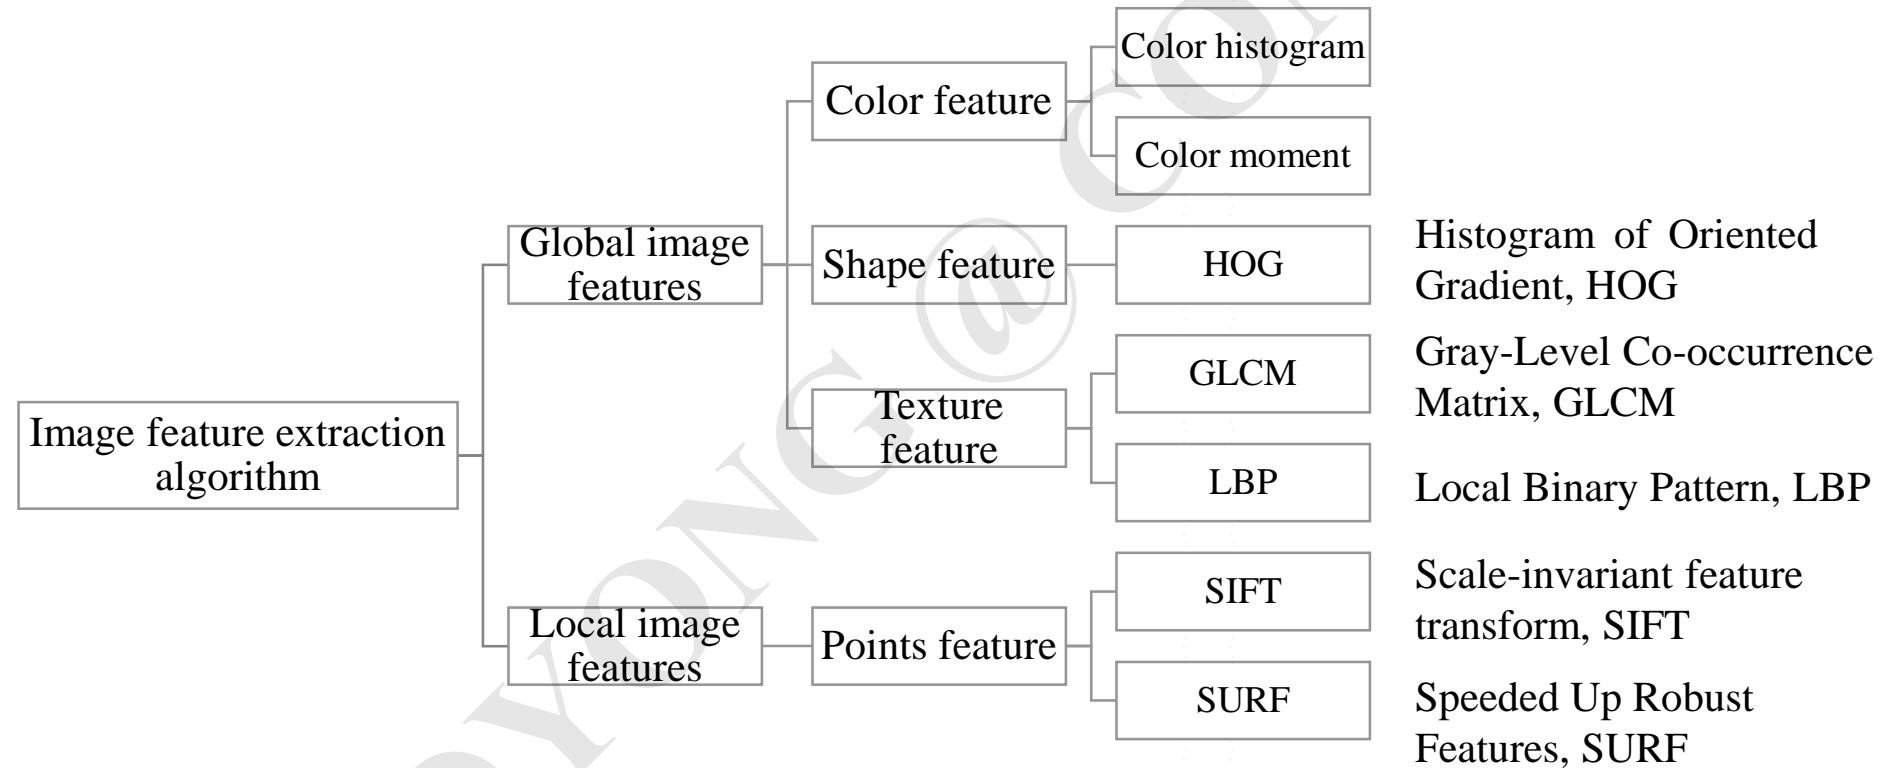
It's too early!

How can we obtain those numbers?

	Hair Color	Skin Color
	Yellow (2.135)	White (1.05)
	Black (3.324)	Yellow (2.008)
	Black (2.945)	Yellow (2.355)
	Yellow (2.035)	White (1.035)
	Black (2.994)	Black (3.208)

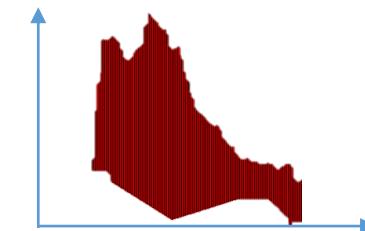
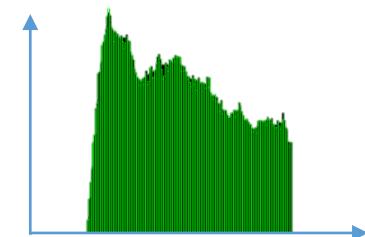
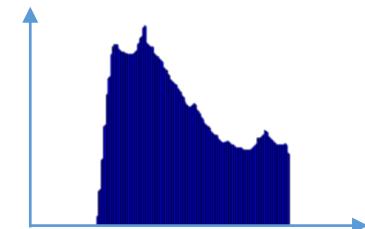
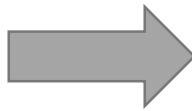
Feature Extraction

Features

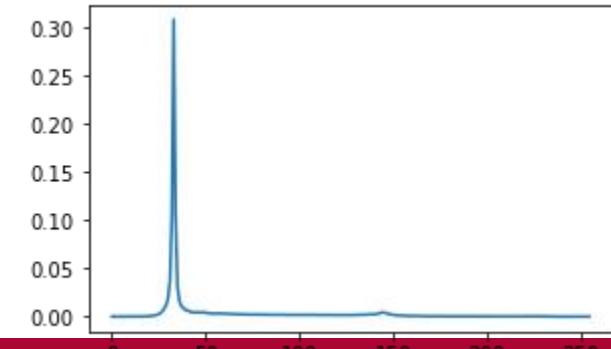
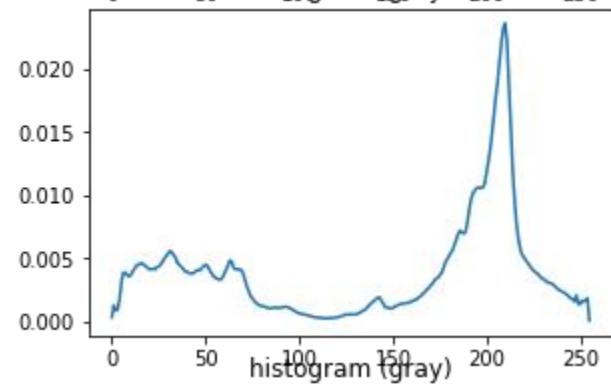
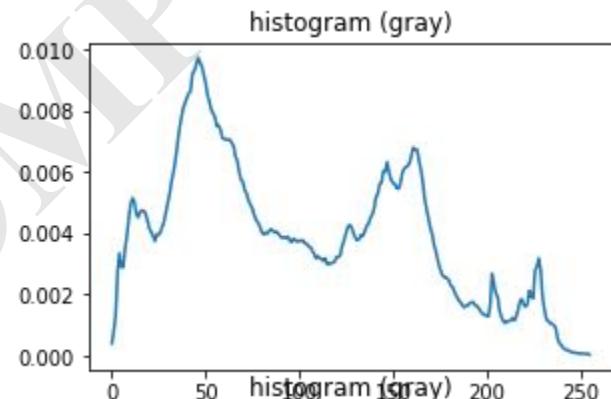
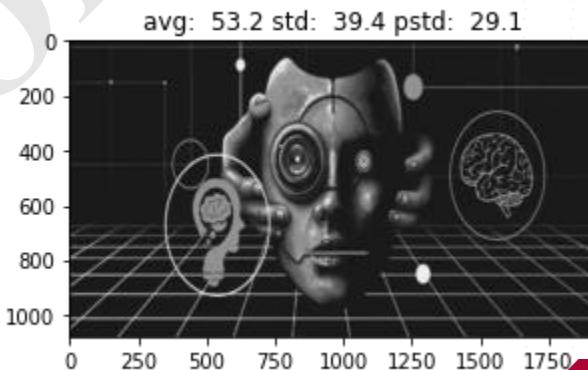
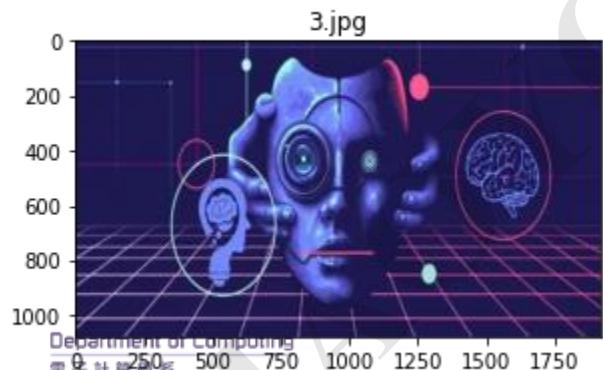
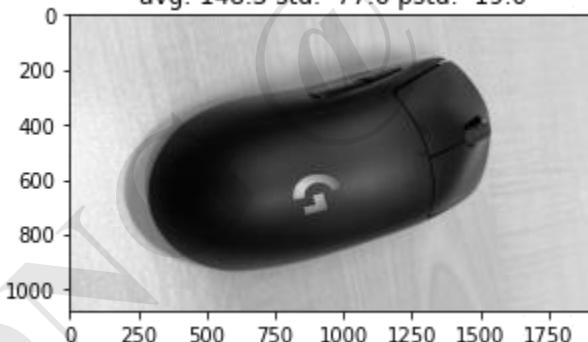
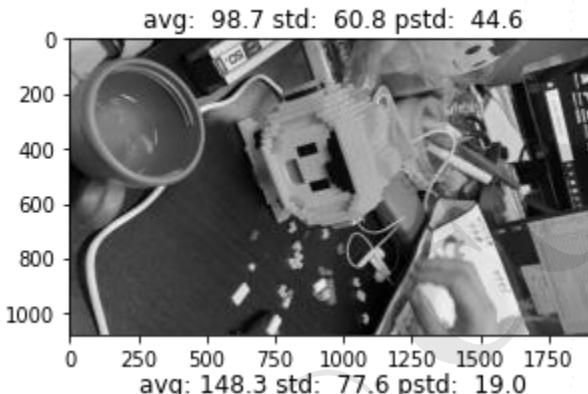
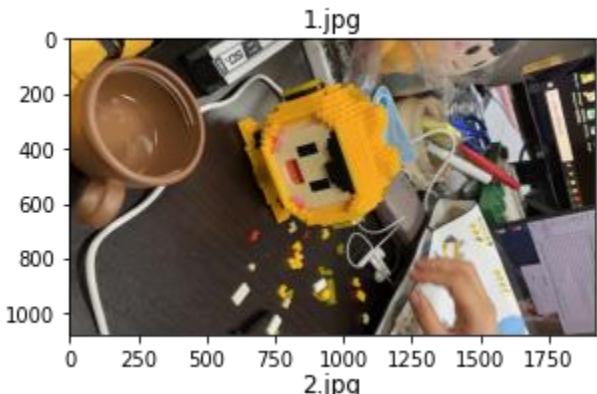


Color Histograms

A color histogram of an image represents the distribution of the composition of colors in the image. It shows different types of colors appeared and the number of pixels in each type of the colors appeared.

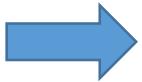


Examples from IMHere



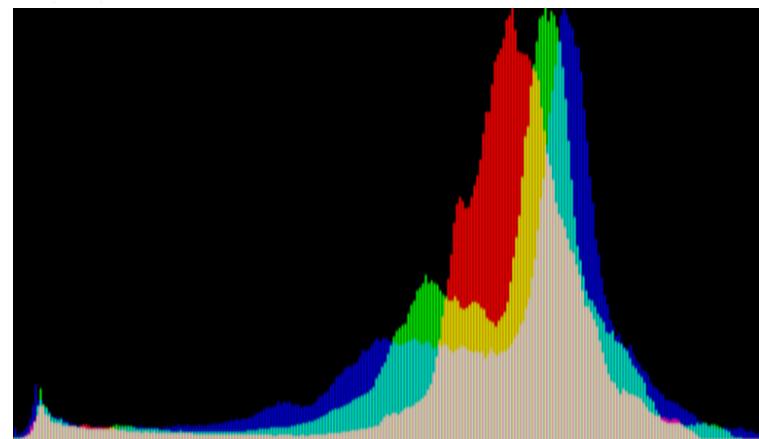
Color Histogram Extraction

To calculate the color histogram, the color space needs to be divided into several small color intervals, and each cell becomes a bin of the histogram. This process is called color quantization. Then, the color histogram can be obtained by calculating the number of pixels in each cell.



Red	Green	Blue	Pixel Count
Bin 0	Bin 0	Bin 0	7414
Bin 0	Bin 0	Bin 1	230
Bin 0	Bin 0	Bin 2	0
Bin 0	Bin 0	Bin 3	0
Bin 0	Bin 1	Bin 0	8
Bin 0	Bin 1	Bin 1	372
Bin 0	Bin 1	Bin 2	88
Bin 0	Bin 1	Bin 3	0
Bin 0	Bin 2	Bin 0	0
Bin 0	Bin 2	Bin 1	0
Bin 0	Bin 2	Bin 2	10
Bin 0	Bin 2	Bin 3	1
Bin 0	Bin 3	Bin 0	0
Bin 0	Bin 3	Bin 1	0

A picture of a cat



Color histogram of the cat picture with x-axis being RGB and y-axis being the frequency.

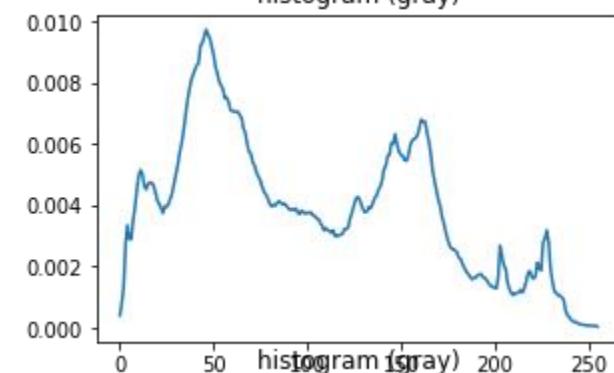
<https://www.cosy.sbg.ac.at/~pmeerw/Watermarking/index.html>

https://en.wikipedia.org/wiki/Color_histogram

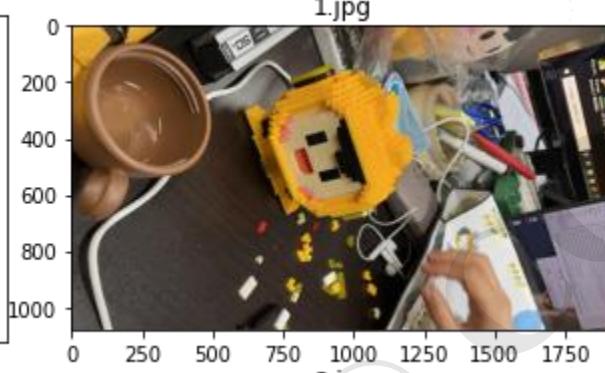
• • •

Examples from IMHere

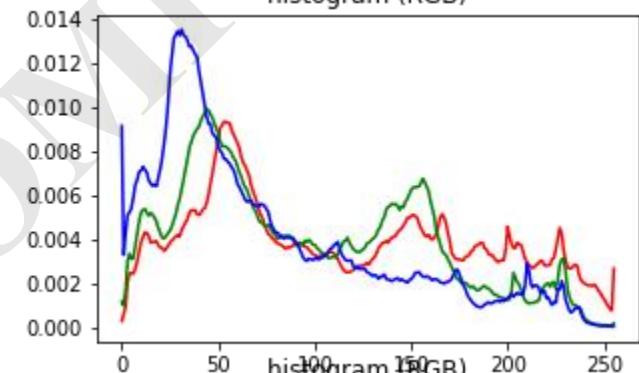
histogram (gray)



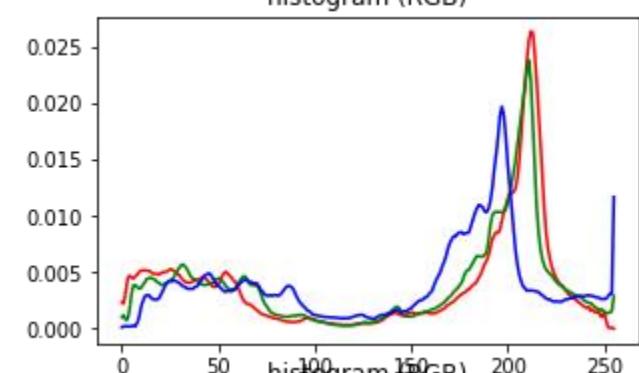
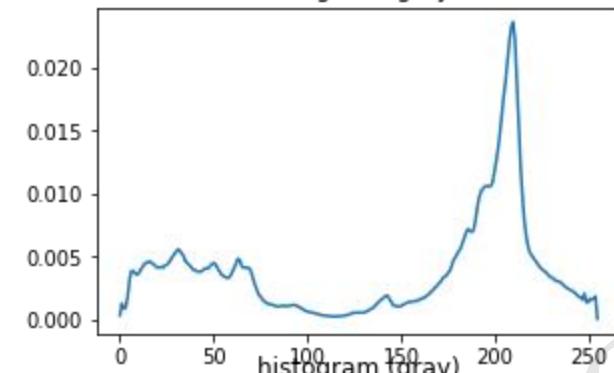
1.jpg



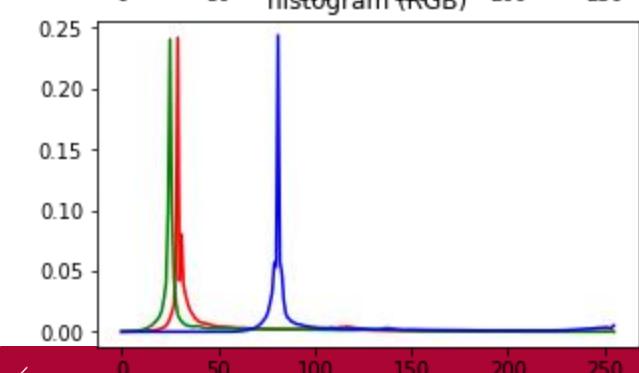
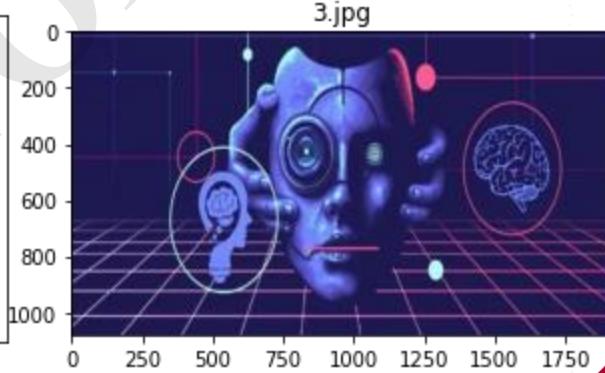
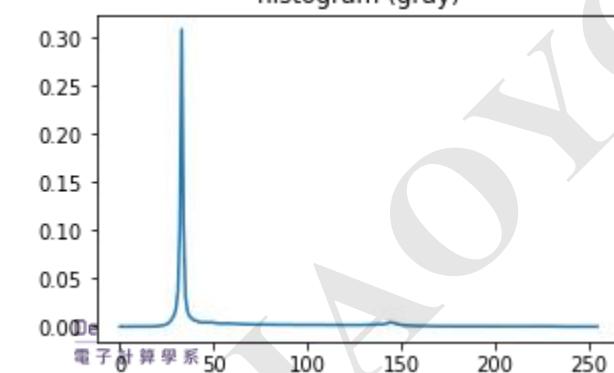
histogram (RGB)



histogram (gray)

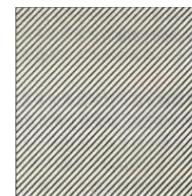
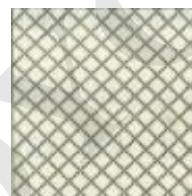
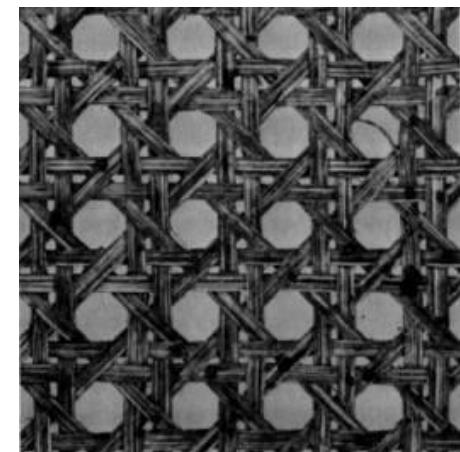
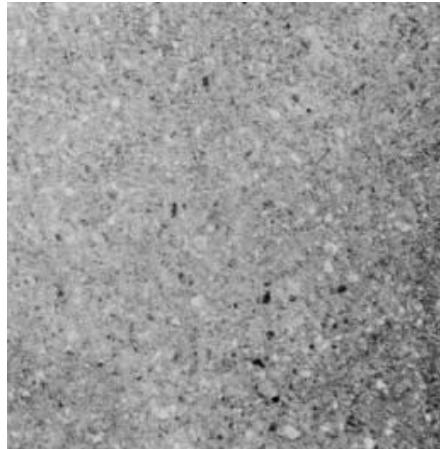


histogram (gray)



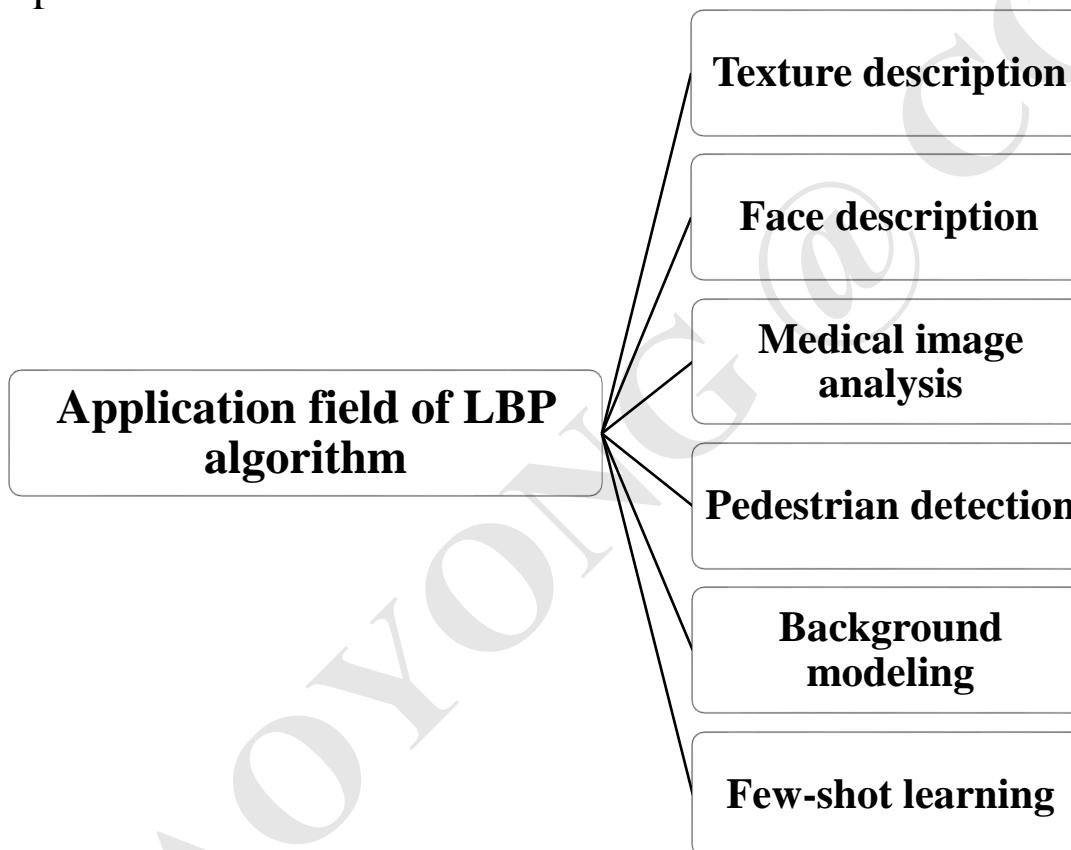
Texture Features

Texture feature refers to the common internal characteristics of the object surface, which contains the important information of the structure and arrangement of the object surface and its relationship with the surrounding objects.



Local Binary Patterns (LBP)

LBP algorithm was originally proposed by Timo Ojala and Matti pietikäinen of machine vision and signal analysis center of Oulu University in Finland in 1996 for texture feature description.



Matti Pietikäinen
IEEE、IAPR Fellow
CMVSA of Oulu University

Local Binary Patterns (LBP)

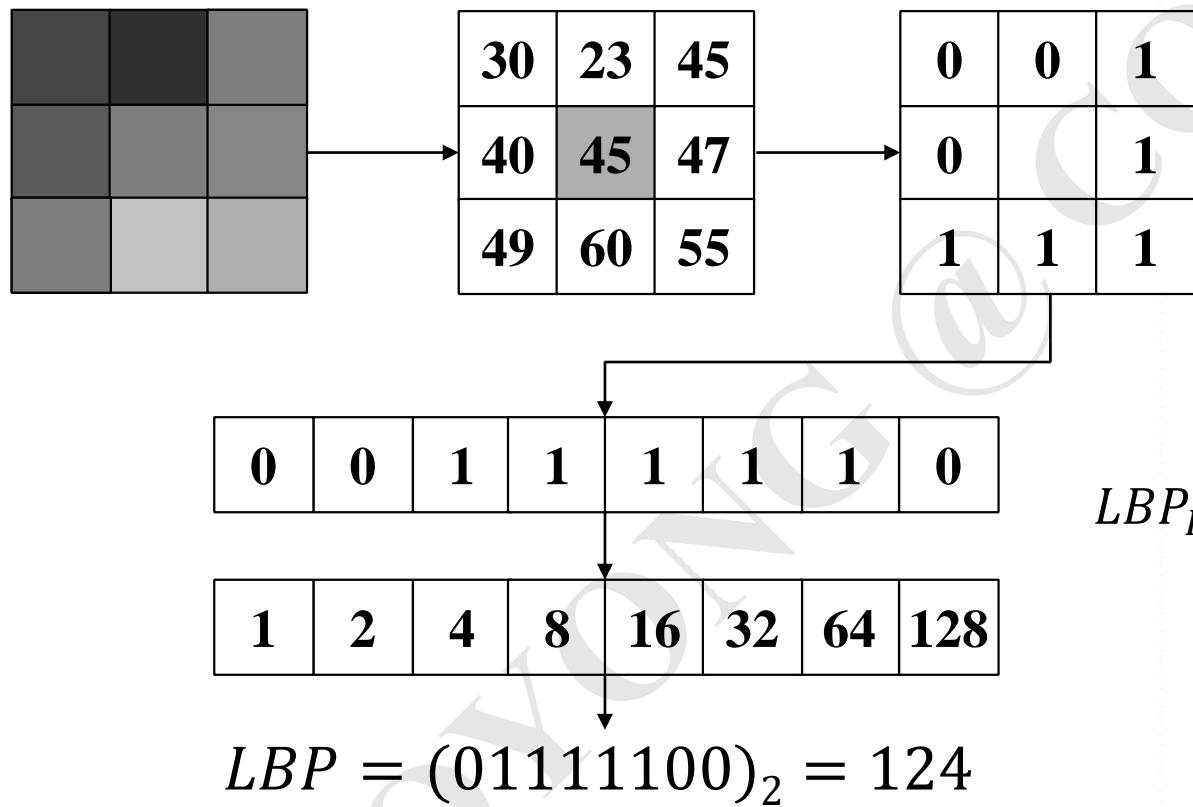


Gray-scale image



LBP feature map

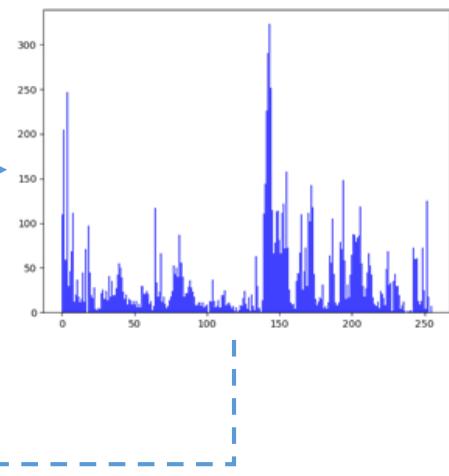
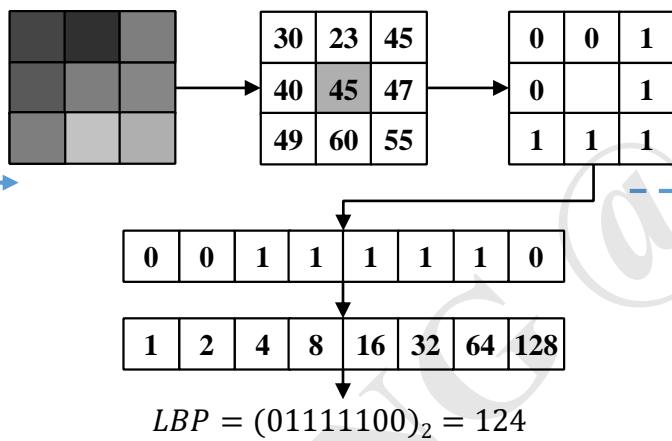
Local Binary Patterns (LBP)



$$LBP_{P,R} = \sum_{p=1}^P S(g_p - g_c) \times 2^{p-1}$$

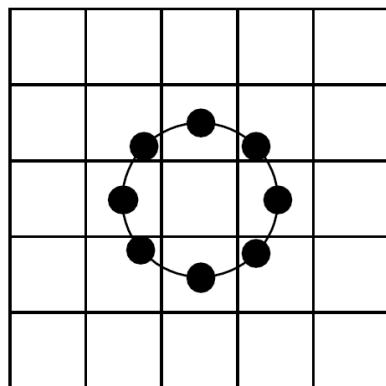
$$S(x) = \begin{cases} 1 & x \geq 0 \\ 0 & x < 0 \end{cases}$$

Local Binary Patterns (LBP)

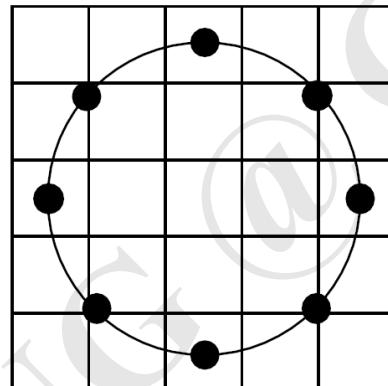


Feature vector $X = [x_1, x_2, x_3, \dots, x_K]$

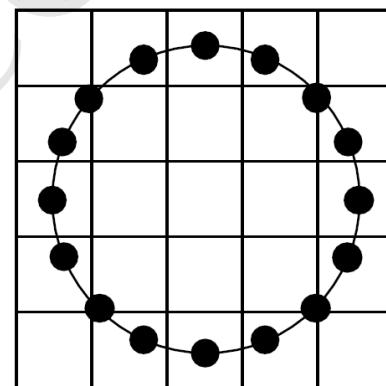
Circular LBP



R=1 P=8

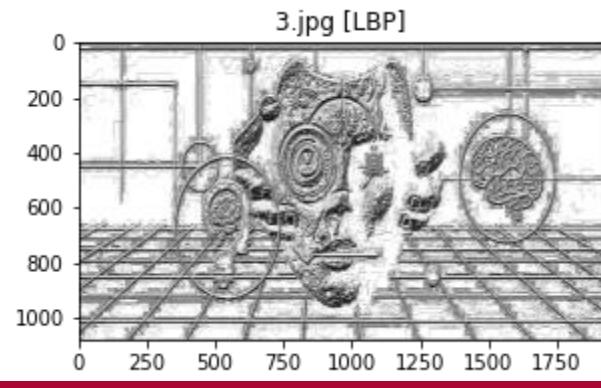
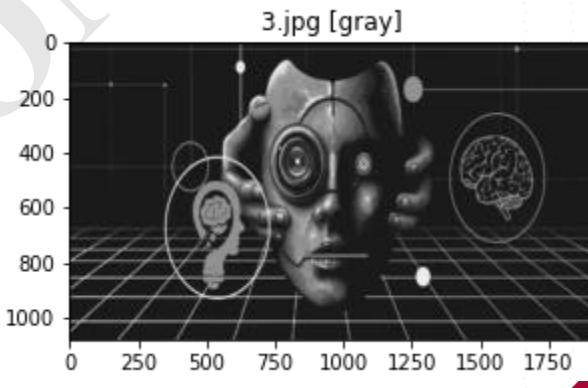
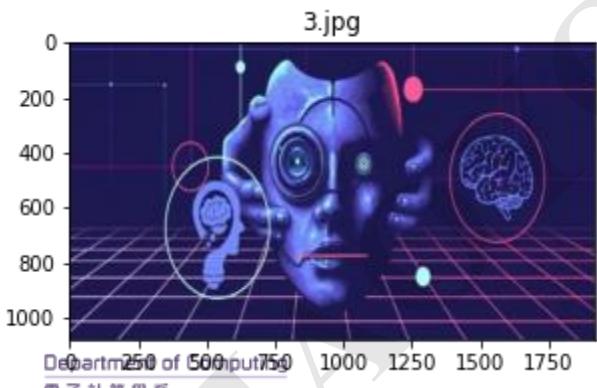
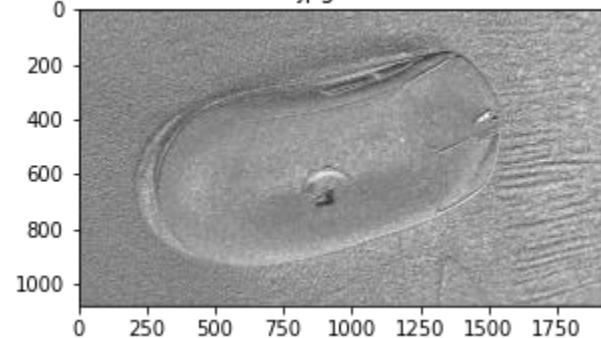
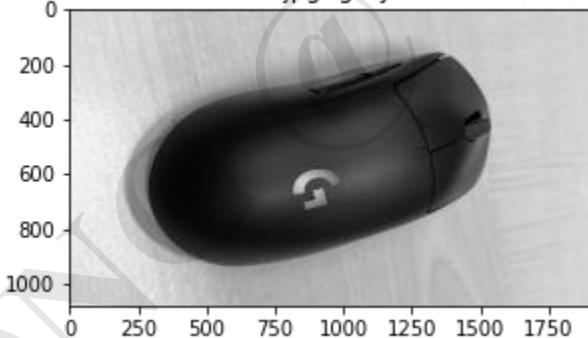
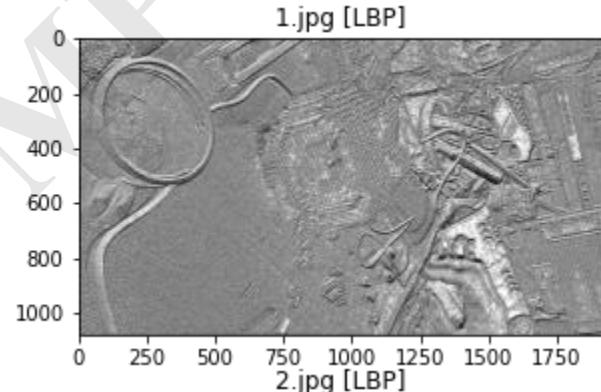
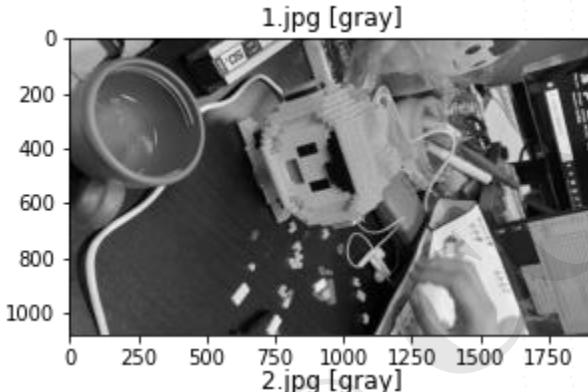
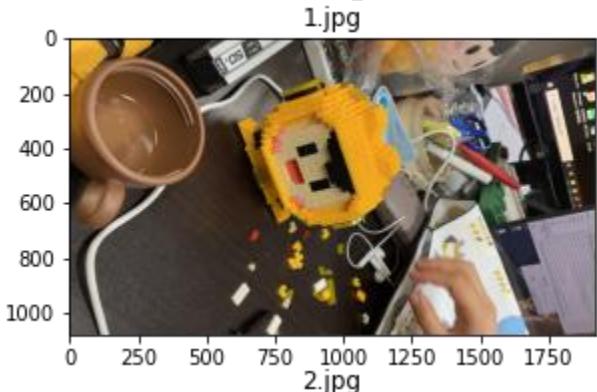


R=2 P=8



R=2 P=16

Examples from IMHere

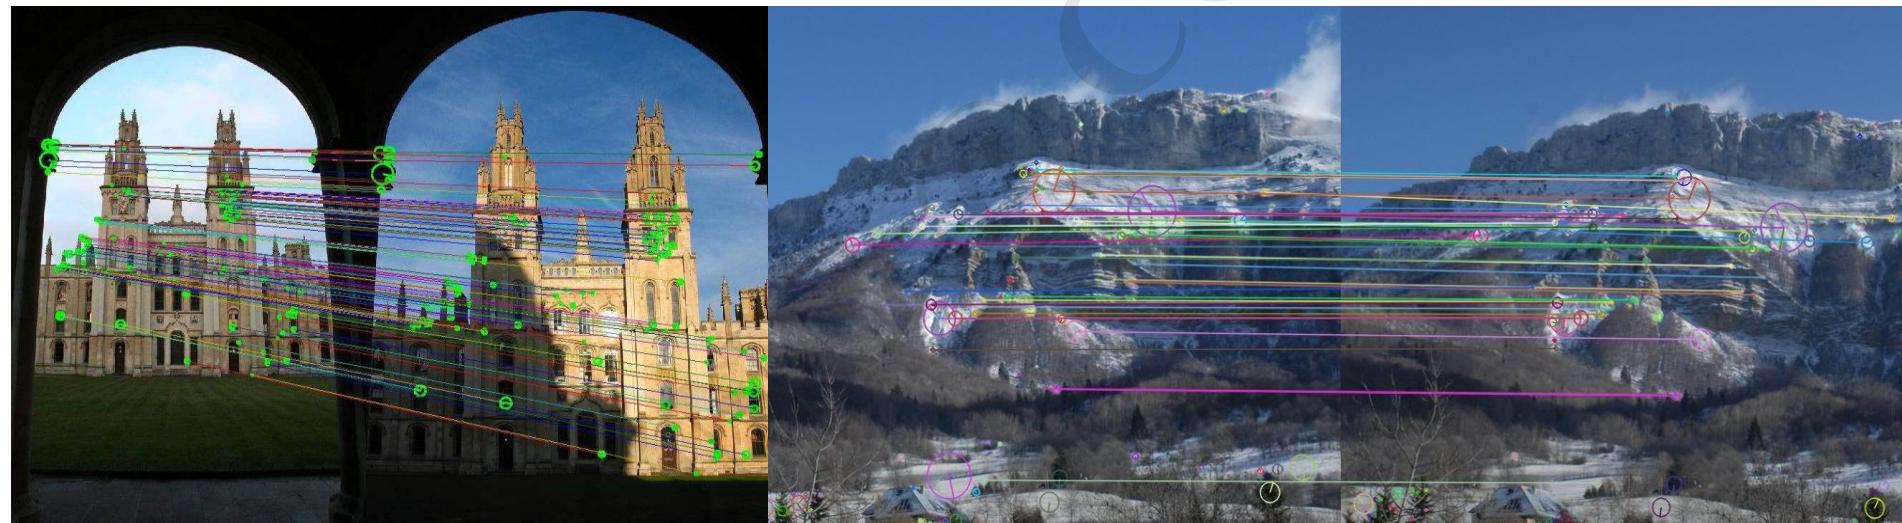


Can computers “see” the difference using color histograms or LBP?



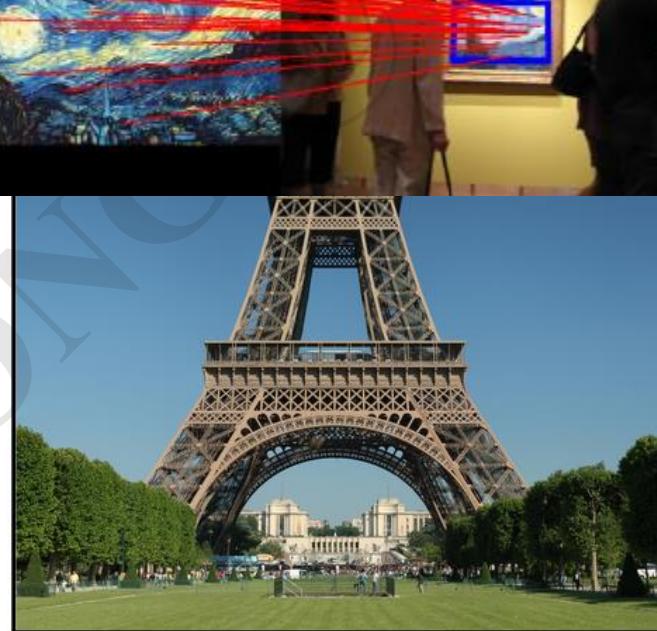
Scale-invariant feature transform (SIFT)

SIFT is a milestone work in the research field of local image feature descriptor. SIFT was first proposed in 1999 and was improved in 2004. SIFT is invariant to image changes such as scale, rotation, certain viewing angle and illumination changes.

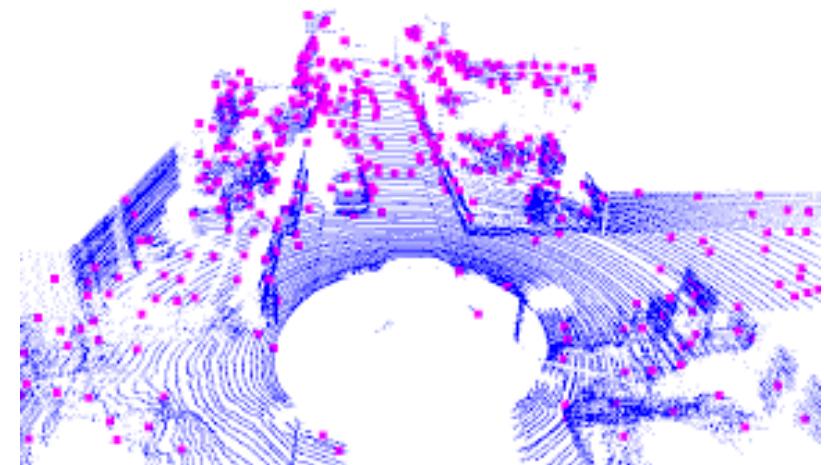
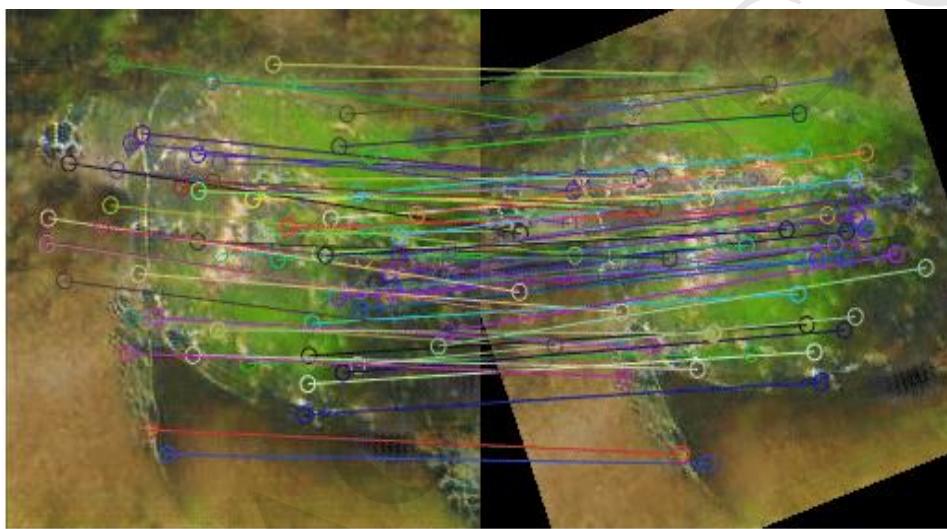
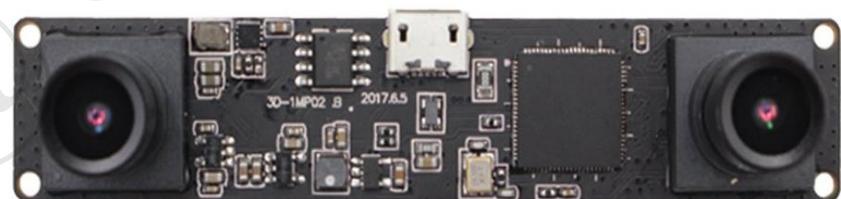
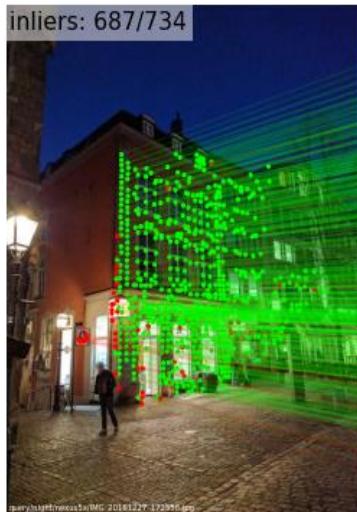


<http://5b0988e595225.cdn.sohucs.com/images/20200318/bacc3866c2eb47b58cde6af3b6512b6.jpeg>

Scale-invariant feature transform (SIFT)



Scale-invariant feature transform (SIFT)



Scale-invariant feature transform (SIFT)



SIFT Detection & Description Steps

- **Candidate Localization:** Find Scale-invariant candidates
- **Refinement:** Keypoint filtering
- **Orientation Assignment:** Estimate orientations for keypoints
- **Description Generation:** Build description

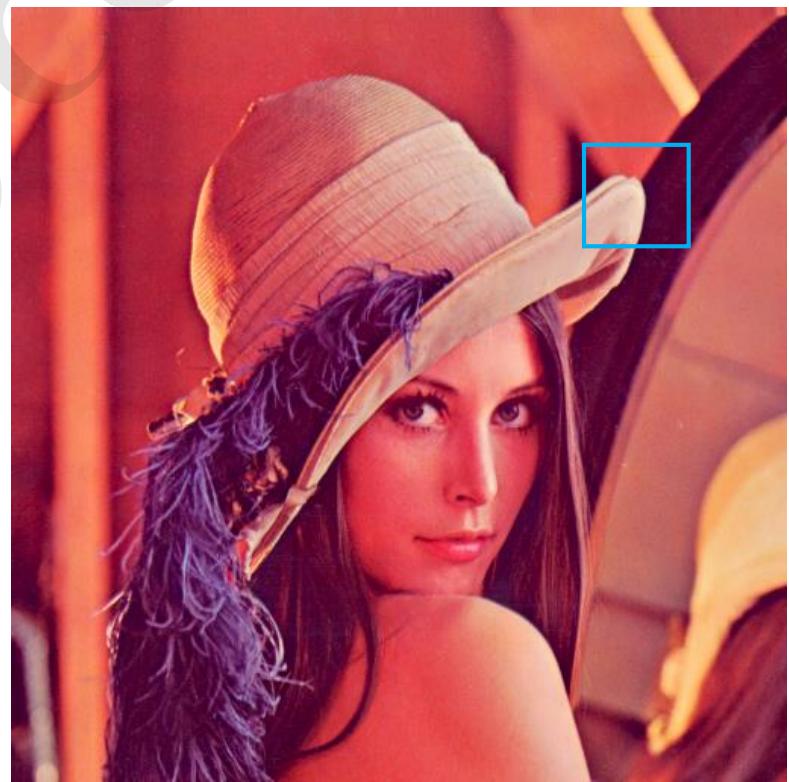
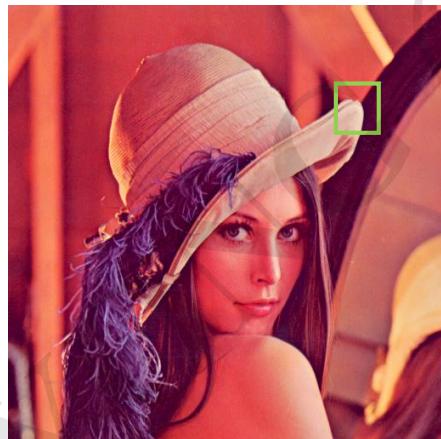
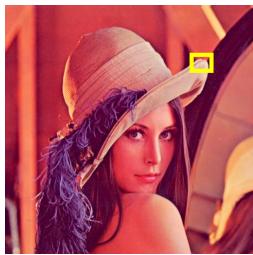
Step 1: Candidate Localization

>What is a scale?



Step 1: Candidate Localization

>What is a scale?



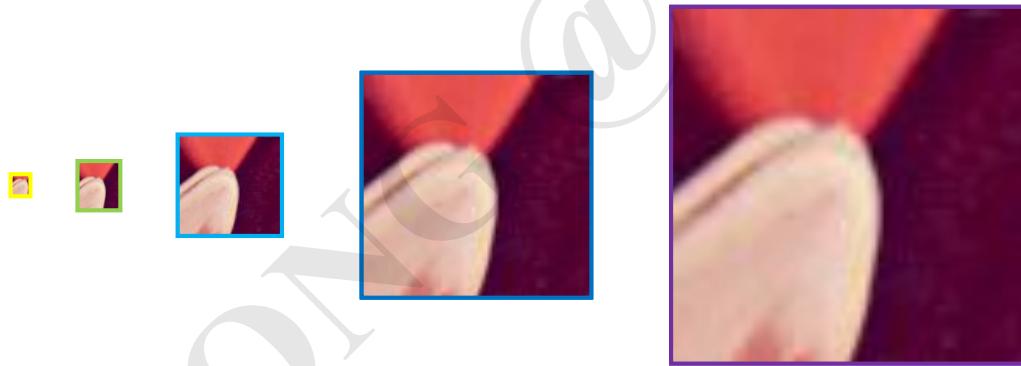
Step 1: Candidate Localization

>What is a scale?



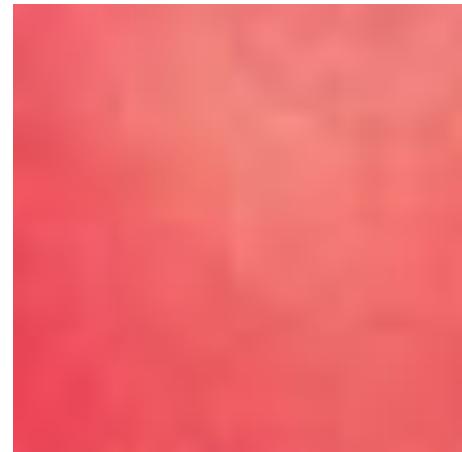
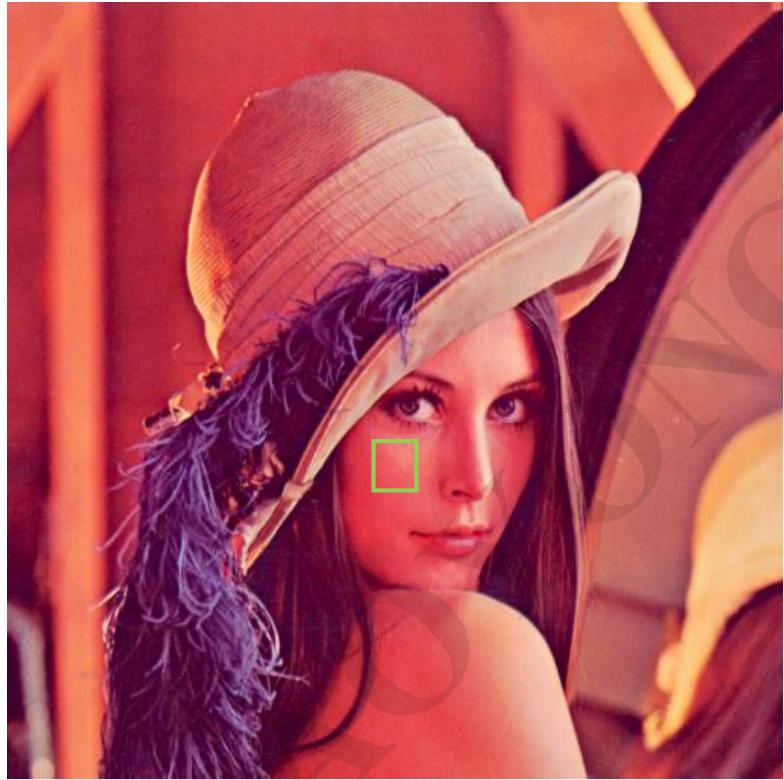
Step 1: Candidate Localization

>What is a scale?



Step 1: Candidate Localization

>What is a scale?



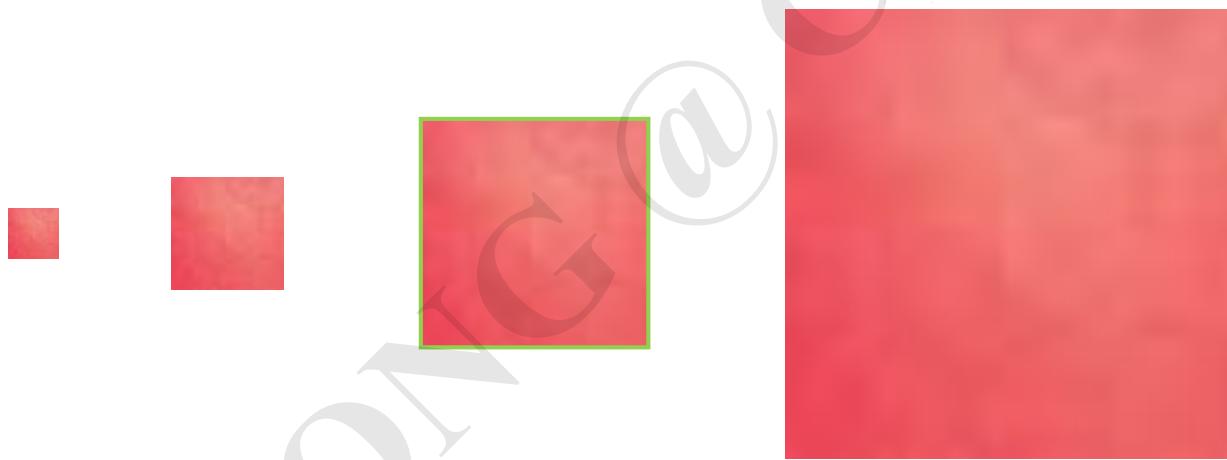
Step 1: Candidate Localization

>What is scale-invariant?



Step 1: Candidate Localization

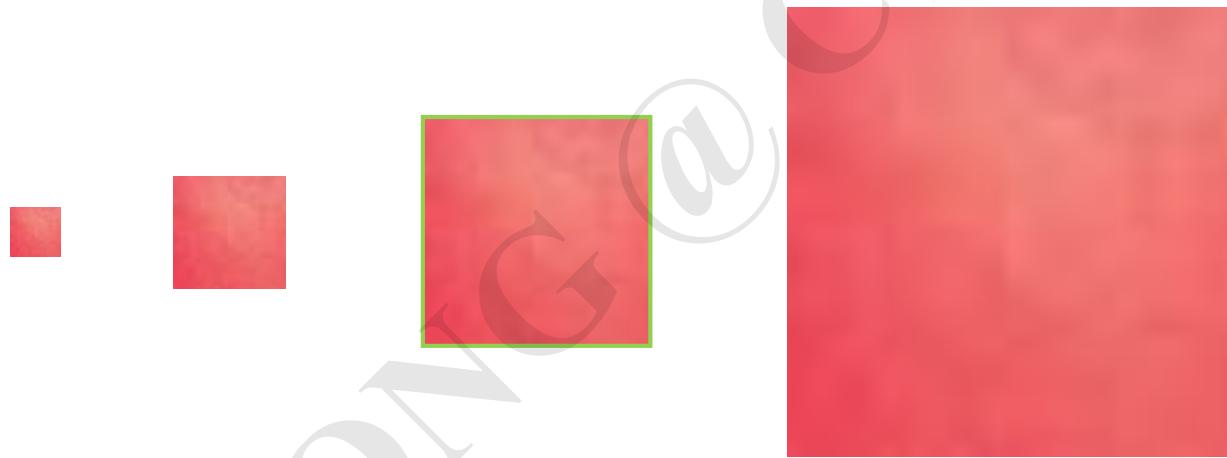
> What is scale-invariant?



To find the scale at which a pattern is represented
the **best** at its center (the key point)

Step 1: Candidate Localization

> What is scale-invariant?

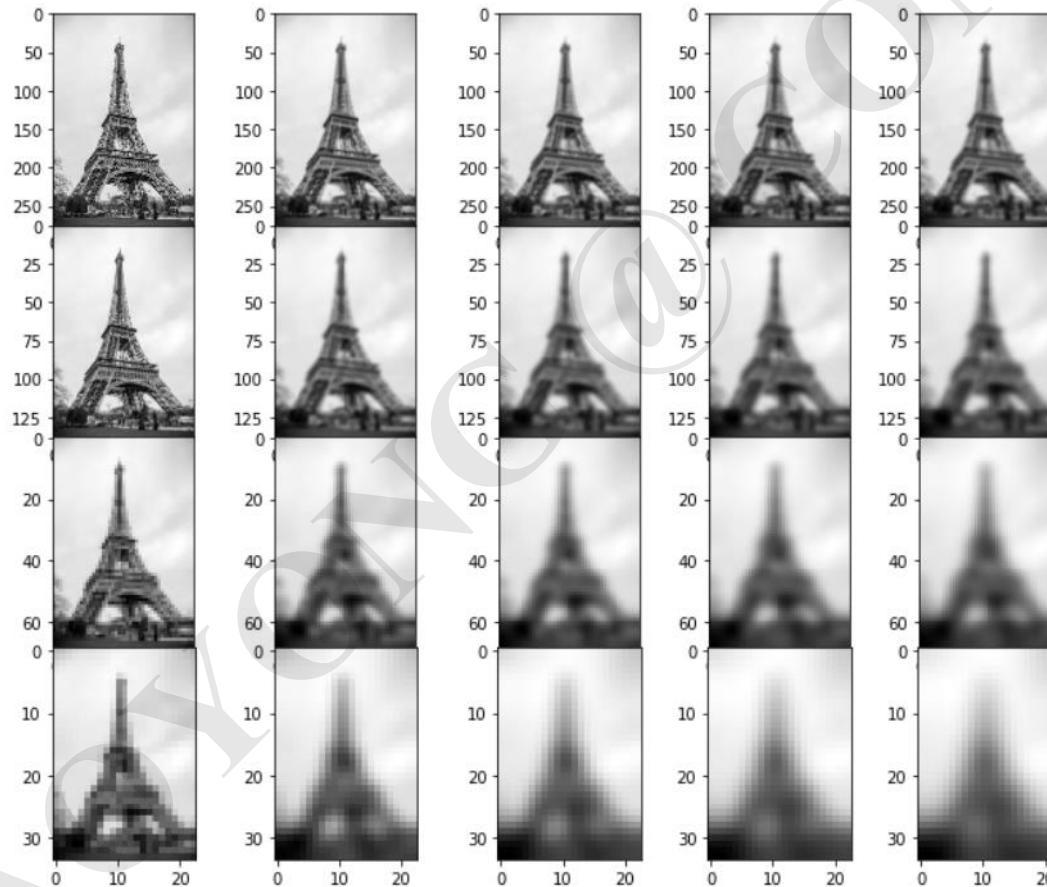


Scale-invariance: no matter what the original scale of the image is, the same **best scale** will be found

How can we do it?

Step 1: Candidate Localization

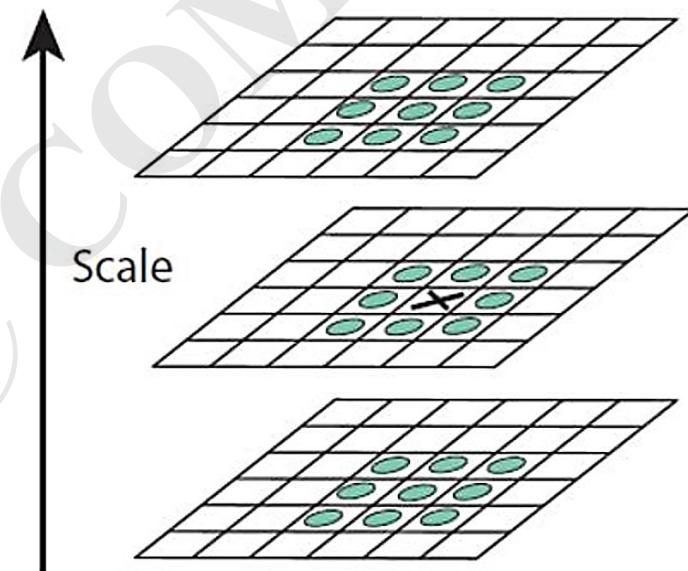
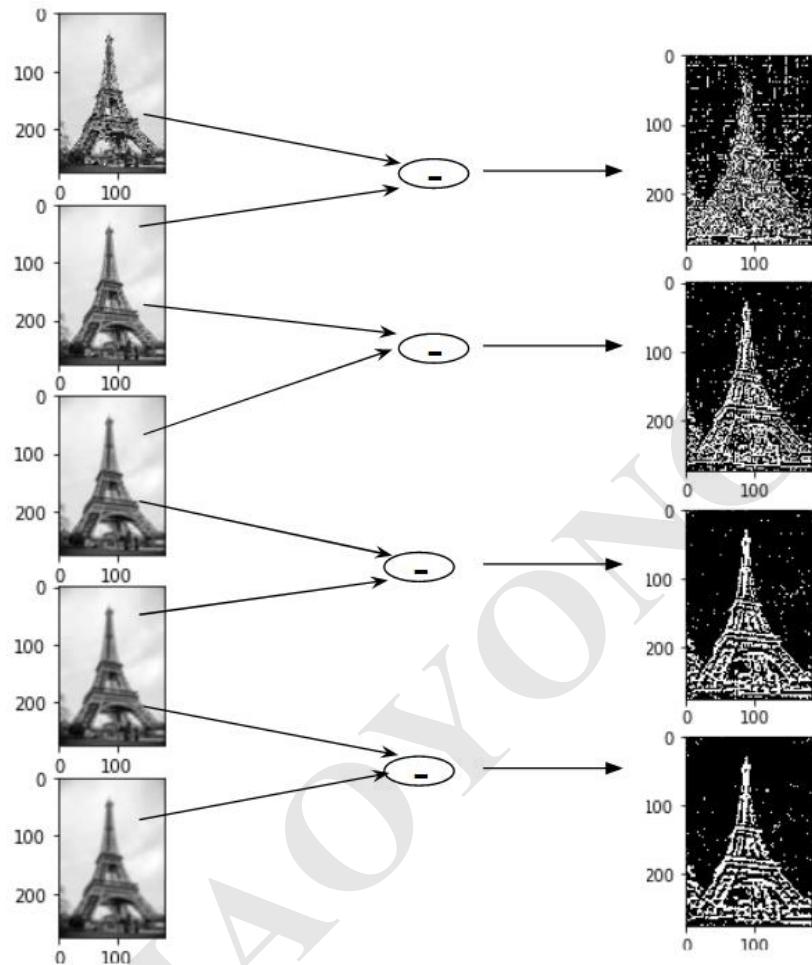
>Rescale and blur



Credit: <https://www.analyticsvidhya.com/blog/2019/10/detailed-guide-powerful-sift-technique-image-matching-python/>

Step 1: Candidate Localization

> Difference of Gaussian

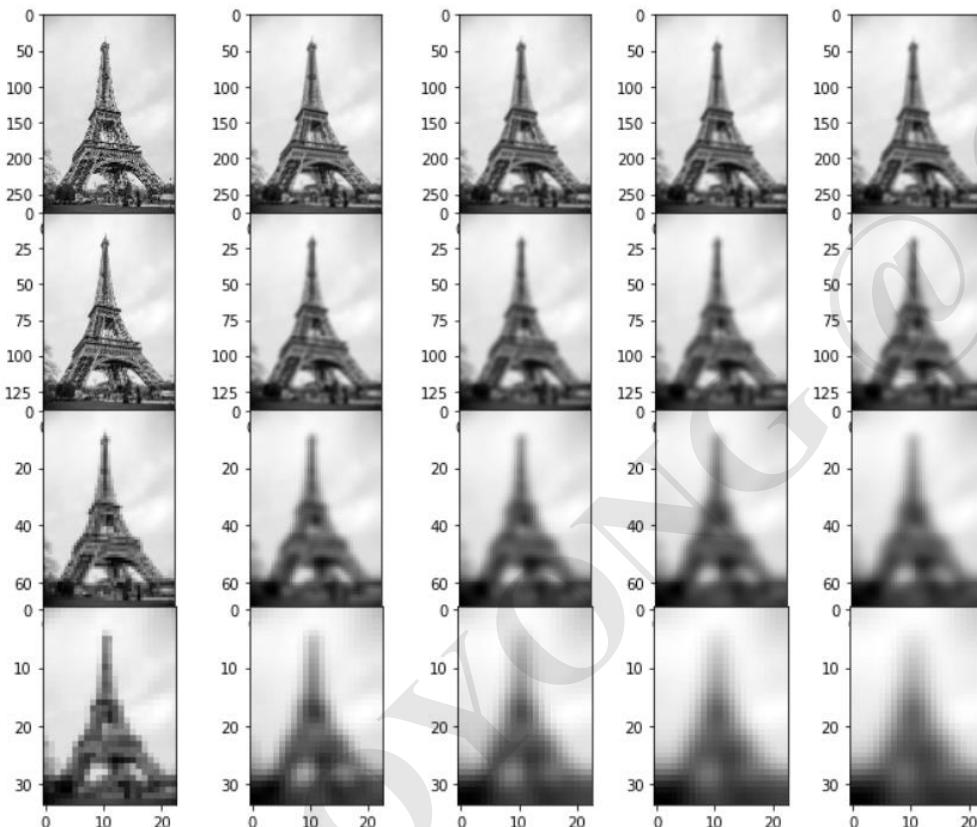


Local maxima and minima are those stand out at neighborhood

Sounds good but less efficient.
We need a little math here.

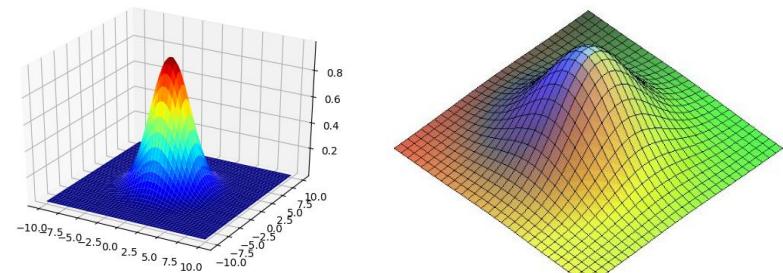
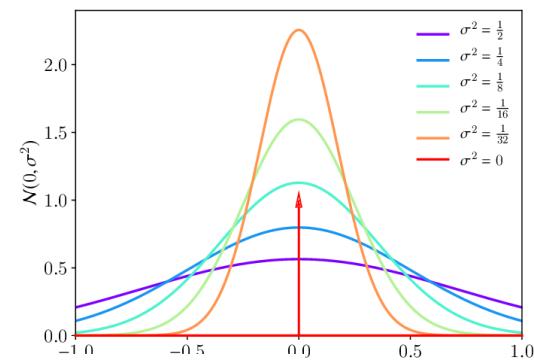
Step 1: Candidate Localization

> Rescale and blur



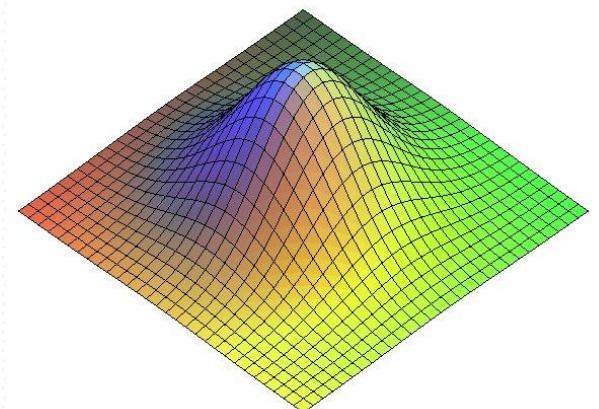
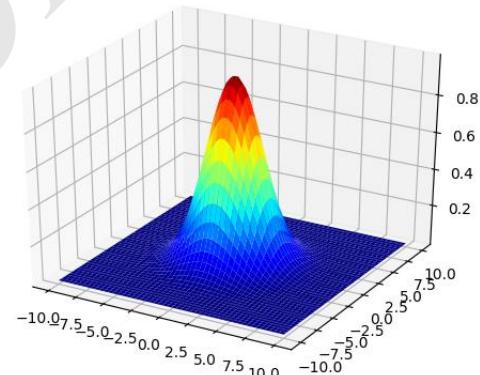
$$G(x, y, \sigma) * I(x, y)$$

$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$



Step 1: Keypoint Localization

0.0049	0.0092	0.0134	0.0152	0.0134	0.0092	0.0049
0.0092	0.0172	0.0250	0.0283	0.0250	0.0172	0.0092
0.0134	0.0250	0.0364	0.0412	0.0364	0.0250	0.0134
0.0152	0.0283	0.0412	0.0467	0.0412	0.0283	0.0152
0.0134	0.0250	0.0364	0.0412	0.0364	0.0250	0.0134
0.0092	0.0172	0.0250	0.0283	0.0250	0.0172	0.0092
0.0049	0.0092	0.0134	0.0152	0.0134	0.0092	0.0049



This is a 7×7 (truncated) Gaussian mask with a mean of zero and a standard deviation $\sigma = 2$

Step 1: Keypoint Localization



Scale = 0



Scale = 1



Scale = 4



Scale = 16



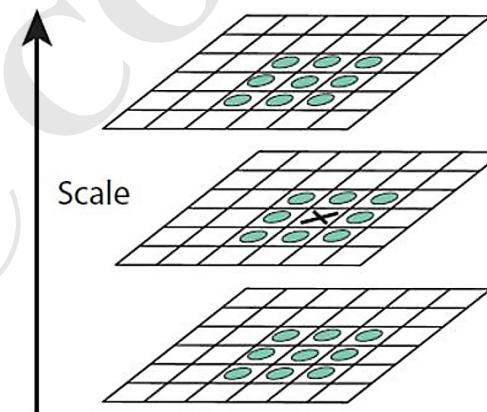
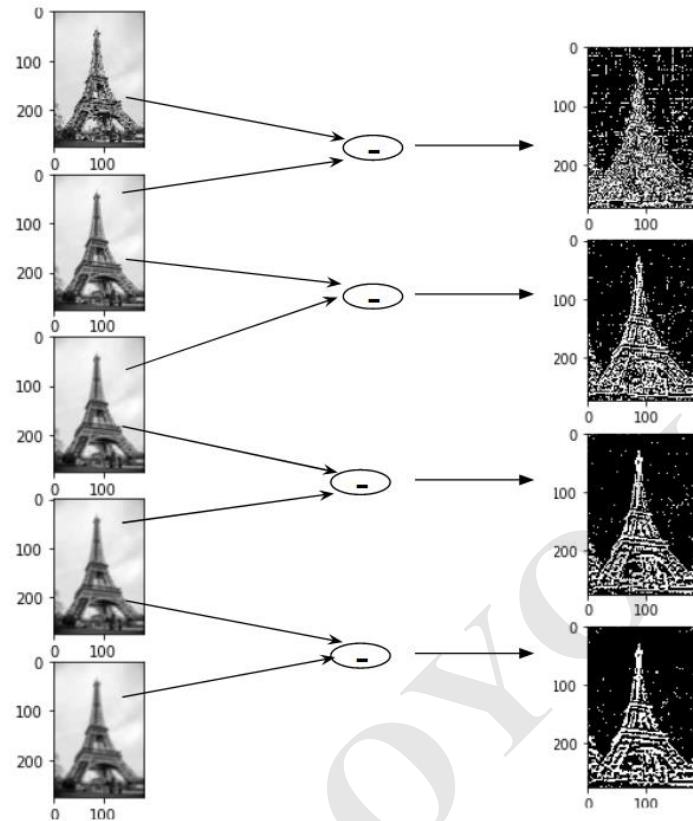
Scale = 64



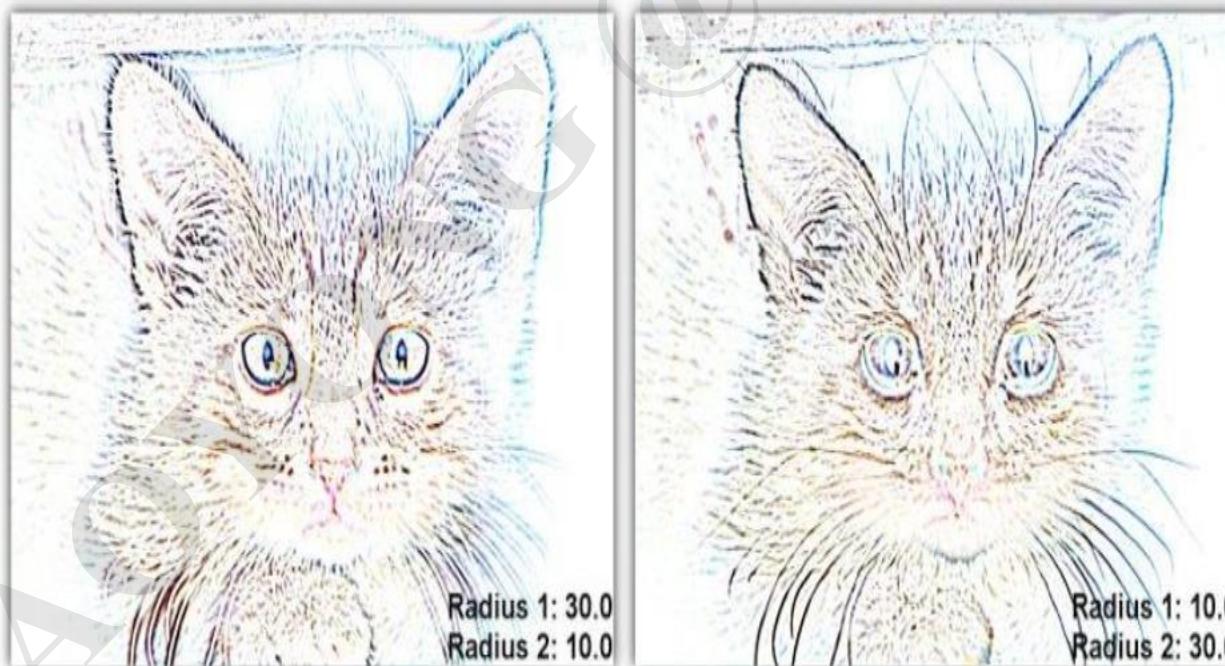
Scale = 256

Step 1: Candidate Localization

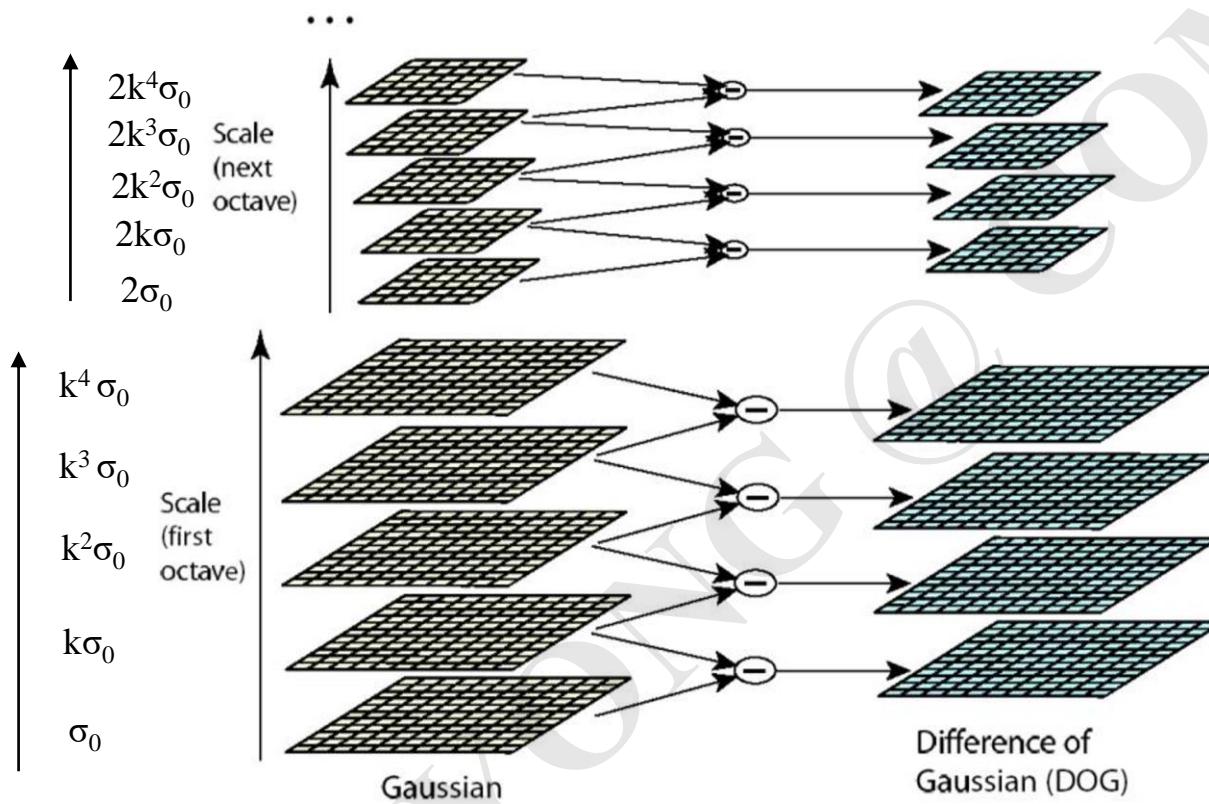
> Difference of Gaussian



$$\begin{aligned}
 L(x, y, \sigma) &= G(x, y, \sigma) * I(x, y) \\
 D(x, y, \sigma) &= L(x, y, \sigma) - L(x, y, k\sigma) \\
 &= (G(x, y, \sigma) - G(x, y, k\sigma)) * I(x, y)
 \end{aligned}$$



Step 1: Candidate Localization



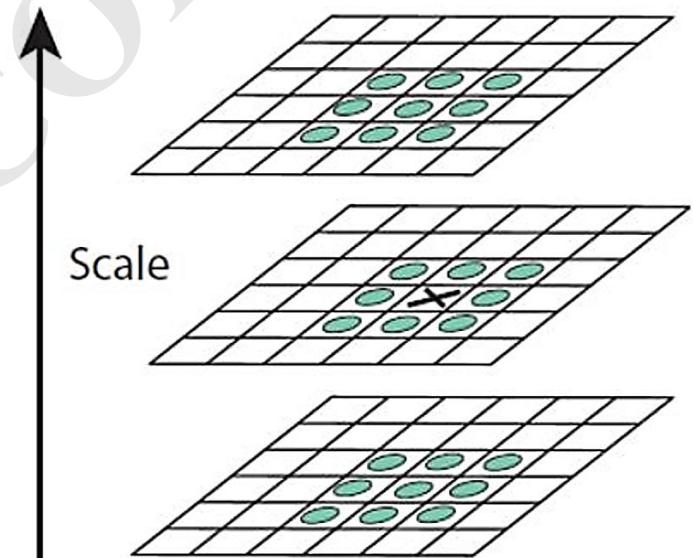
$$D(x, y, \sigma) = L(x, y, k\sigma) - L(x, y, \sigma)$$

$$L(x, y, \sigma) = G(x, y, \sigma)^* I(x, y)$$

Such a sequence of images convolved with Gaussians of increasing σ_0 constitute a so-called **scale space**.

Step 1: Candidate Localization

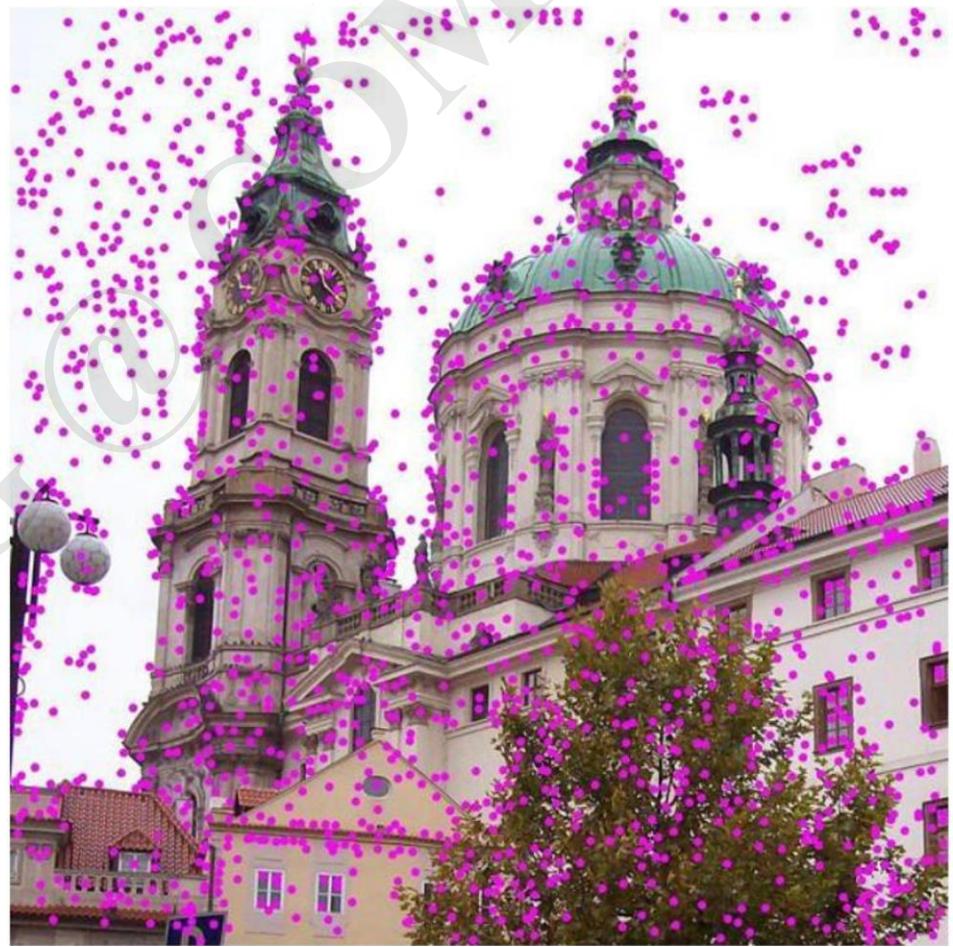
- The keypoints are maxima or minima in the “scale-space-pyramid”.
- Maxima and minima of the difference-of-Gaussian images are detected by comparing a pixel (marked with X) to its 26 neighbors in 3×3 regions at the current and adjacent scales (marked with circles)



Lowe, D.G. Distinctive Image Features from Scale-Invariant Keypoints. International Journal of Computer Vision 60, 91–110 (2004). <https://doi.org/10.1023/B:VISI.0000029664.99615.94>

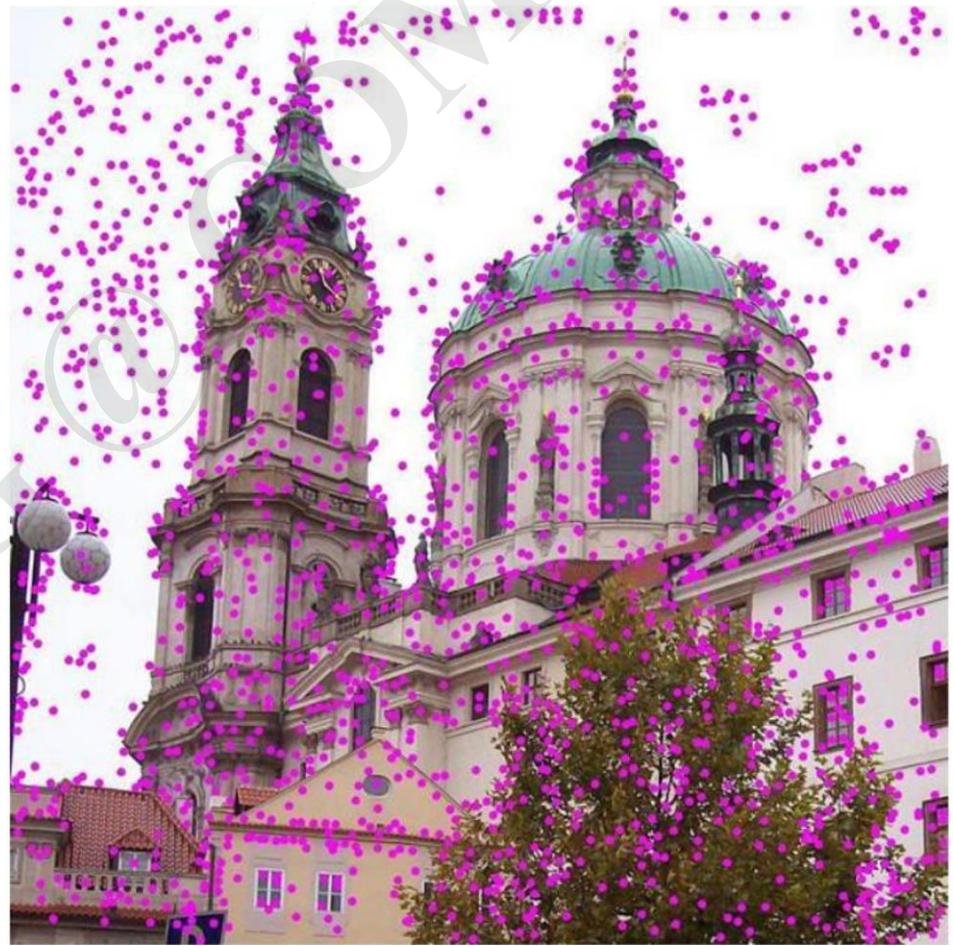
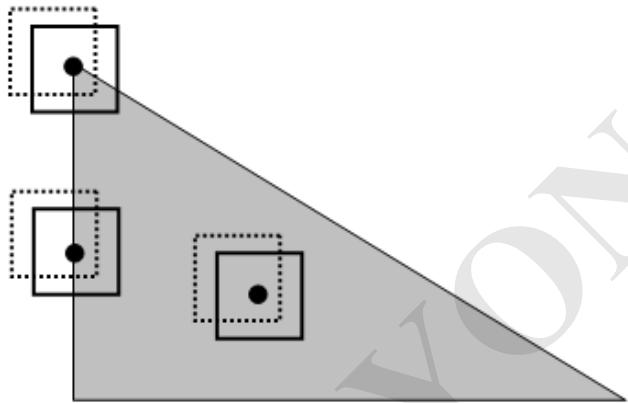
Step 1: Candidate Localization

Keypoint candidates



Step 2: Refinement

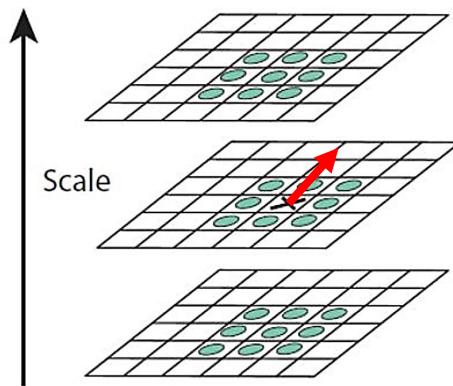
Idea: keypoint candidates at low-contrast positions or on edges are not stable enough. We have to remove them.



Point-wise calculation is straightforward but inefficient.
We need to introduce a little math again

Step 2: Refinement

Let us denote the DoG function in a small 3D neighborhood around a keypoint by a second-order Taylor-series:



$$\mathbf{x} = (x, y, \sigma)^T$$

Offset to the origin rather than the coordinates

$$D(\mathbf{x}) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$

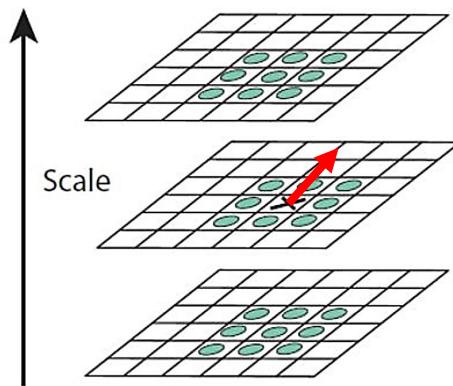
Gradient
Vector

Hessian
Matrix

By setting this to zero, we will have a \hat{x} , at which we can find the **theoretical** local extremum.
We discard current candidate if $\hat{x} > 0.5$. Why?

Step 2: Refinement

Let us denote the DoG function in a small 3D neighborhood around a keypoint by a second-order Taylor-series:



$$\mathbf{x} = (x, y, \sigma)^T$$

Offset to the origin rather than the coordinates

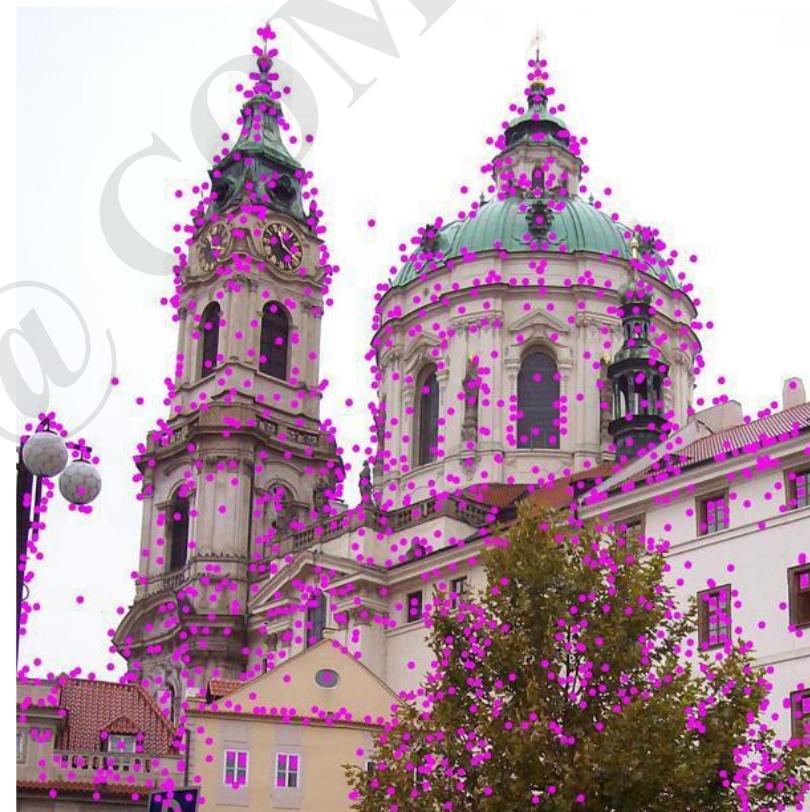
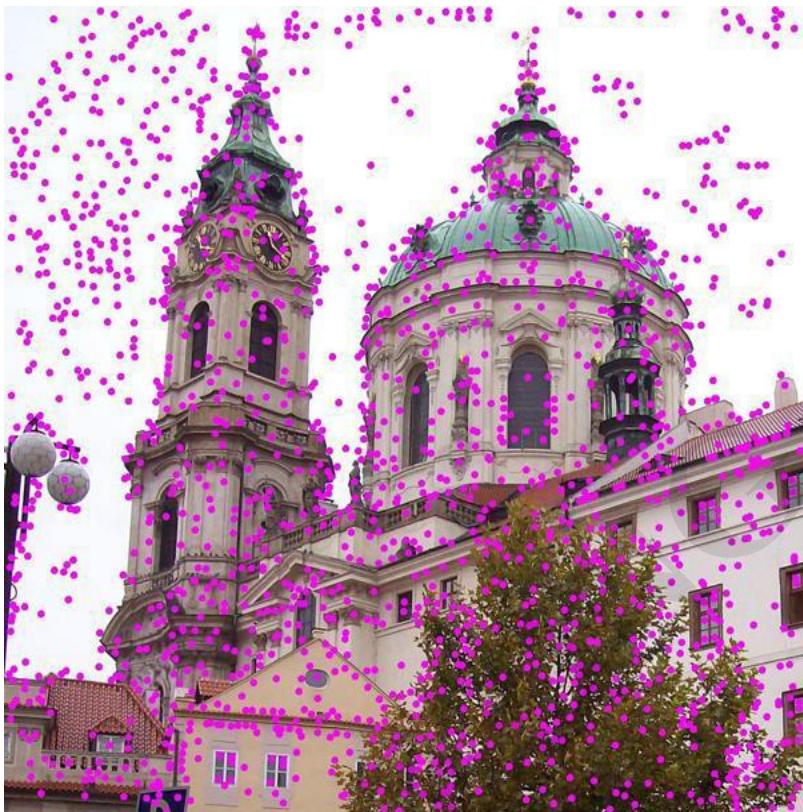
$$D(\mathbf{x}) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$

Gradient
Vector

Hessian
Matrix

To remove the low-contrast candidates, we calculate the second order Taylor at \hat{x} and discard the candidate if the result less than 0.03. Why?

Step 2: Refinement

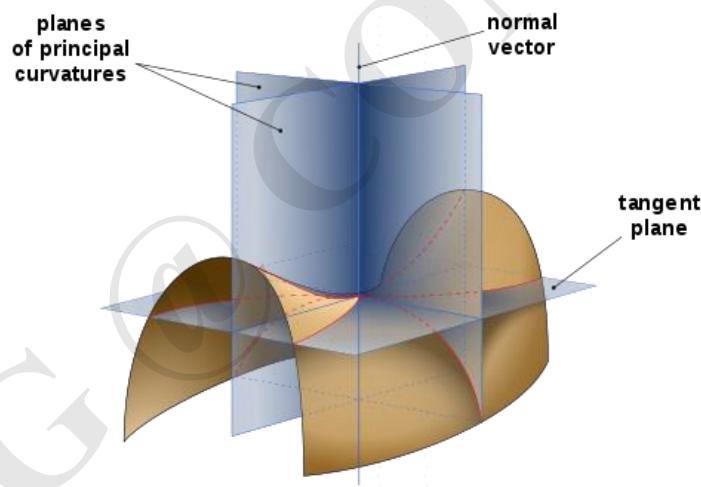


Removal of low-contrast keypoints

Step 2: Refinement

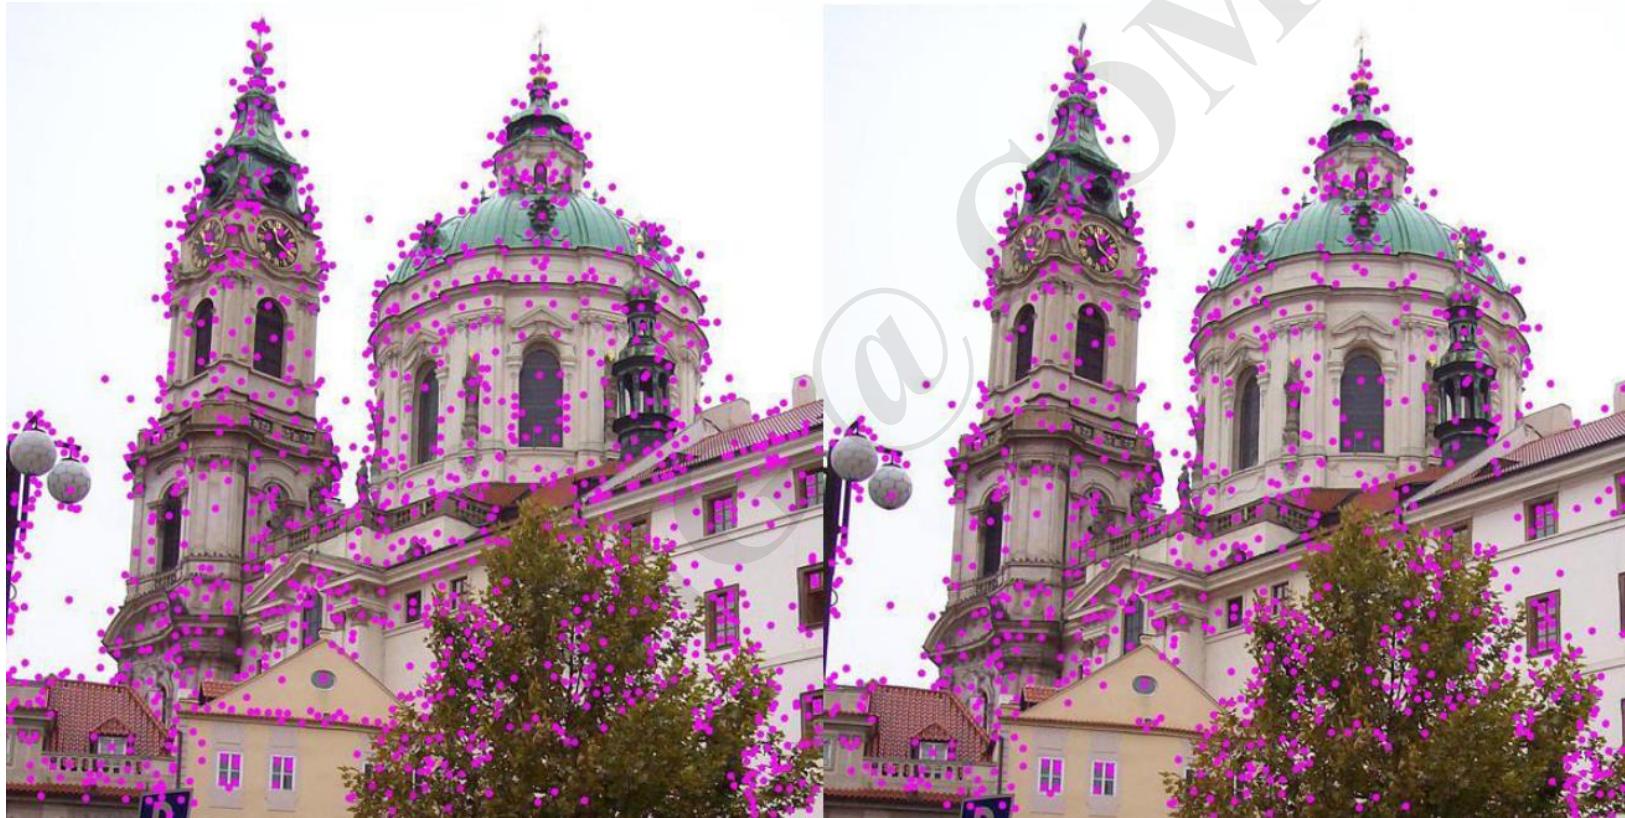
$$\mathbf{H} = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$$

Hessian
Matrix



At an edge candidate, the principal curvature across the edge is much larger than that along the edge. To remove the edge candidates, we calculate **Hessian Matrix**. The eigenvalues of \mathbf{H} are proportional to the principal curvatures of D .

Step 2: Refinement



Removal of high-contrast keypoints residing on edges

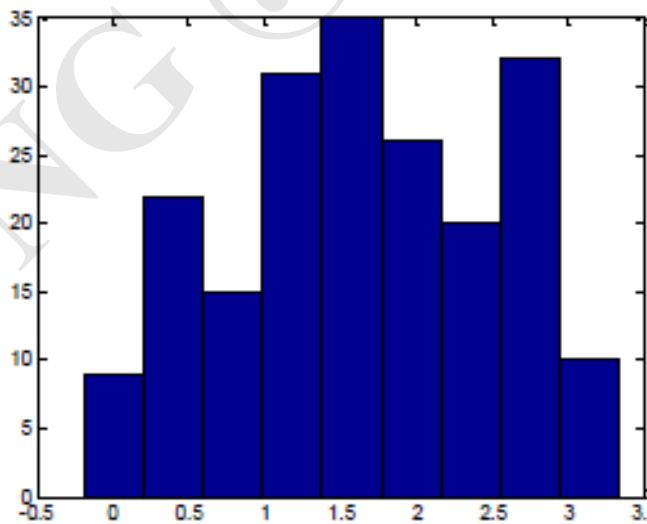
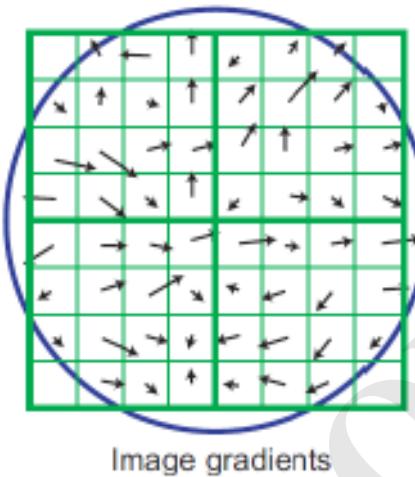
Step 3: Orientation Assignment

- Compute the gradient magnitudes and orientations in a small window around the keypoint – at the appropriate scale.

$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$$

$$m(x, y) = \sqrt{(L(x + 1, y) - L(x - 1, y))^2 + (L(x, y + 1) - L(x, y - 1))^2}$$

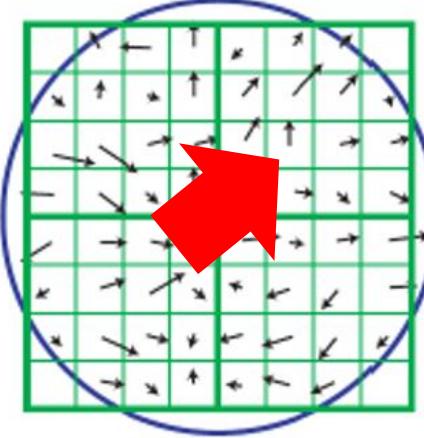
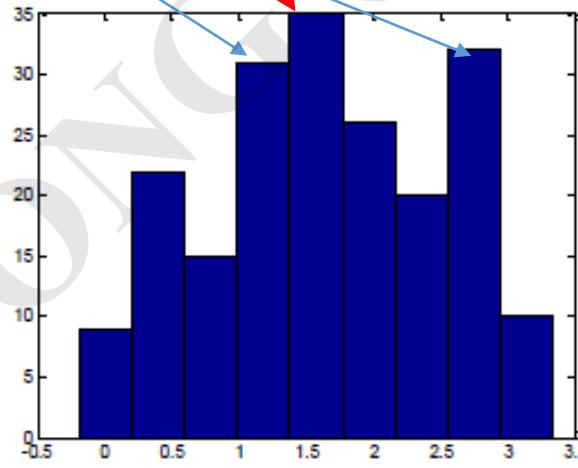
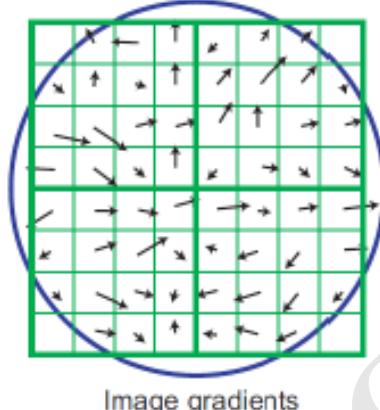
$$\theta(x, y) = \tan^{-1}((L(x, y + 1) - L(x, y - 1)) / (L(x + 1, y) - L(x - 1, y)))$$



Histogram of gradient orientation – the bin-counts are weighted by gradient magnitudes and a Gaussian weighting function. Usually, 36 bins are chosen for the orientation.

Step 3: Orientation Assignment

- Assign the **dominant orientation** as the orientation of the keypoint.
- In case of multiple peaks or histogram entries more than 0.8 times the peak, create a **separate** descriptor for each orientation (they will all have the same scale and location)



Step 3: Orientation Assignment



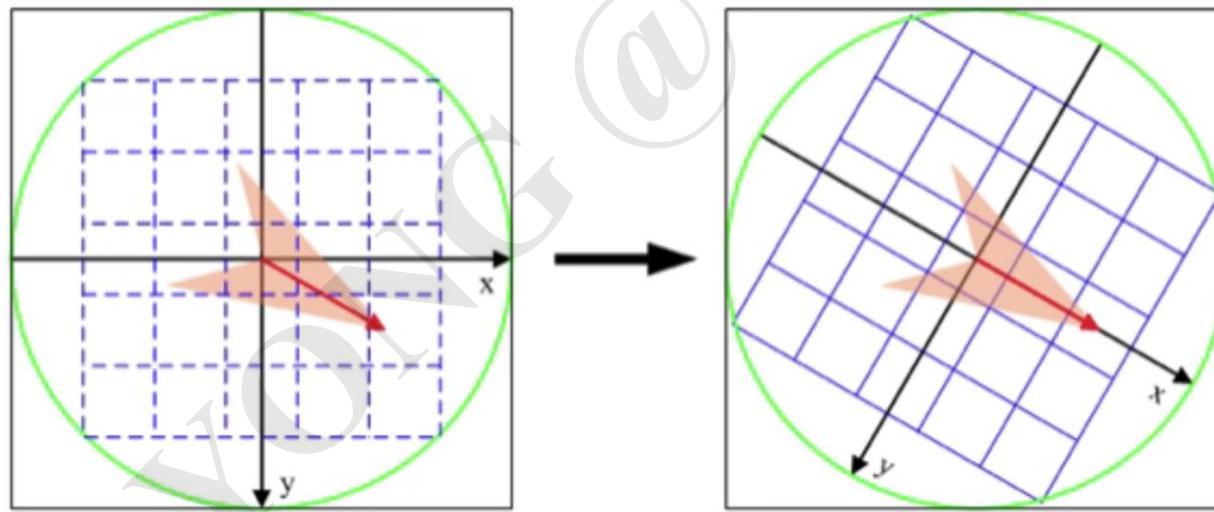
Original image



Keypoints and their orientations

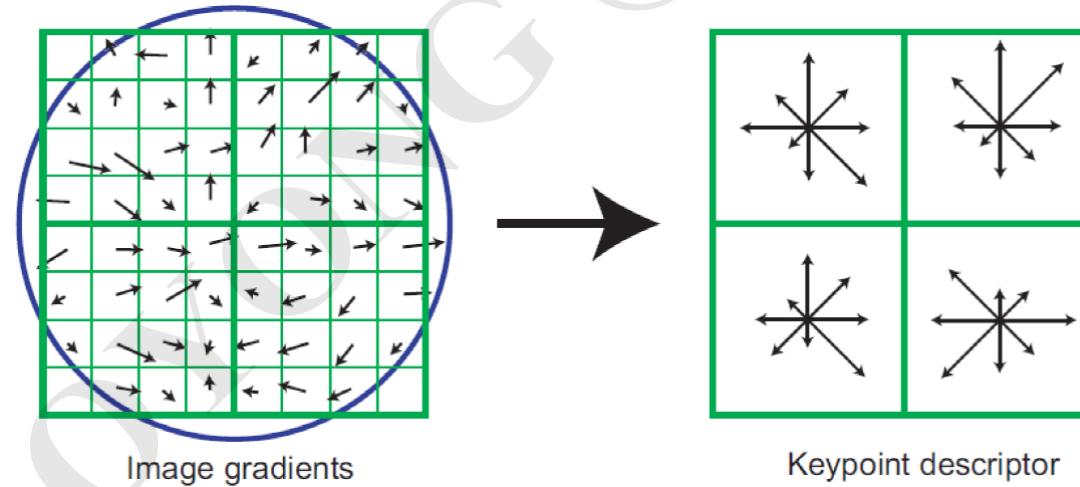
Step 4: Build descriptions

- In order to achieve orientation invariance, the coordinates of the descriptor and the gradient orientations are rotated relative to the keypoint orientation



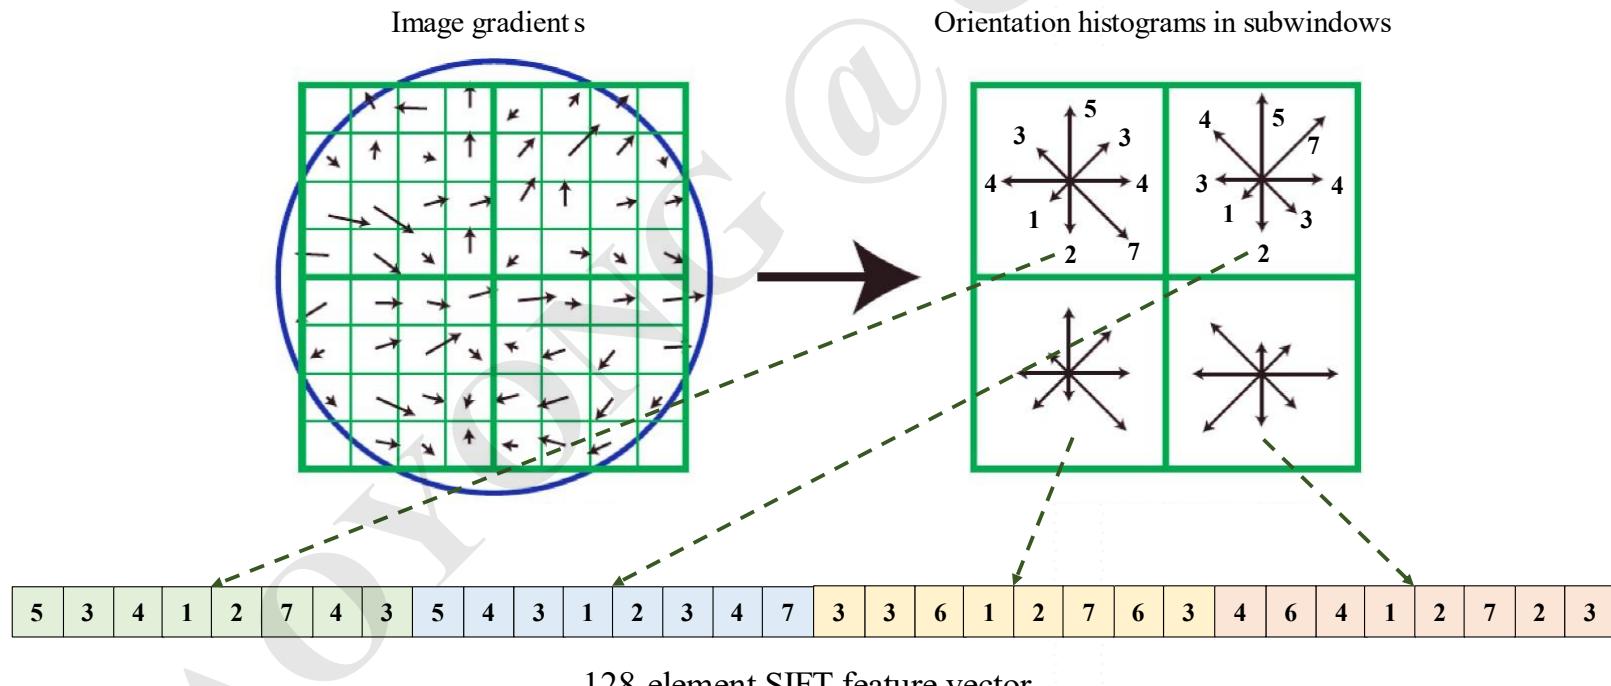
Step 4: Build descriptions

- Consider a small region around the keypoint. Divide it into $n \times n$ cells (for example, $n = 2$, Each cell is of size 4 x 4).
- Build a **gradient orientation histogram** in each cell. Each histogram entry is weighted by the *gradient magnitude* and a *Gaussian weighting function* with $\sigma = 0.5$ times window width.



Step 4: Build descriptions

- We now have a descriptor of size rn_2 if there are r bins in the orientation histogram.
- Typical case used in the SIFT paper: $r = 8, n = 4$, so length of each descriptor is 128.



Lowe, D.G. Distinctive Image Features from Scale-Invariant Keypoints. International Journal of Computer Vision 60, 91–110 (2004).
<https://doi.org/10.1023/B:VISI.0000029664.99615.94>

Step 4: Build descriptions

- To achieve scale-invariance, the size of the window should be adjusted as per scale of the keypoint. Larger scale = larger window.

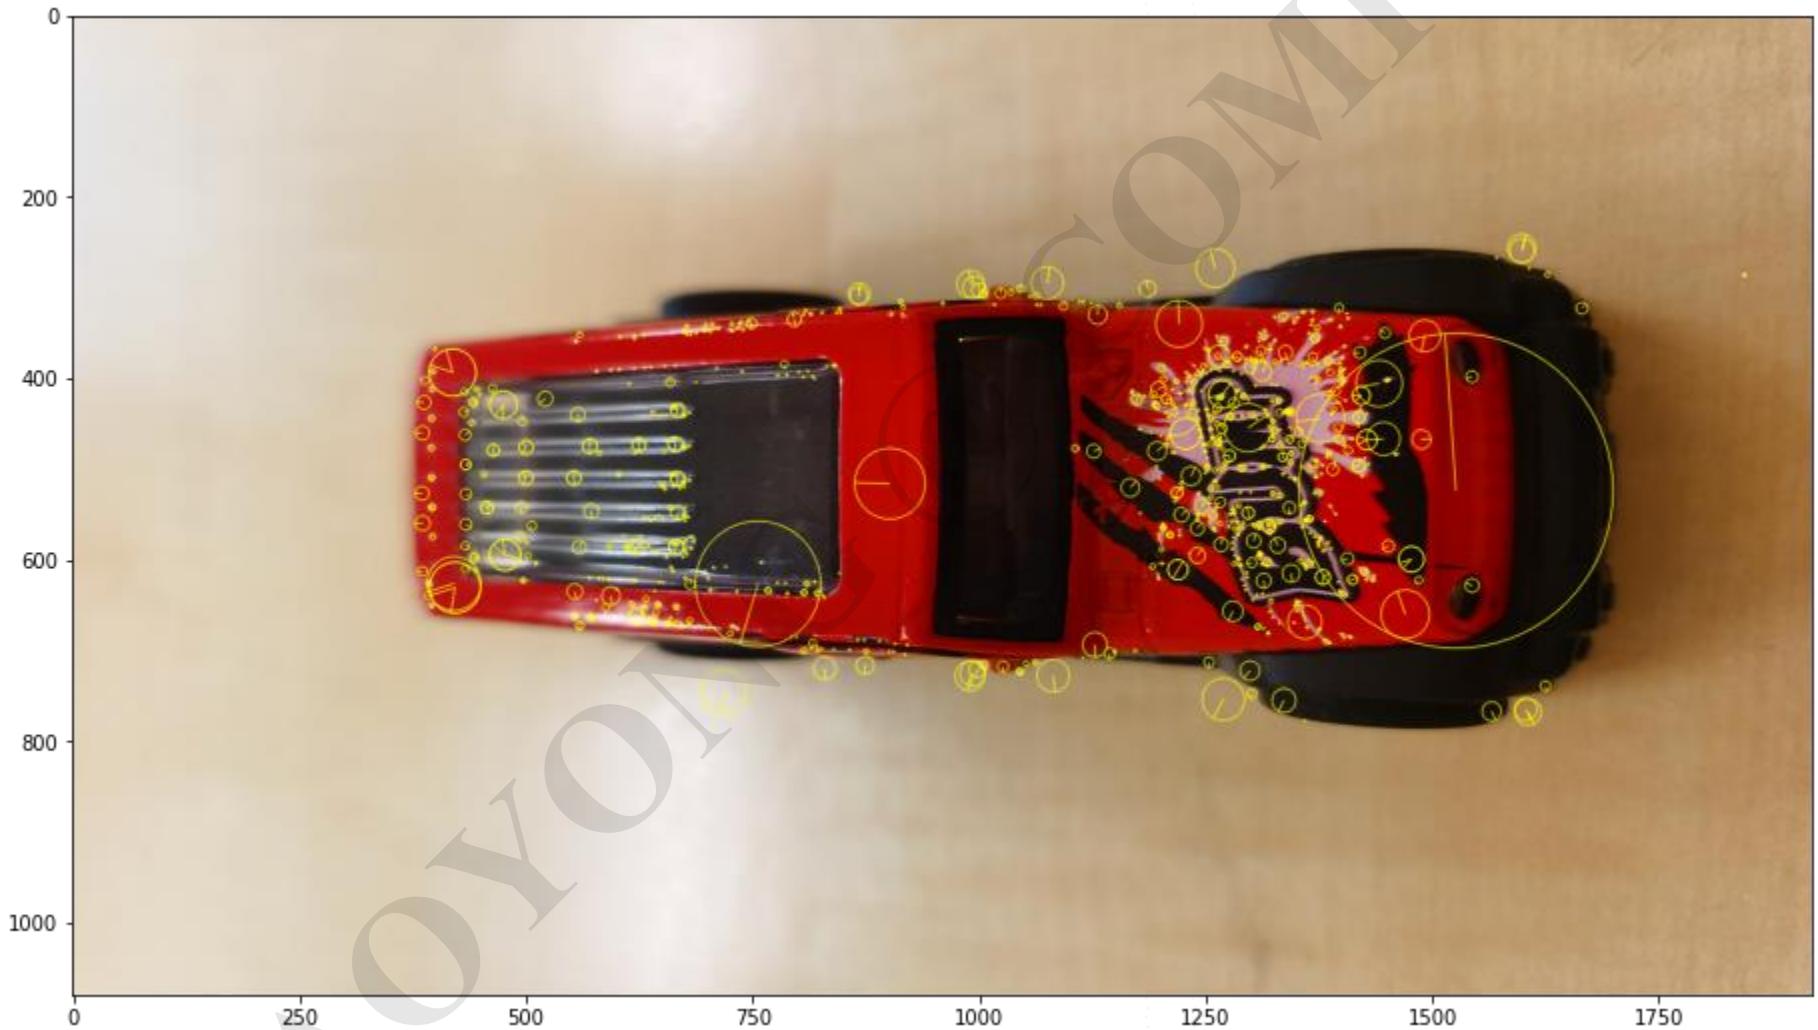


Lowe, D.G. Distinctive Image Features from Scale-Invariant Keypoints. International Journal of Computer Vision 60, 91–110 (2004).
<https://doi.org/10.1023/B:VISI.0000029664.99615.94>

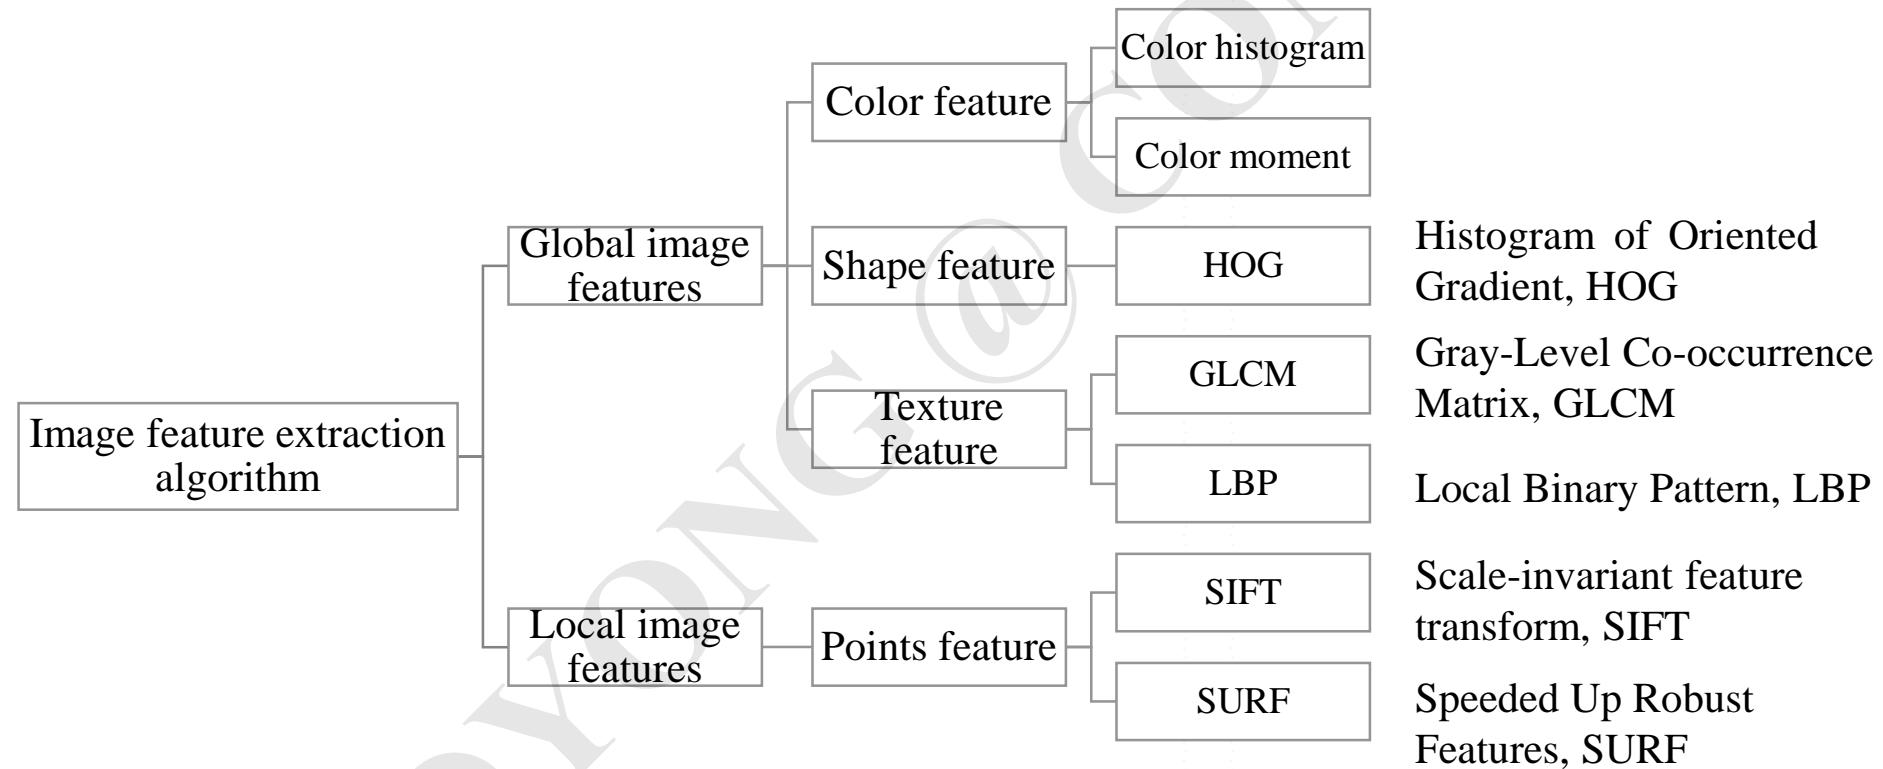
Examples from IMHere



Examples from IMHere



Features



The New Toy



The New Toy

```
1 import cv2, numpy as np, copy
2 # machter reference https://docs.opencv.org/4.x/dc/dc3/tutorial\_py\_matcher.html
3
4 # declare a video capture object
5 vid = cv2.VideoCapture(0)
6 inversion_on=False
7 keypoints_on=False
8
9 sift = cv2.SIFT_create()
10 detector=sift
11
12 bf_matcher=cv2.BFMatcher_create()
13 FLANN_INDEX_KDTREE = 1
14 index_params = dict(algorithm=FLANN_INDEX_KDTREE, trees=5)
15 search_params = dict(checks=100)
16 flann = cv2.FlannBasedMatcher(index_params, search_params)
17 matcher=flann
18 kp_old,des_old=None,None
19 traces={}
20
21 def distance(p1, p2):
22     dis = ((p2[0] - p1[0]) ** 2 + (p2[1] - p1[1]) ** 2) ** 0.5
23     return dis
24
```

The New Toy

```
25 while True:  
26     # capture the video frame by frame  
27     ret, frame_raw = vid.read()  
28     frame = cv2.flip(frame_raw, 1) # optional  
29     # get dimensions of the frame  
30     h, w, c = frame.shape  
31     print(h,w,c)  
32  
33     #  
34     gray = cv2.cvtColor(frame, cv2.COLOR_RGB2GRAY)  
35     kp, des = detector.detectAndCompute(gray, None)  
36     #print(type(kp))  
37  
38     if kp_old!=None:  
39         matches = matcher.knnMatch(des,des_old, k=2)  
40         good = []  
41         for m,n in matches:  
42             if m.distance < 0.8*n.distance:  
43                 good.append(m)  
44     #matches=bf_matcher.match(des,des_old)  
45  
46     new_traces={}  
47     for match in good:  
48         pt1,pt2=kp[match.queryIdx].pt,kp_old[match.trainIdx].pt  
49         pt1,pt2=[int(pt1[0]),int(pt1[1])],[int(pt2[0]),int(pt2[1])]  
50         if distance(pt1,pt2)>200: continue
```

The New Toy

```
52         if (match.trainIdx in traces):
53             len_pts=len(traces[match.trainIdx])
54             new_traces[match.queryIdx]=[pt1]+copy.deepcopy(traces[match.trainIdx])[:20 if len_
55             pts=np.array(new_traces[match.queryIdx],dtype='int32')
56             #print(pts)
57             cv2.polyline(frame,[pts],False,(255, 255, 0),1)
58         else:
59             cv2.line(frame,pt1,pt2,(255, 255, 0),1)
60             new_traces[match.queryIdx]=[pt1,pt2]
61             traces=new_traces
62
63 kp_old,des_old=kp,des
64 if keypoints_on: cv2.drawKeypoints(frame, kp, frame, (255, 255, 0), cv2.DRAW_MATCHES_FLAGS_DI
65
66 # display the resulting frame
67 cv2.imshow('COMP 4423', 255-frame if inversion_on else frame)
68
69 key=cv2.waitKey(1) & 0xFF
70 # quit when 'q' is pressed
71 if key== ord('q'):
72     break
73 if key==ord('i'):
74     inversion_on=not inversion_on
75 if key==ord('k'):
76     keypoints_on=not keypoints_on
77
78 # releases the cap object
```

Thank you!

