Image Retrieval – COMP4423 Computer Vision

Xiaoyong Wei (魏驍勇)

x1wei@polyu.edu.hk

Department of Computing
電子計算學系

The Hong Kong Polytechnic University
香港理工大學

Opening Minds • Shaping the Future
啟迪思維 • 成就未來

New Toy

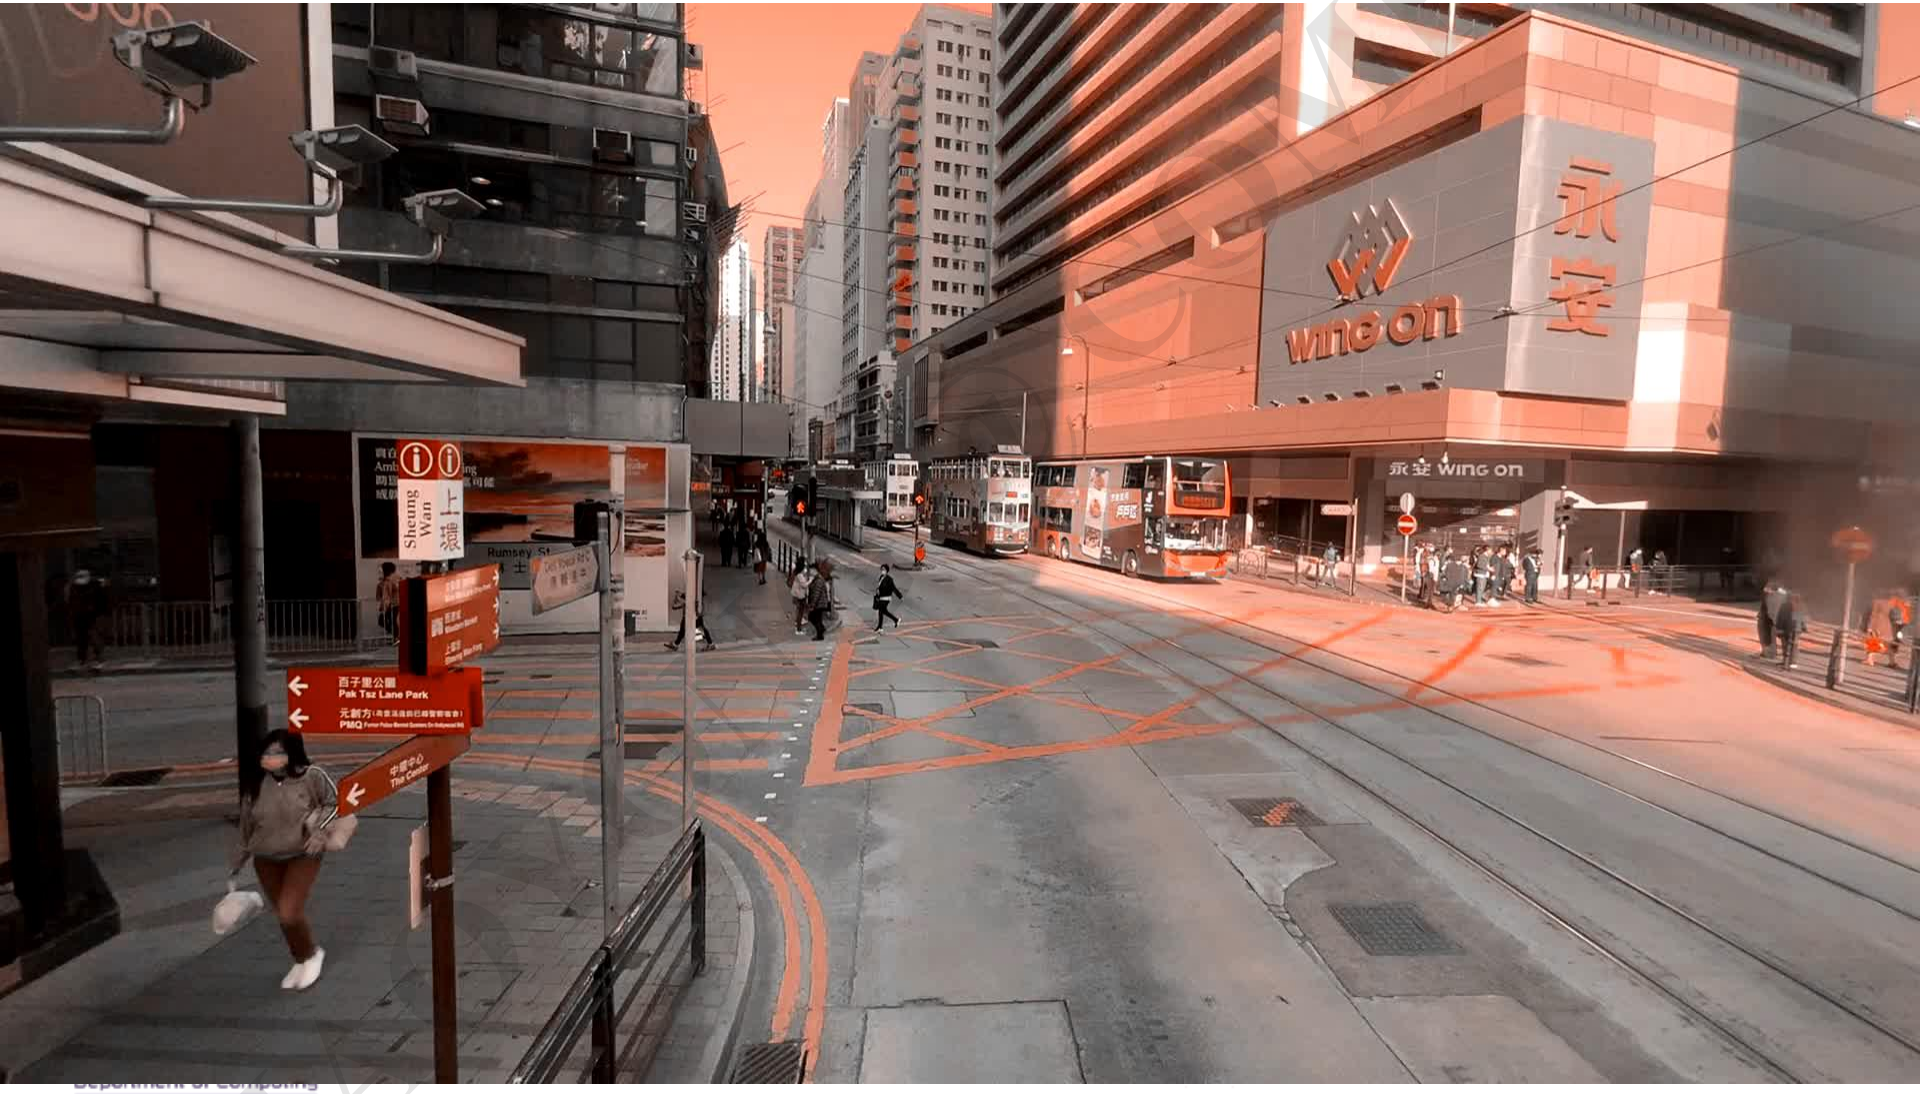
wing on
永安
永安 wing on
Sheung Wan
上環
Rumsey St
百子里公園
Pak Tsz Lane Park
中環中心
The Center

New Toy

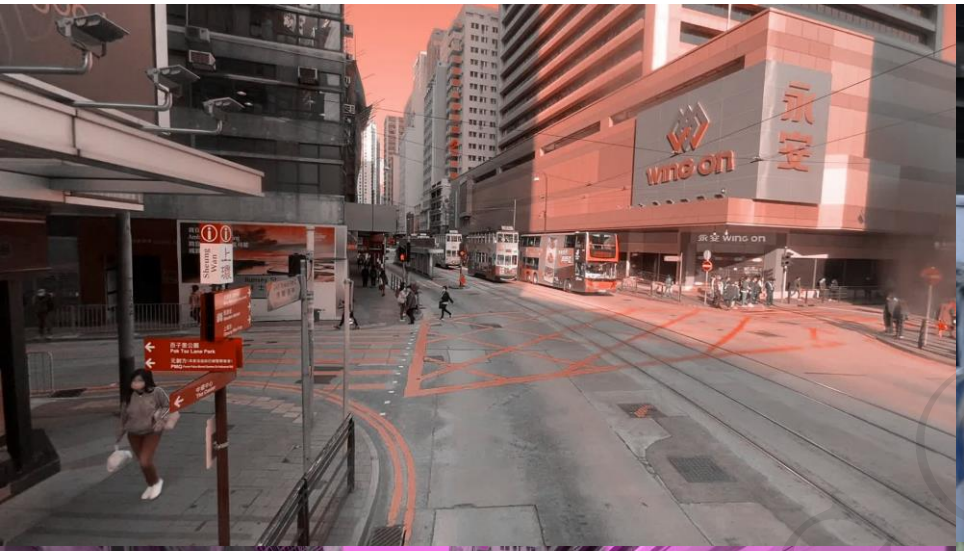

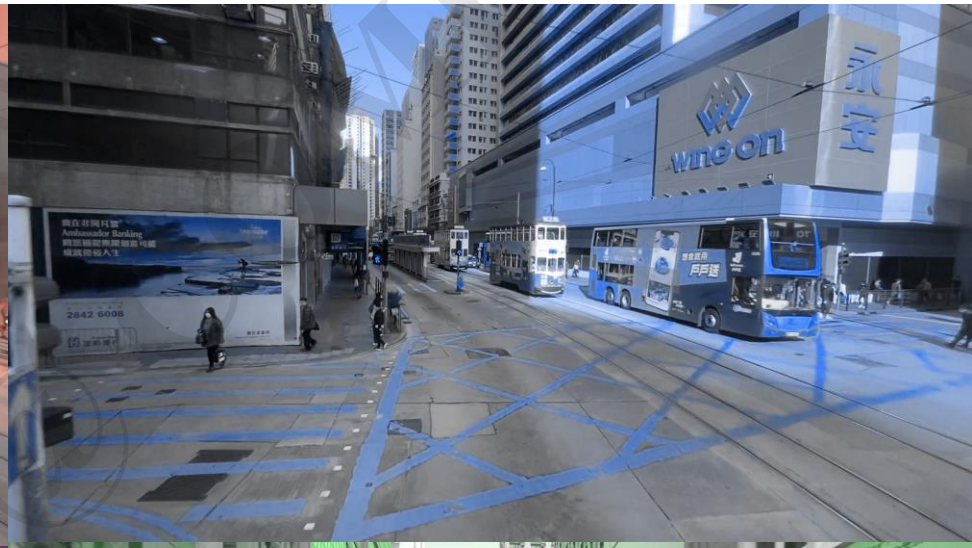

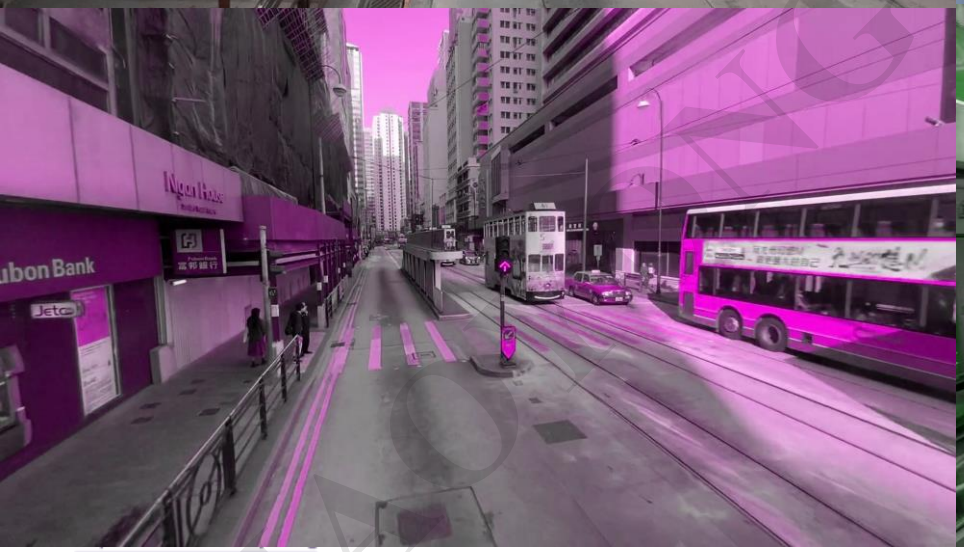

Outline

>Clustering

>K-Means

>Content-based image retrieval (CBIR)

>Bag of Visual Words (BoVW)

In Feature Extraction, we teach the computers to represent the "content" of the images.

How can we use these "content"?

Recall where we started

How do you group them?

Feature Space

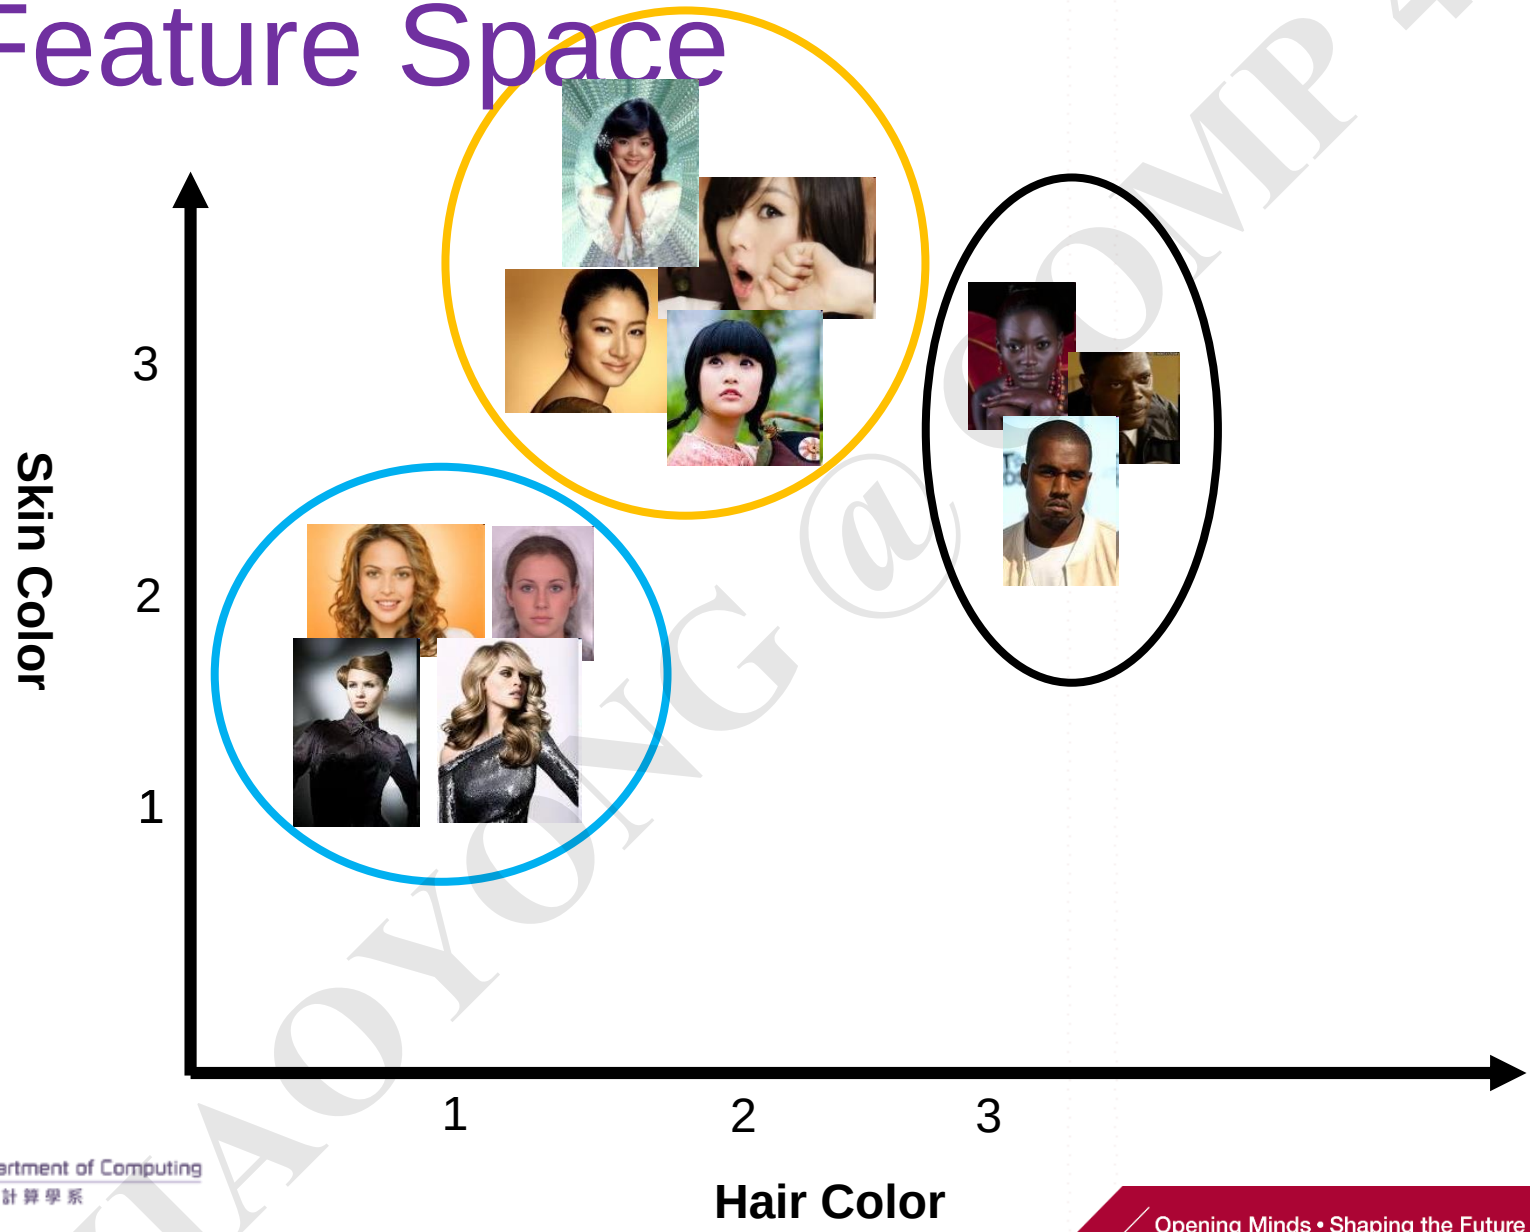

Images become “numbers” (feature vectors) in the **feature space** after the feature extraction.

To use the “content” is to play those “numbers”. This holds to nearly everything we’re going to learn in the rest of this course.

In a more general sense, this applies to texts, audios, videos, and a wide range of other media/information.

We're all trying to represent things as "numbers" in the feature space, making them "readable" for computers.

Feature Space

Can we teach computers to draw circles for grouping?

We can answer the question now.

To group the images is to make the circles as far as possible from each other, while the images inside the same circle as close as possible.

This is the idea of **clustering**: to maximize the **inter-cluster** distance while to minimize the **intra-cluster** distance.

Think in the feature space

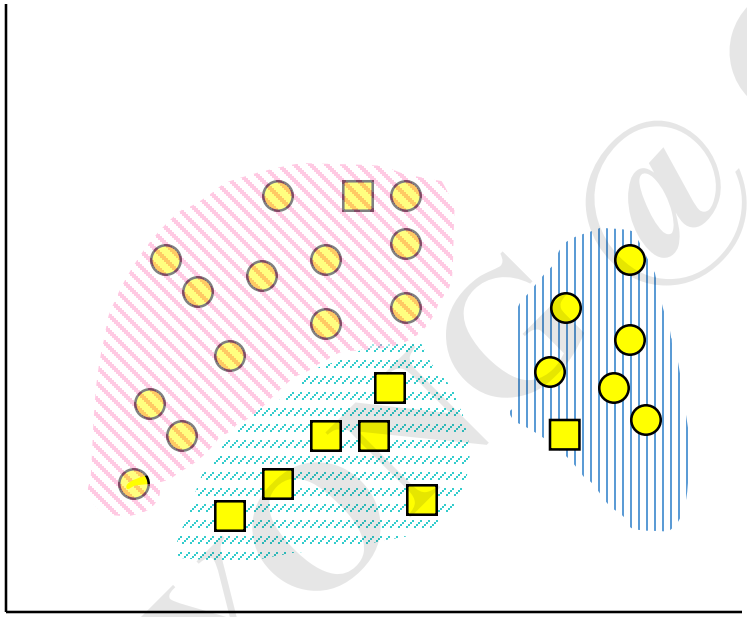

Think in the feature space

- K-Means
  - 1: Pick a number (K) of cluster centers (at random)
  - 2: Assign every item to its nearest cluster center (e.g. using Euclidean distance)
  - 3: Move each cluster center to the **mean** of its assigned items
  - 4: Repeat steps 2,3 until convergence (change in cluster assignments less than a threshold)

1. Captains assigned randomly

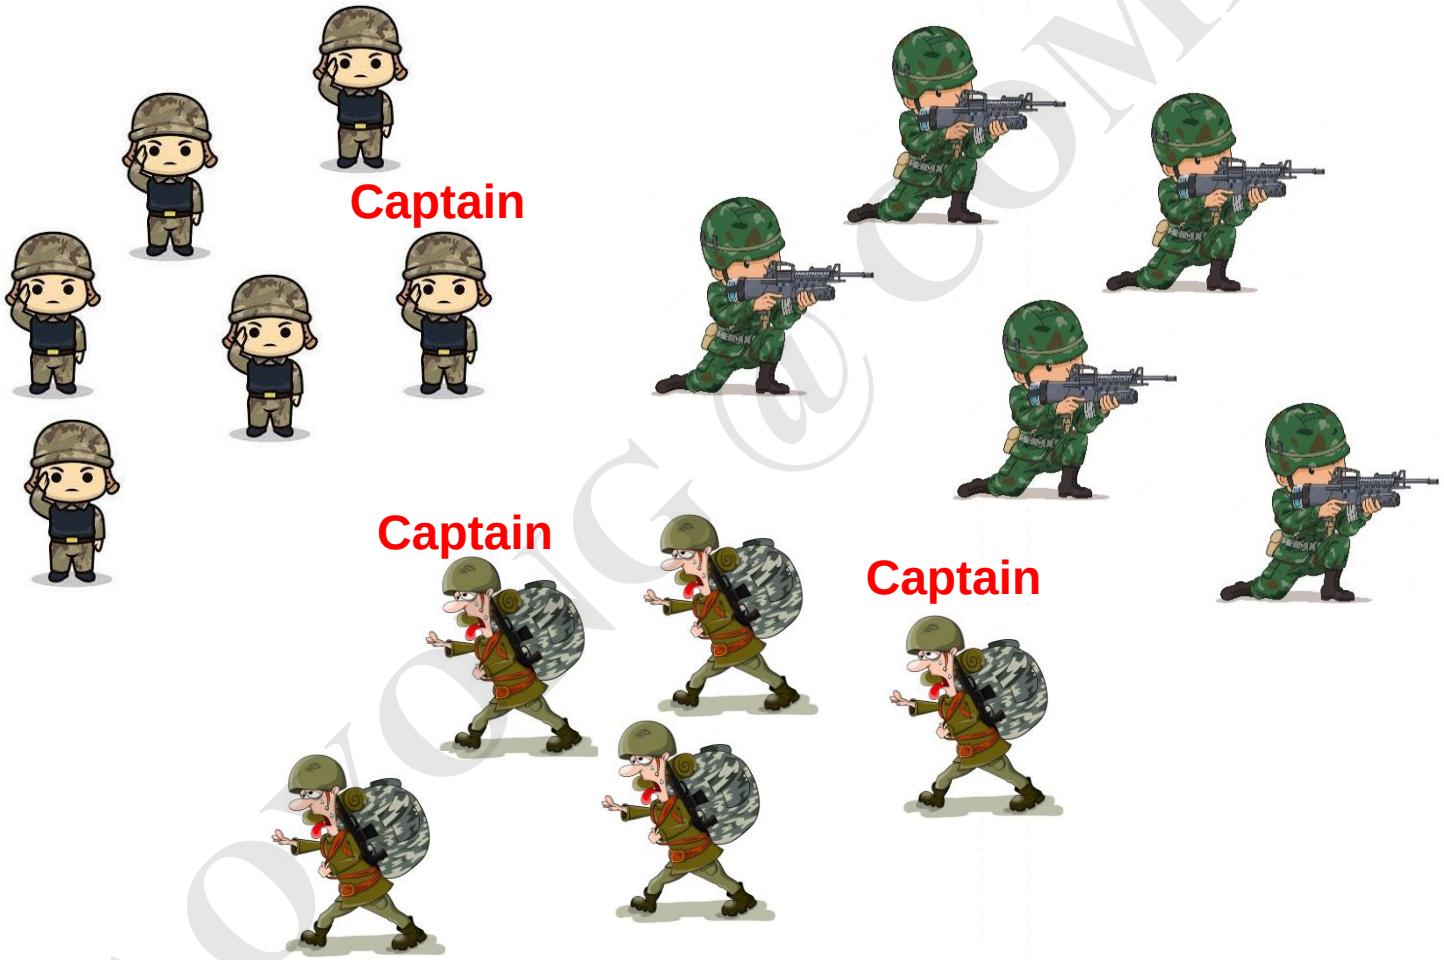

2. Soldiers report to the nearest captains

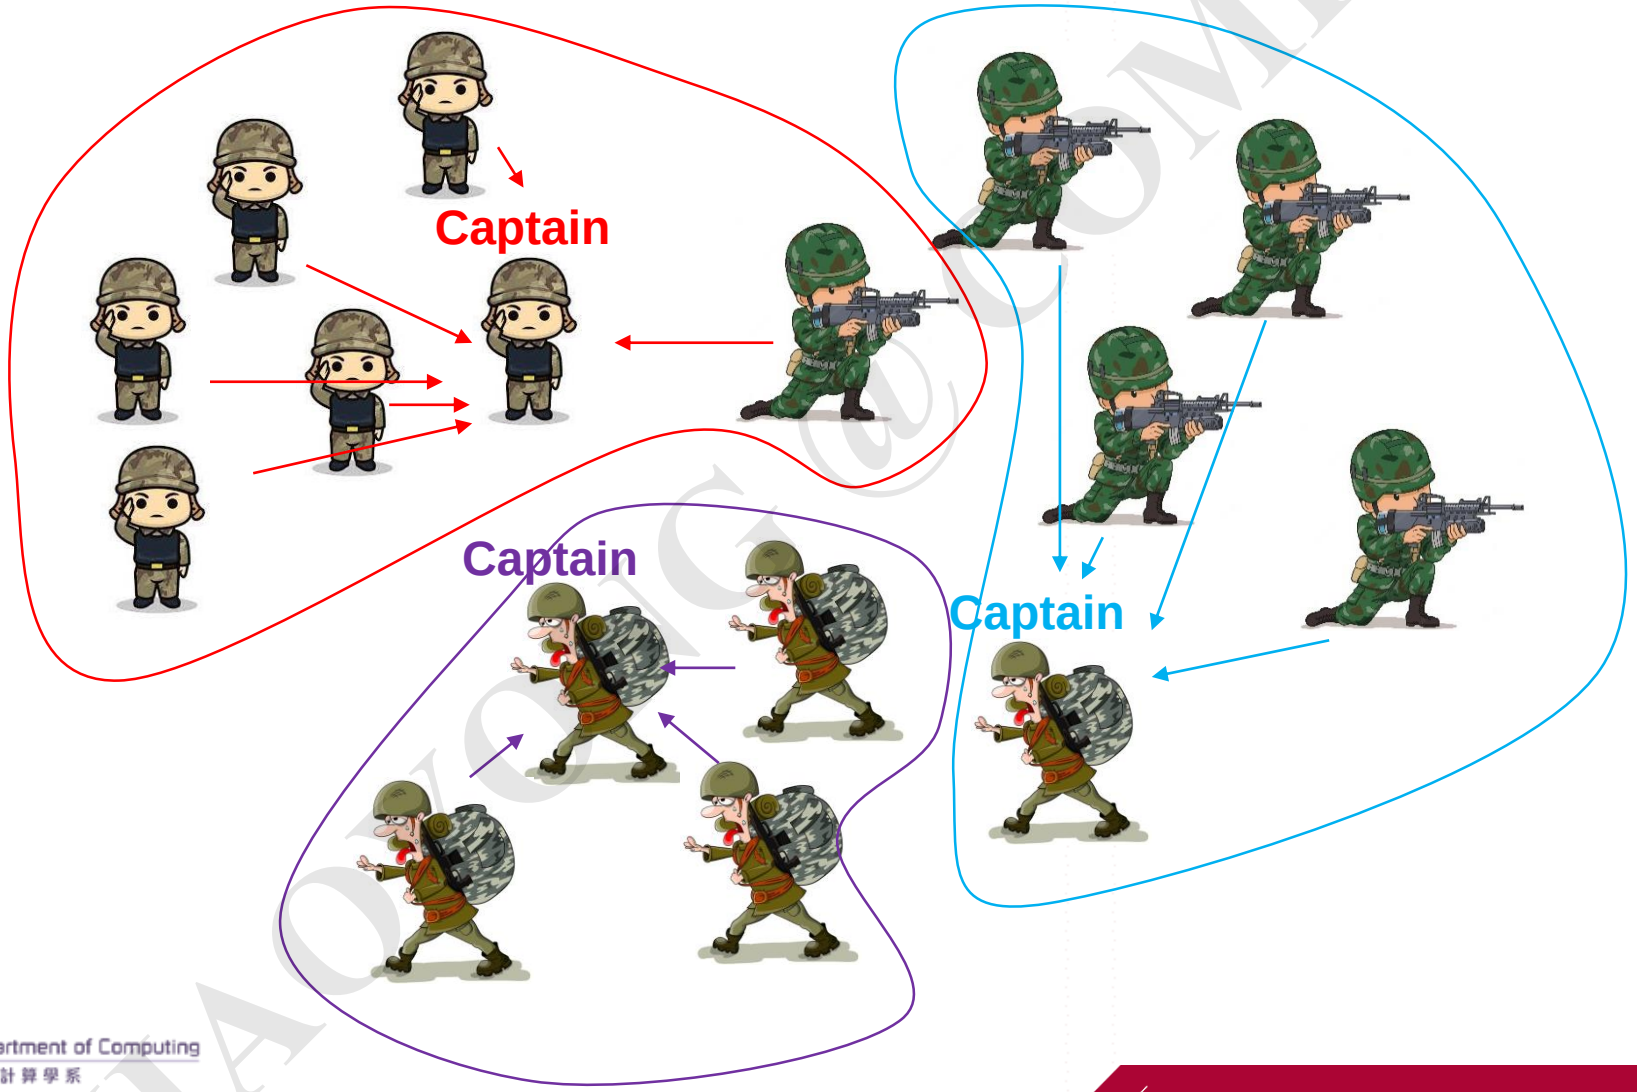

3. Re-election of captains

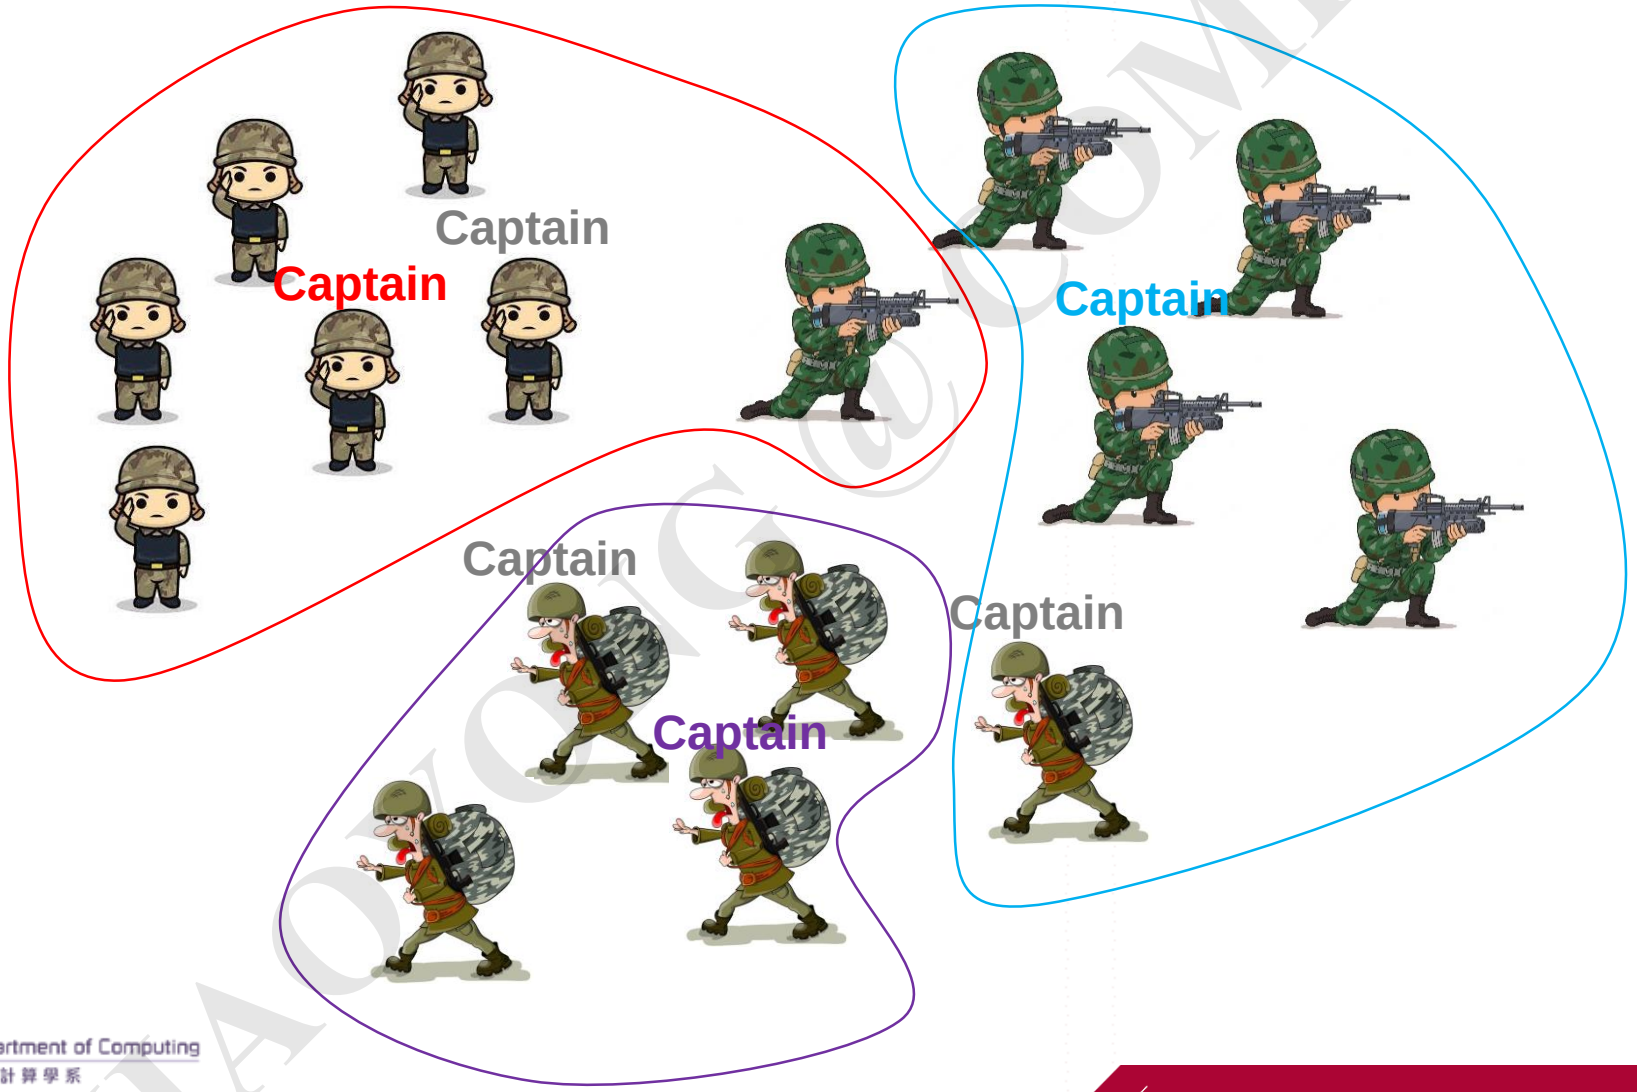

4. Report again

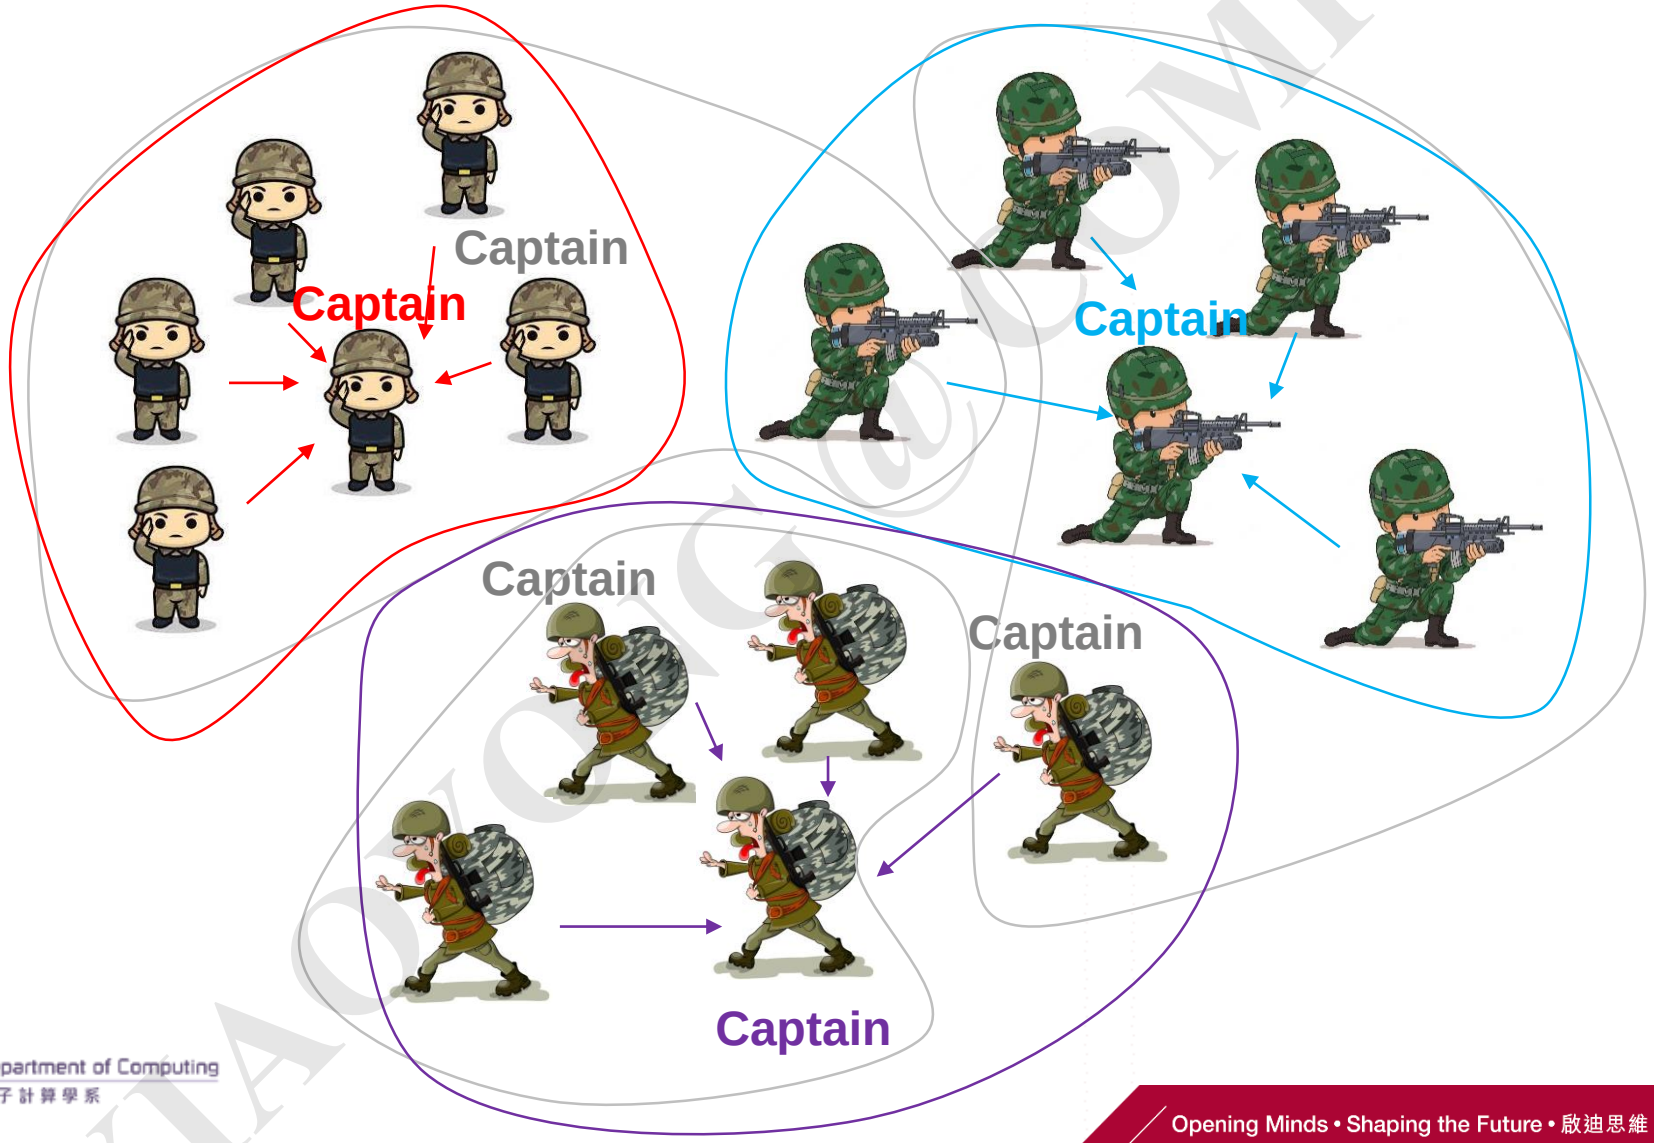

K-Means Step 1

Pick 3 initial cluster centers (randomly)

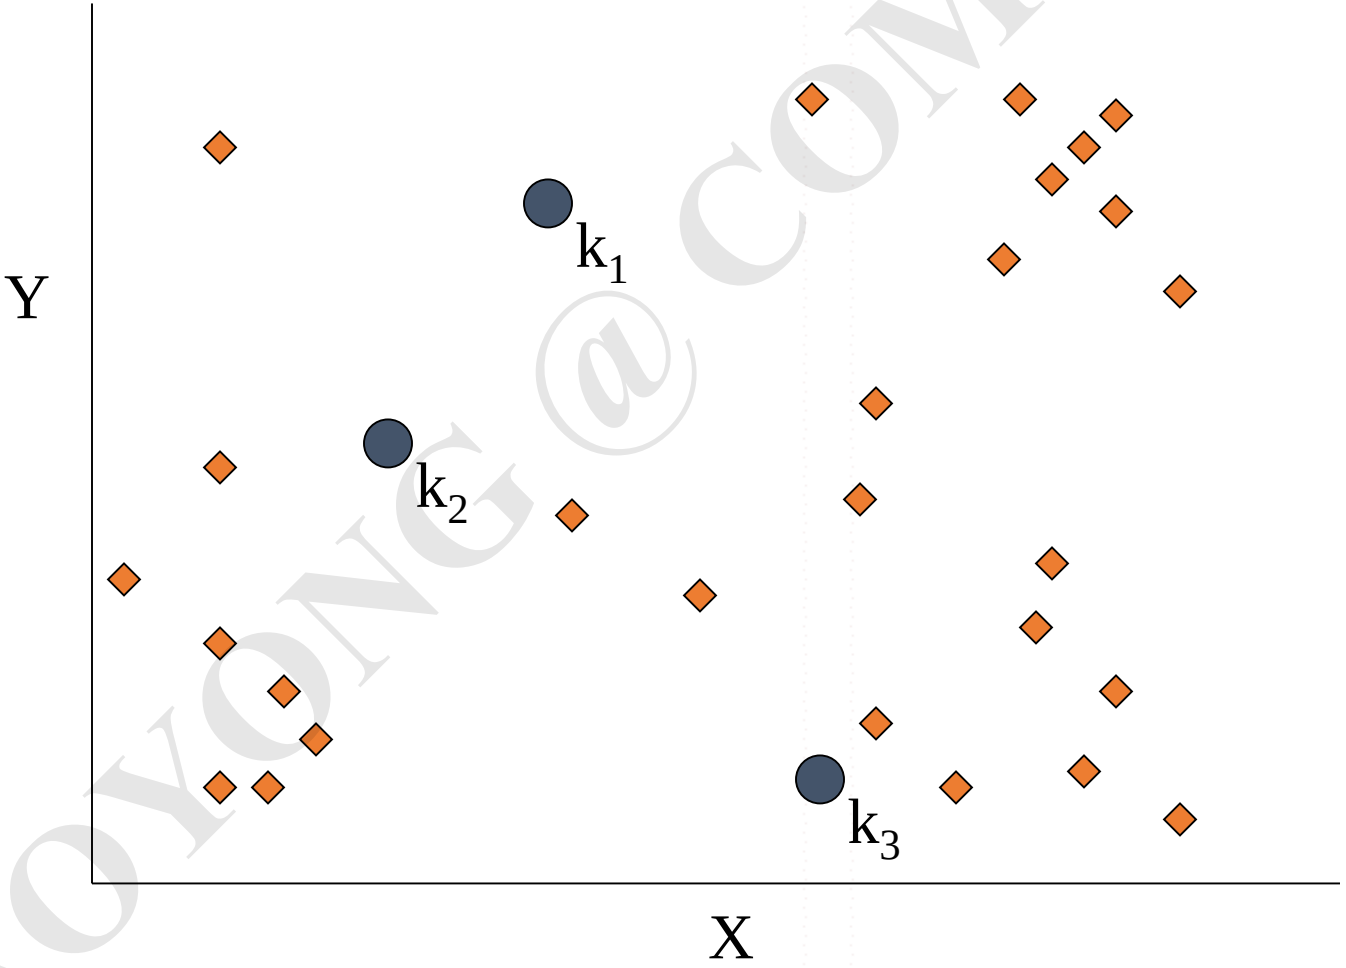

K-Means Step 2

Assign each point to the nearest cluster center

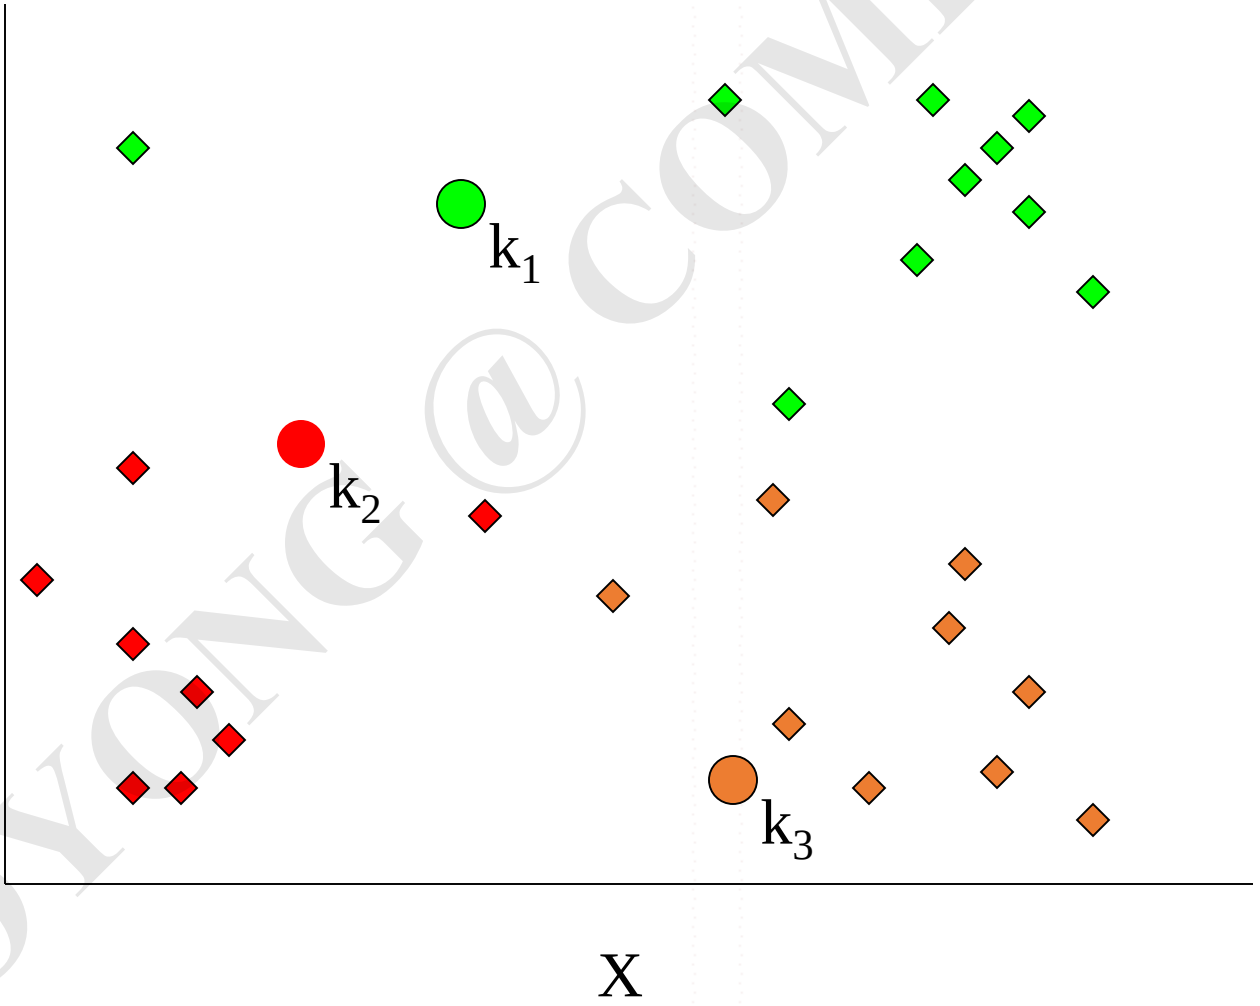

K-Means Step 3

Move
each cluster
center
to the **mean**
of each cluster

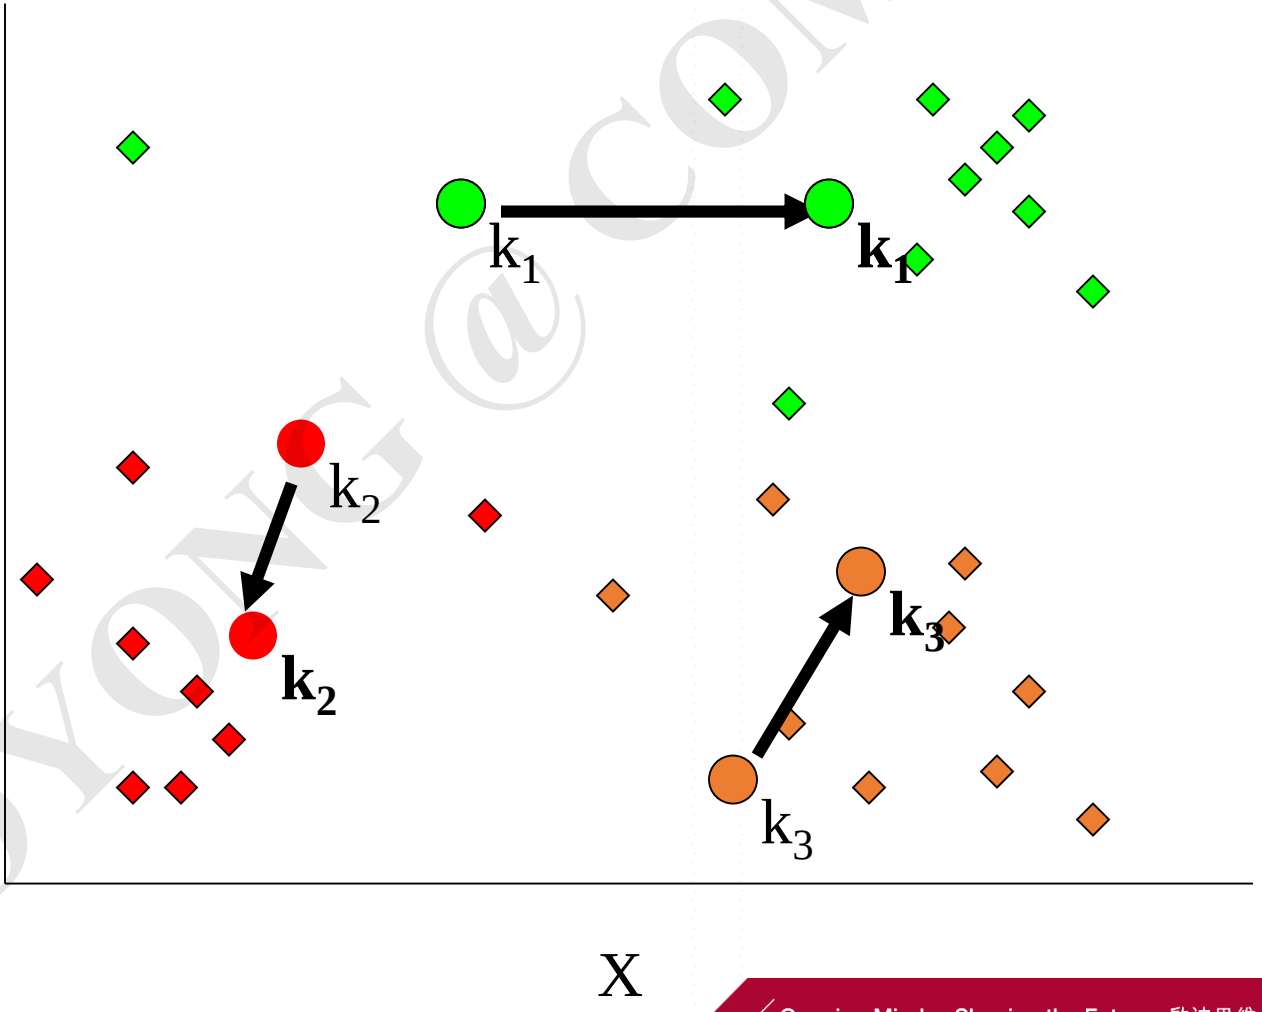

K-Means Step 4

Reassign points to the new (nearest) cluster center

Q: Which points are reassigned?

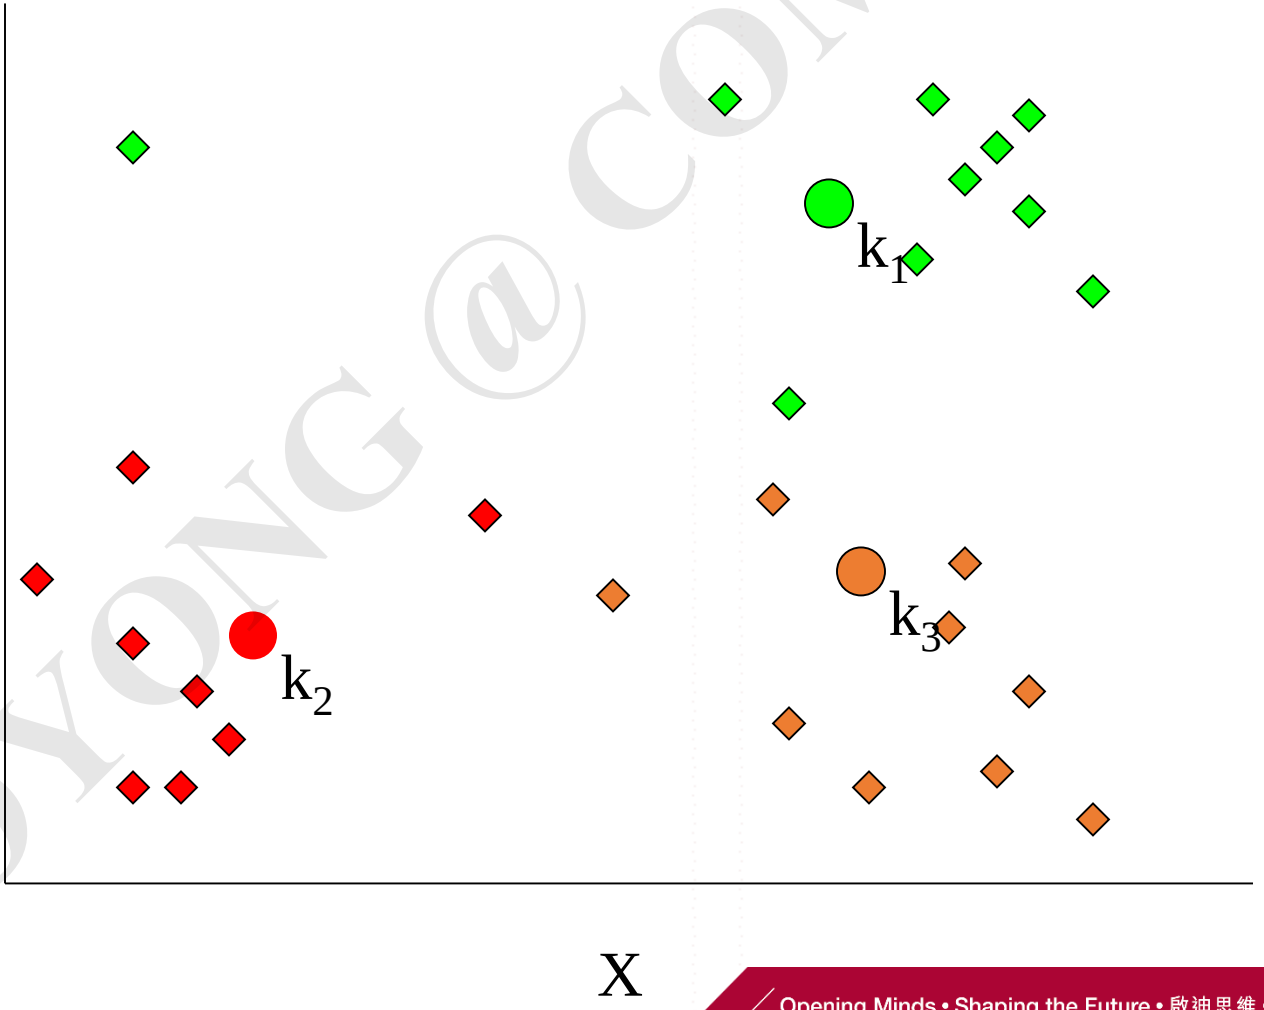

K-Means Step 4 …

A: three points with animation

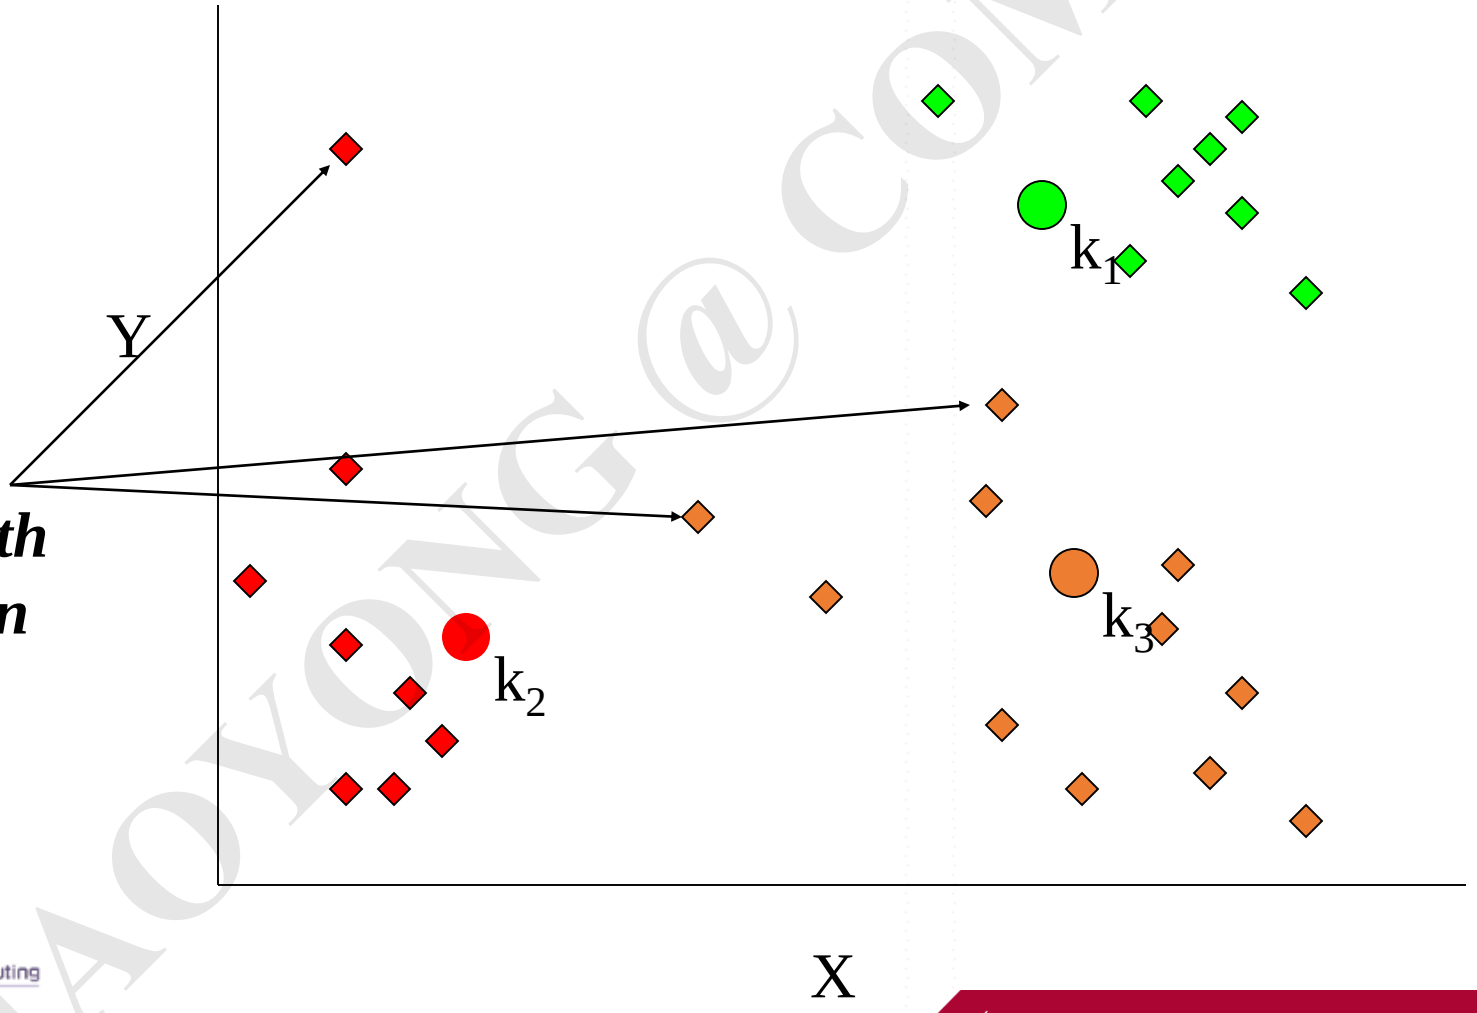

K-Means Step 4 …

Re-compute cluster means

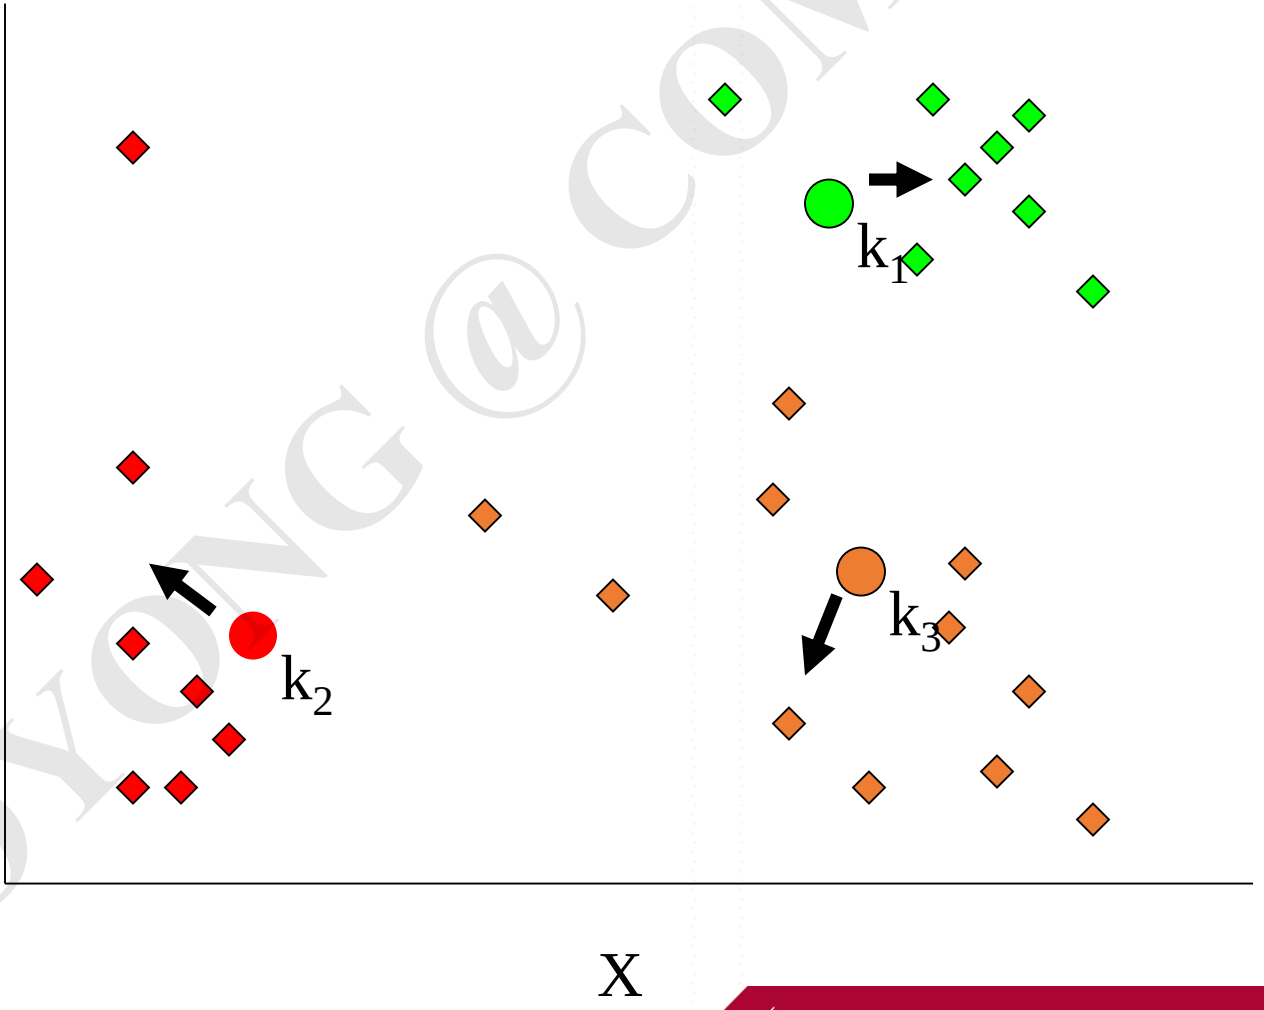

K-Means Step 5

Move cluster centers to cluster means

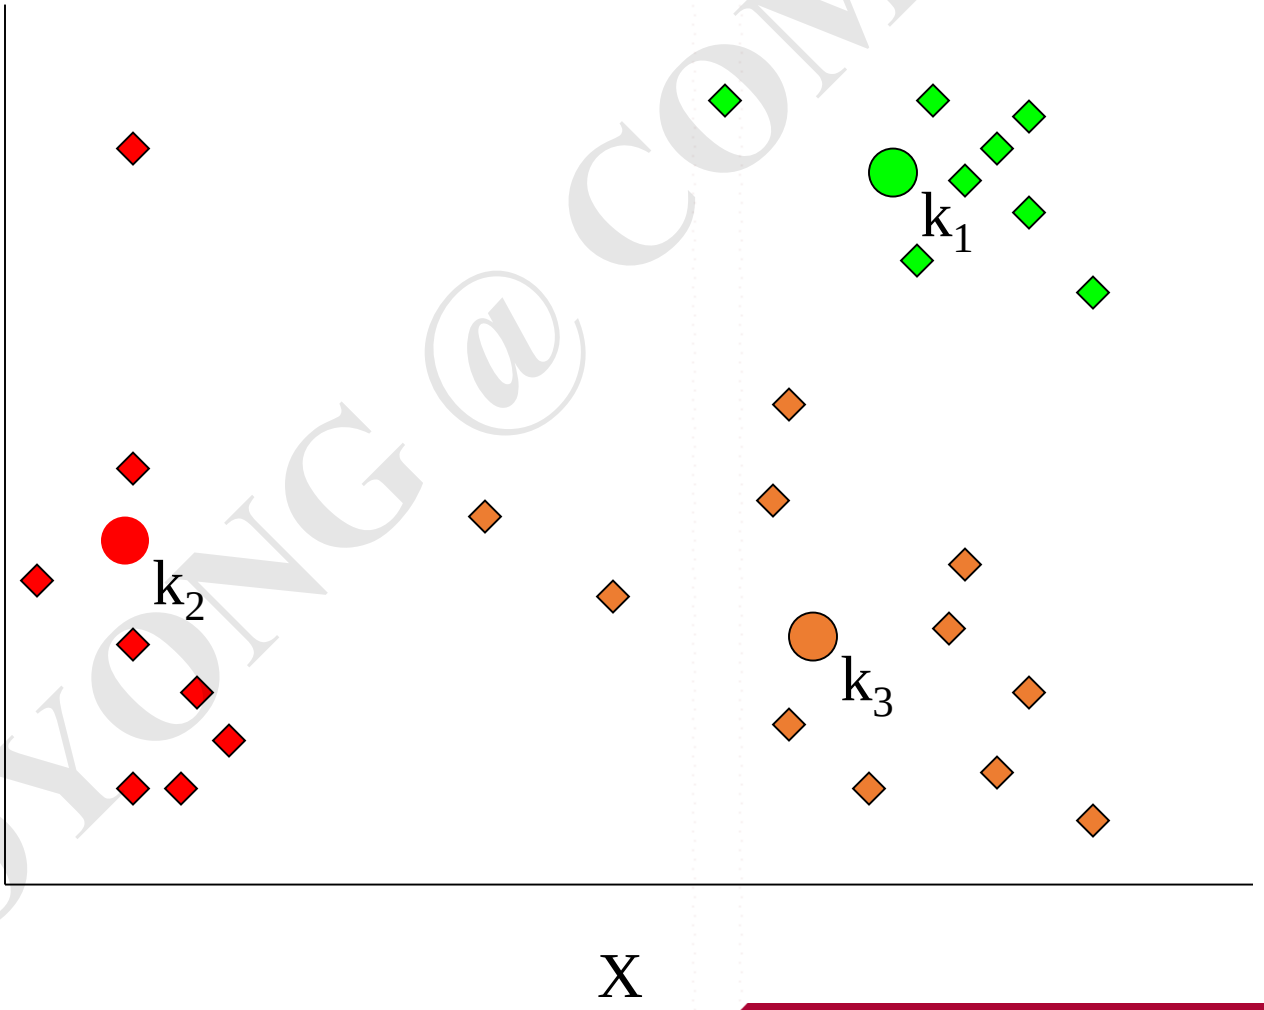

Let's take images from IMHere as examples

Let's take images from IMHere as examples

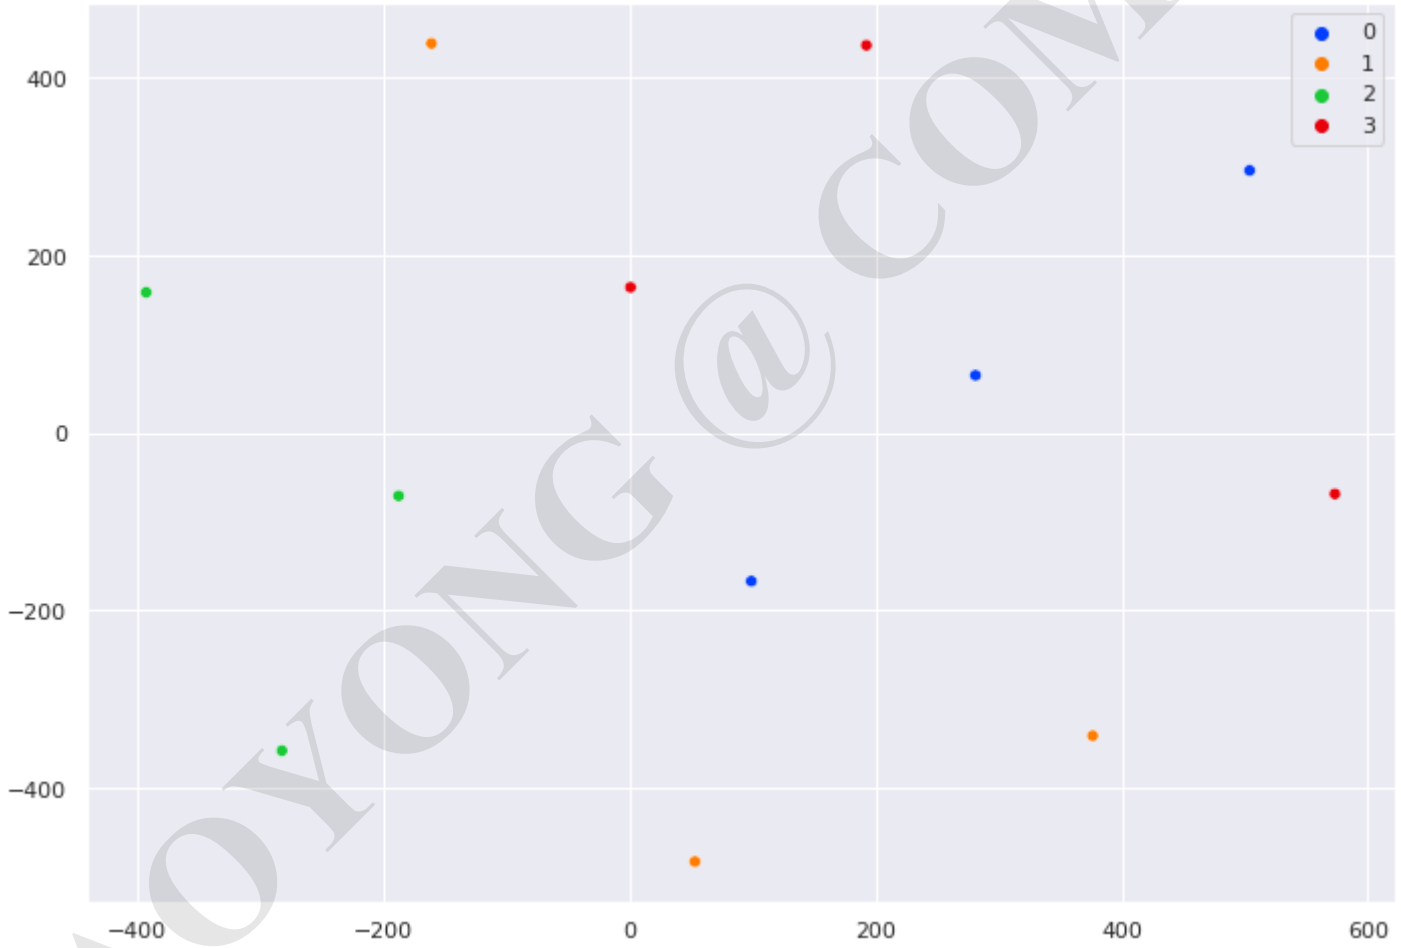

Metrics – Euclidian Distance

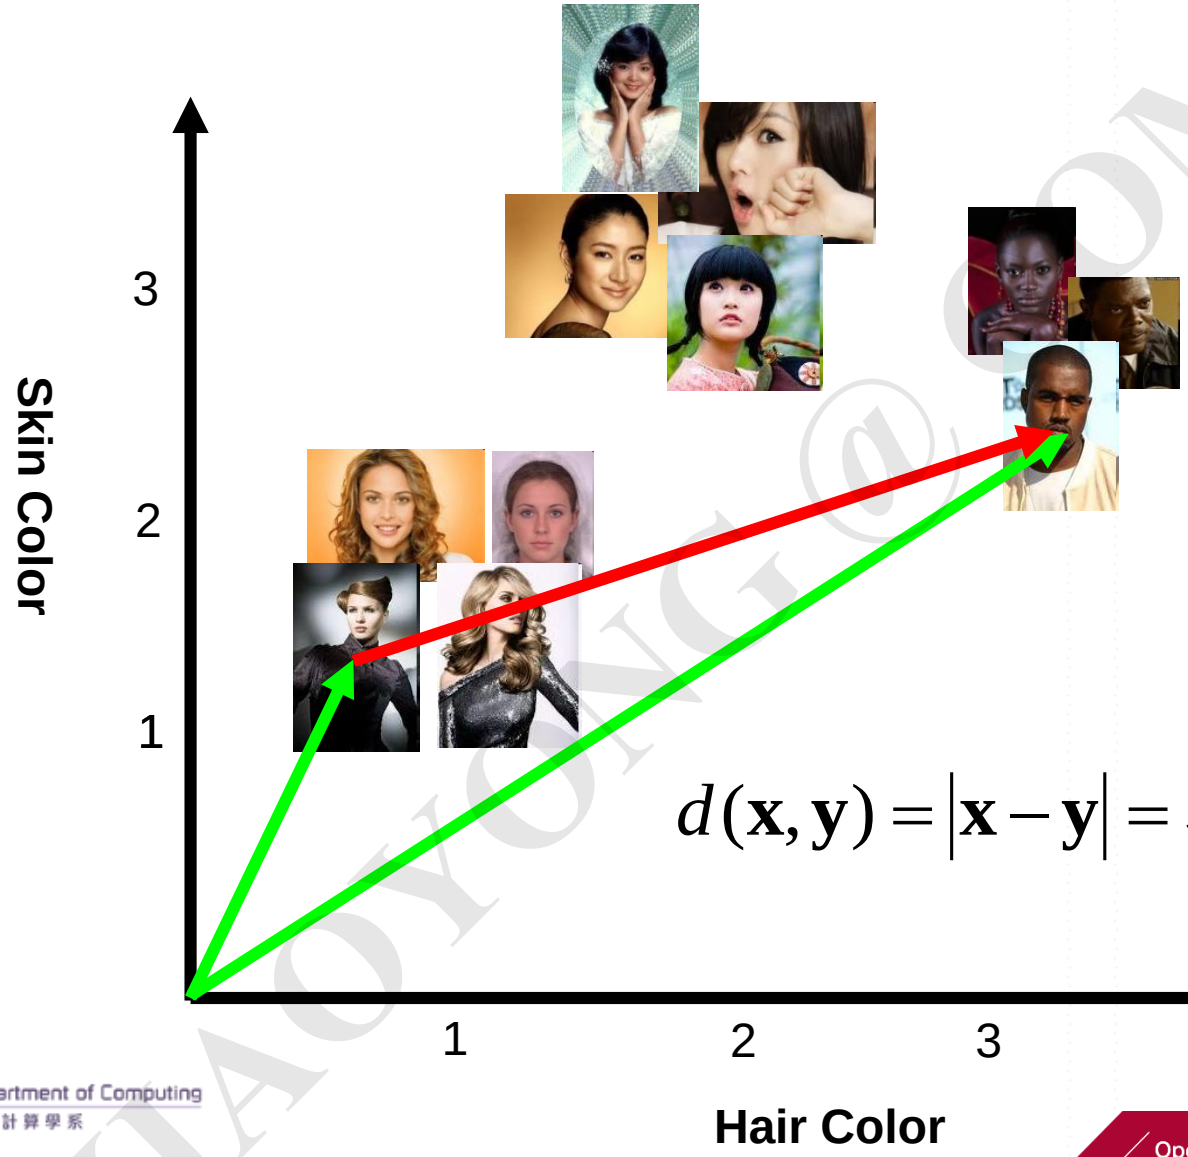

$$d(\mathbf{x},\mathbf{y}) = |\mathbf{x}-\mathbf{y}| = \sqrt{<\mathbf{x}-\mathbf{y},\mathbf{x}-\mathbf{y}>}$$

Clustering is one of the most representative examples of **Unsupervised Learning**.

We will get back later.

Grouping images is fun.

However, that's not the way we employed in IMHere.

IMHere – Token-based attendance checking

In IMHere, we are **looking for** images that are with **similar content** with the one you uploaded. You are checked-in if those are what you used for registrations.

This is called **Content-Based Image Retrieval (CBIR)**.

General Steps of CBIR

- 0. Extract features vectors of all images on file
- 1. Extract the feature vector for the **query** image
- 2. Compare it to all the (**target**) images on file by calculating the query-target similarities
- 3. Sort the similarities in a descending order with which we generate a **ranked list** of the targets
- 4. Present the ranked list to the searcher

Metrics – Cosine Similarity

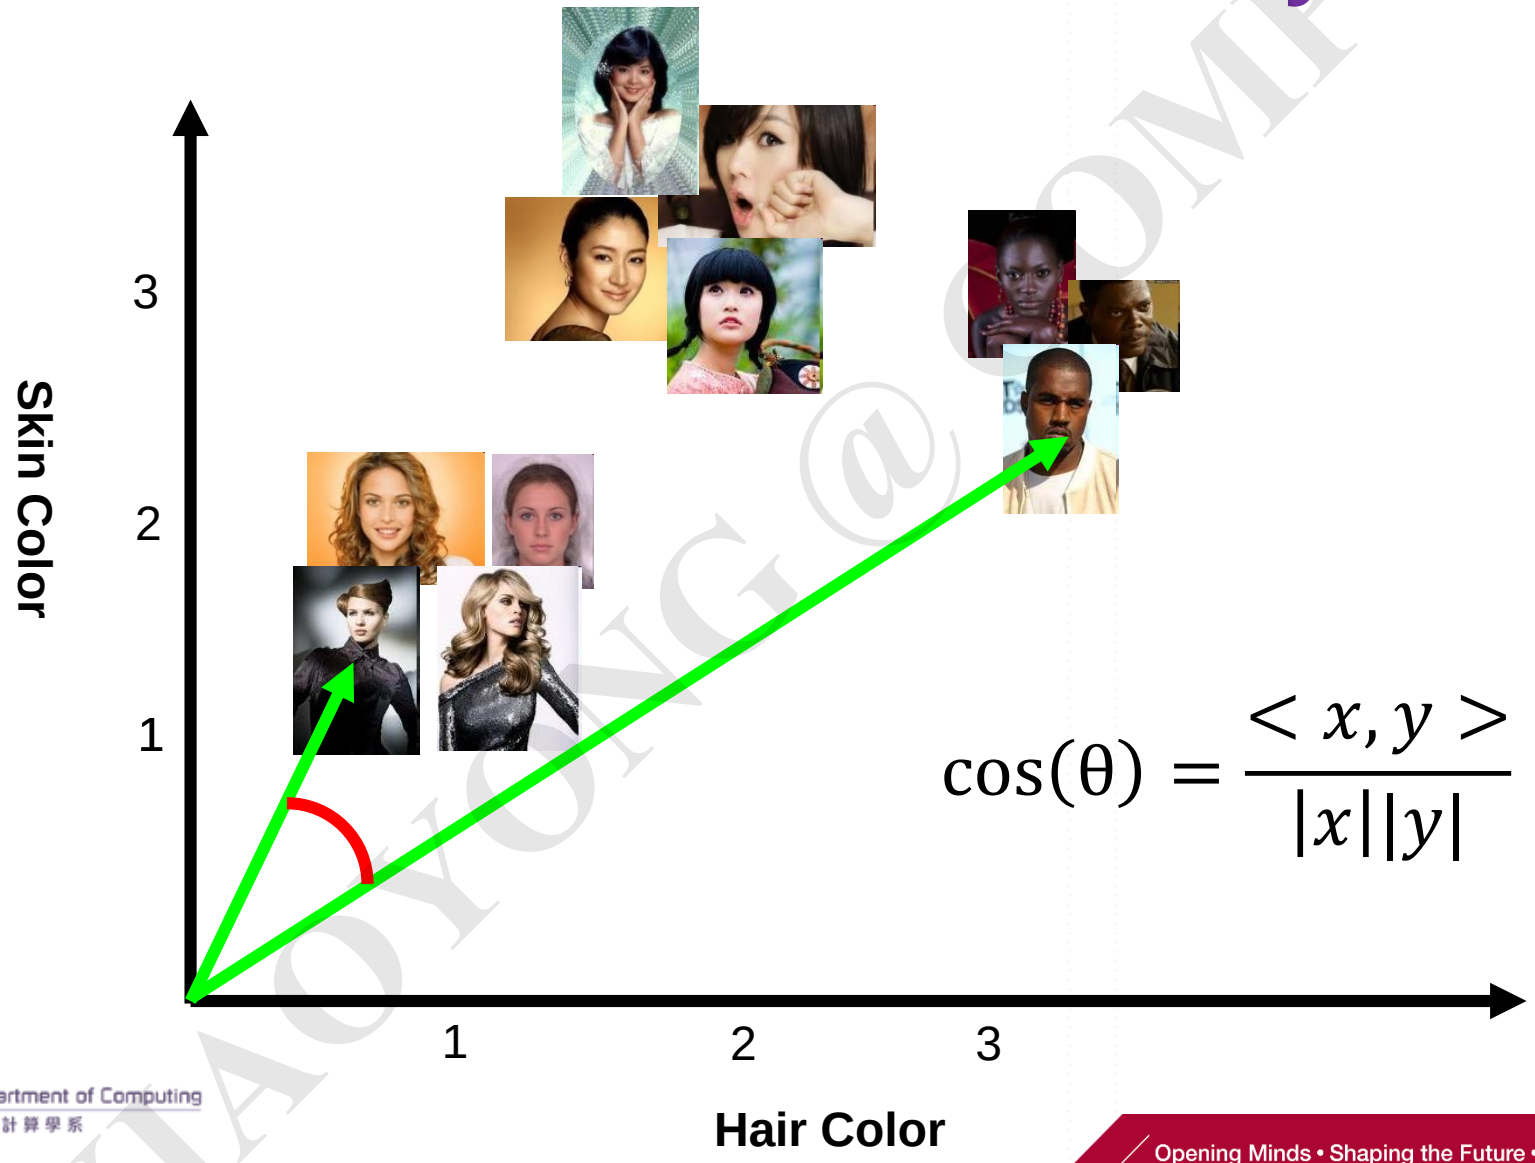

$$\cos(\theta) = \frac{< x, y >}{|x||y|}$$

Examples from IMHere

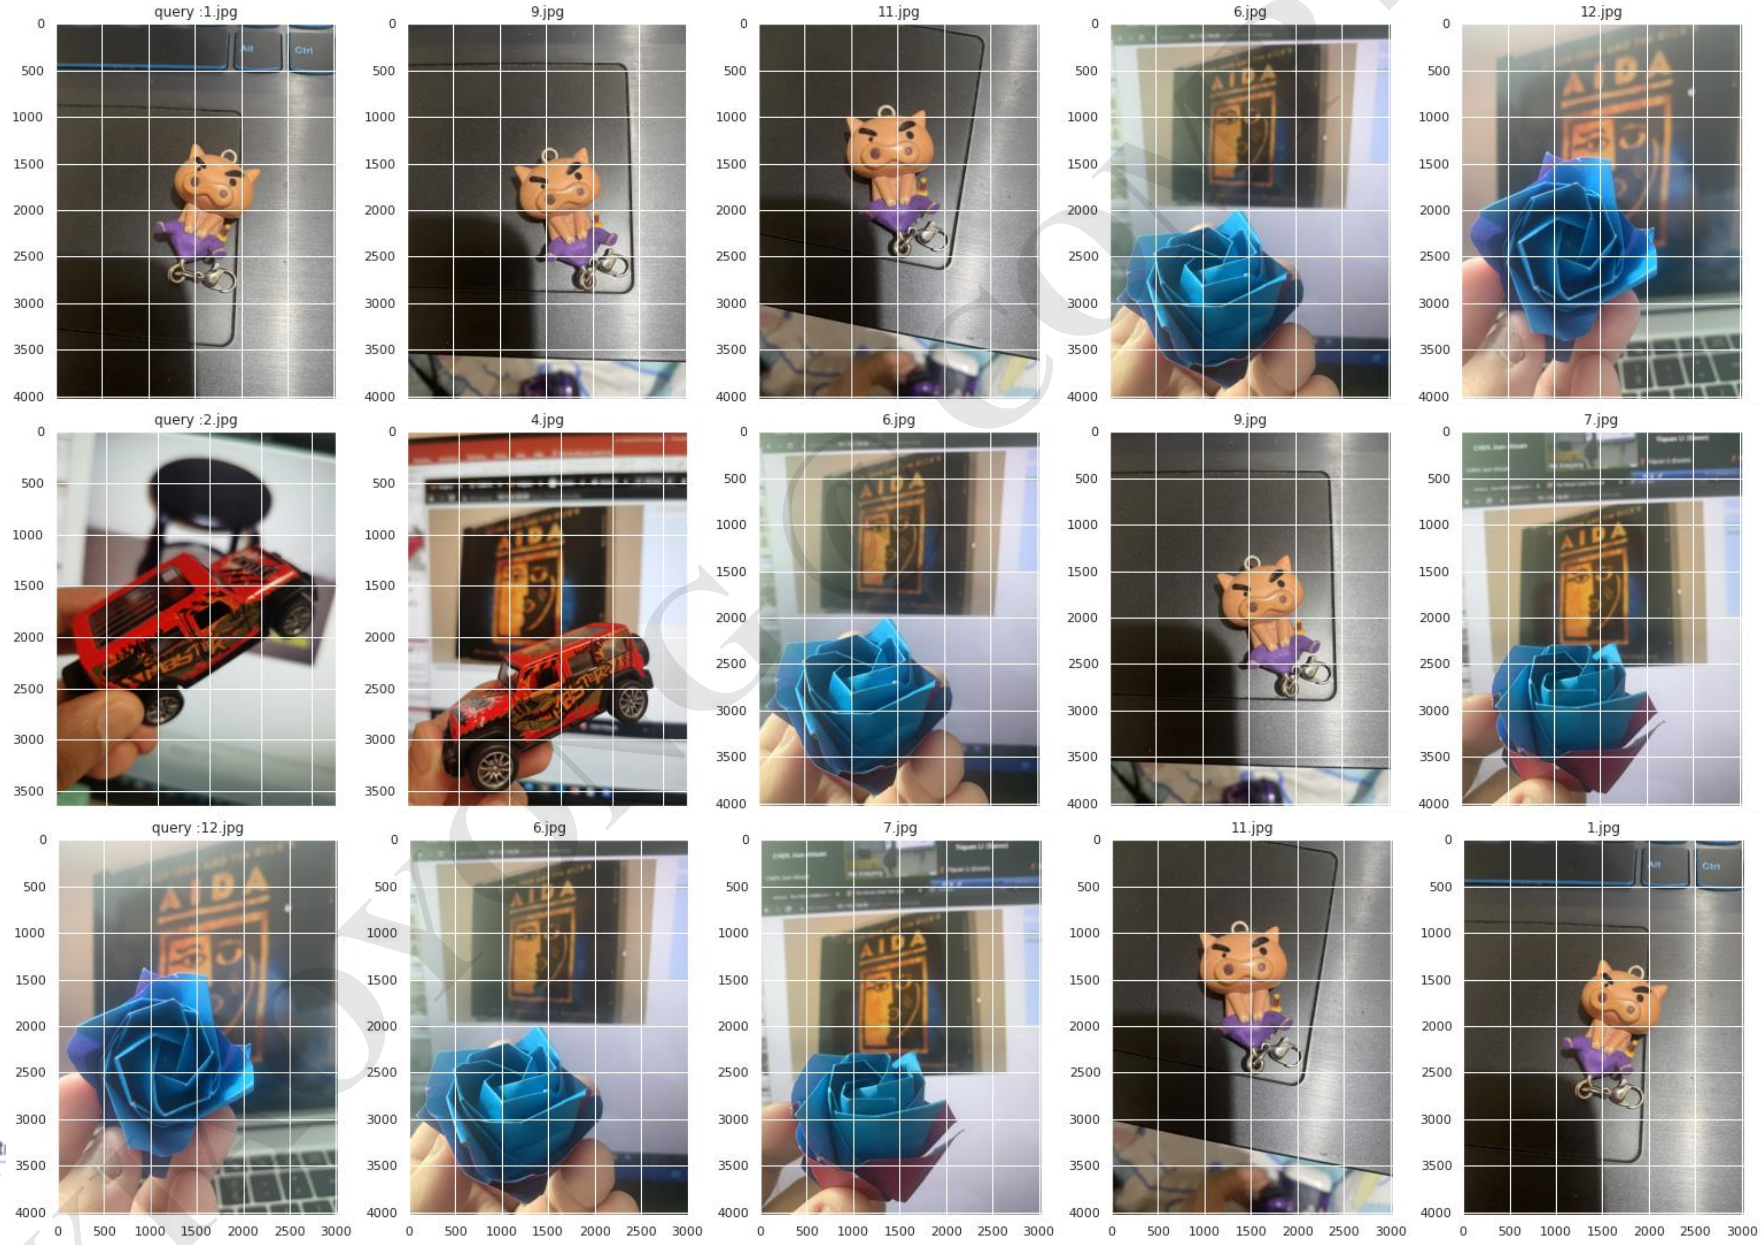

It's straightforward to compare the query to the targets in an sequential manner (one by one), but it's not efficient.

Let's assume we have 1 million targets in IMHere. It may take 27 hours per query (100 milliseconds for each target).

Is there a better way?

Clustering comes to play!

Indexing: group the images as clusters and pick one from each cluster as its representative

Coarse: compare to the representatives only and find top-k

Fine: compare to the member images of the top-k clusters

Index the images as a tree

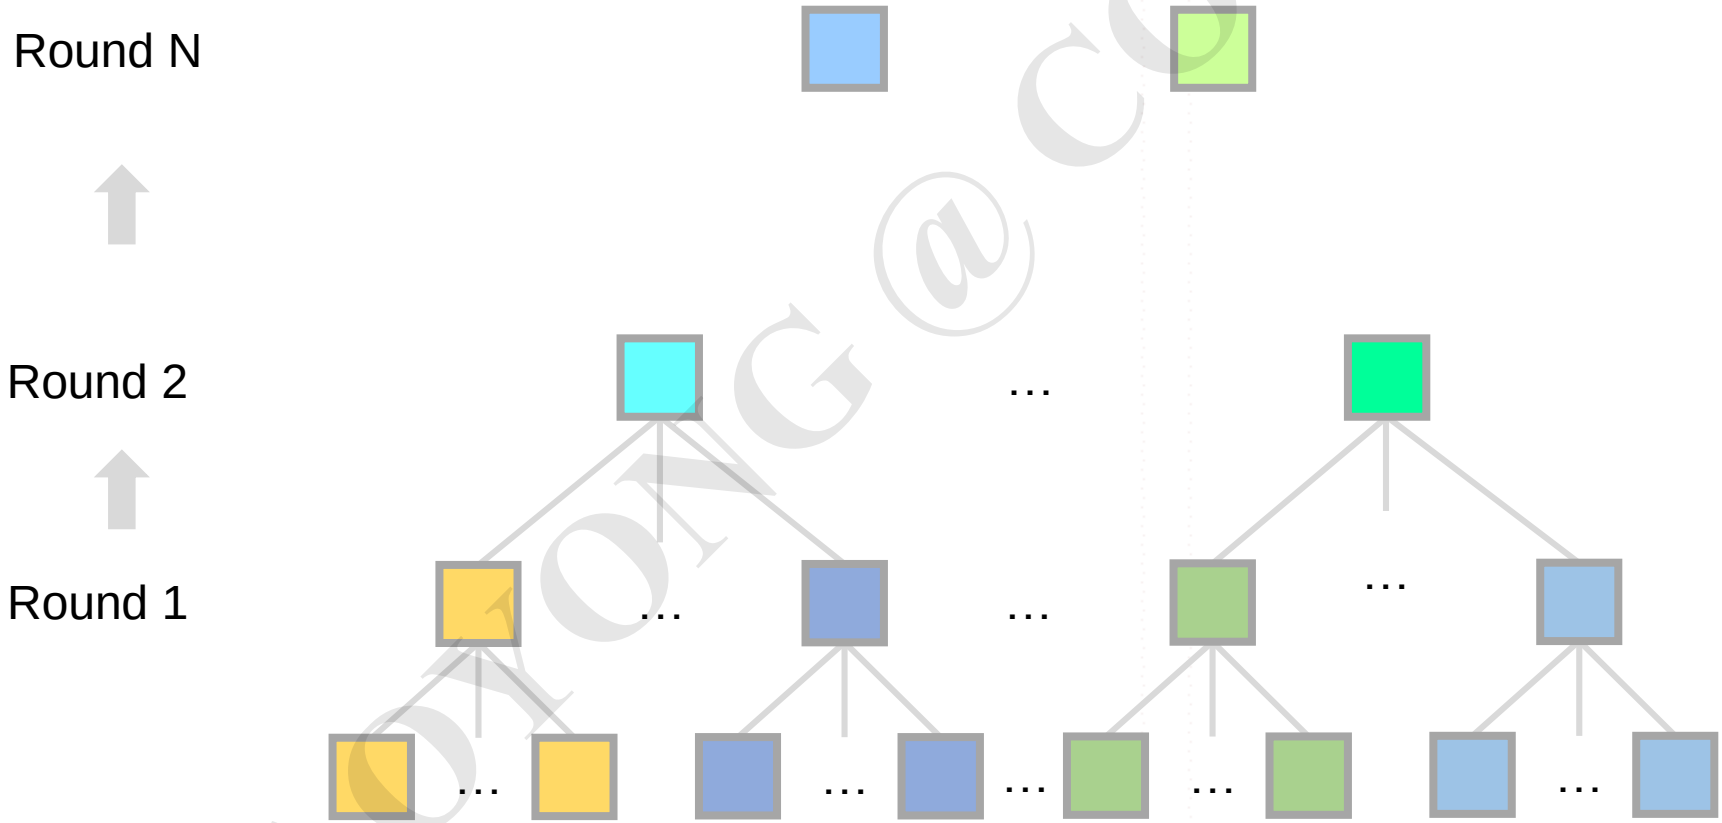

Search from Coarse to Fine

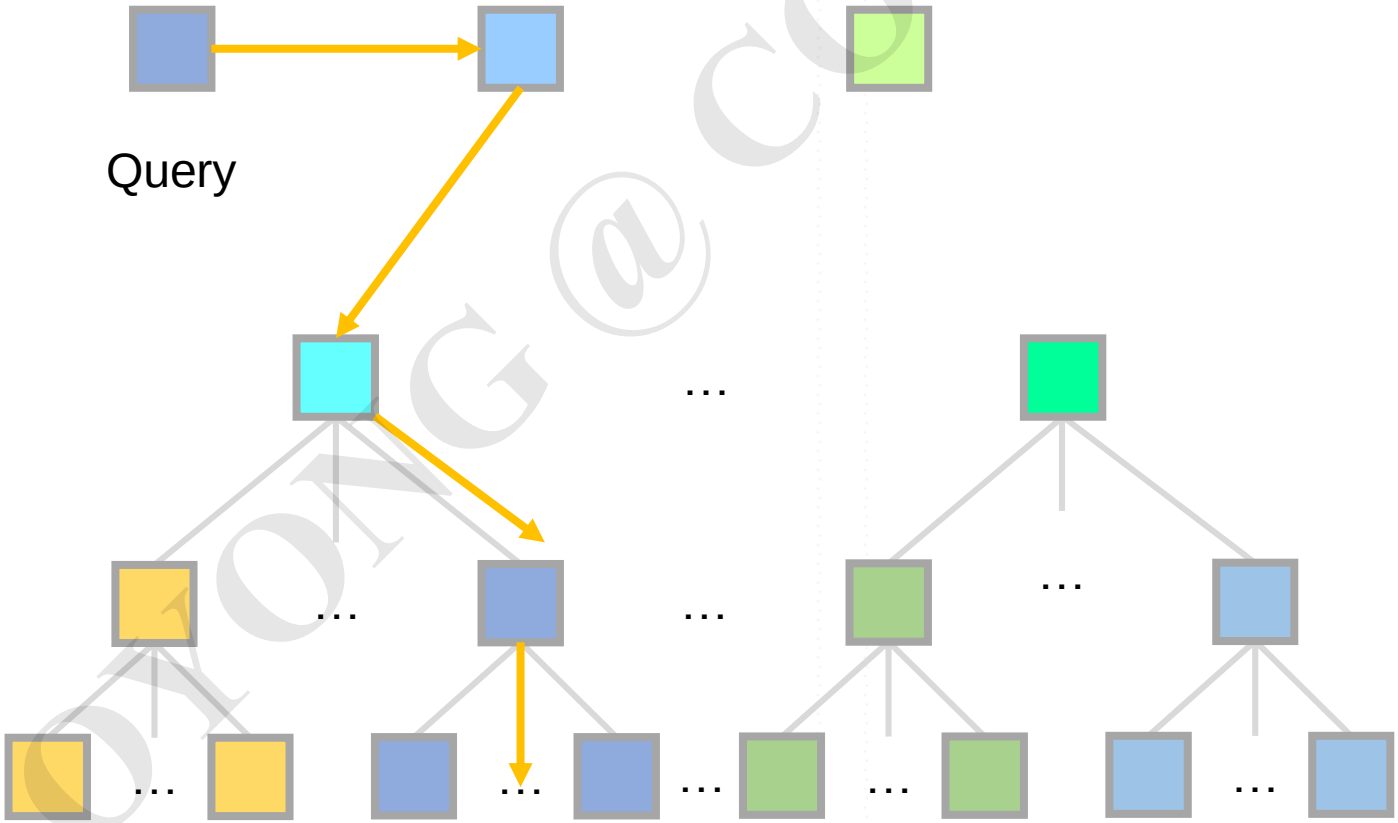

A typical CBIR system

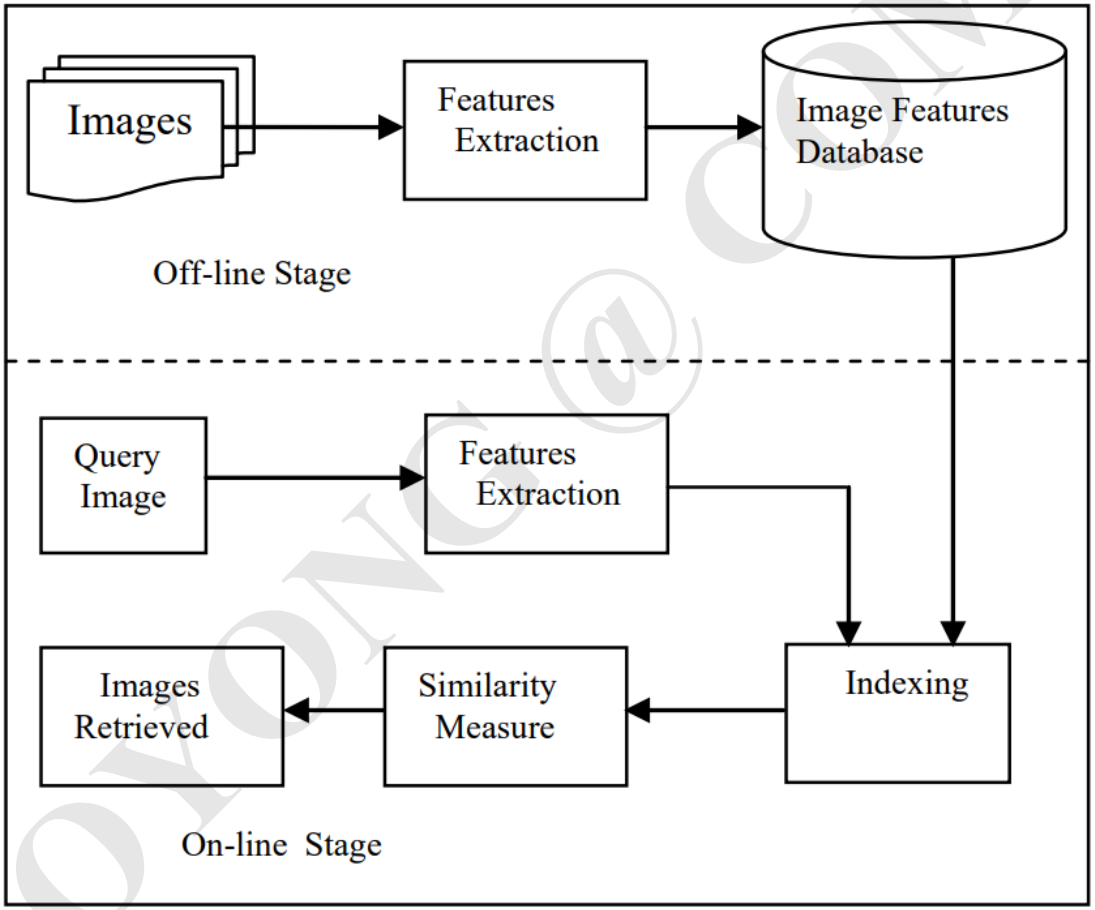

Alkhawlani M, Elmogy M, El Bakry H. Text-based, content-based, and semantic-based image retrievals: A survey[J]. Int. J. Comput. Inf. Technol, 2015, 4(01): 58-66.

Do you notice?

What we introduced are based on an assumption of ***one-feature-per-image***.

How can we deal with images with multiple feature vectors?

Multiple (Local) Feature Vectors

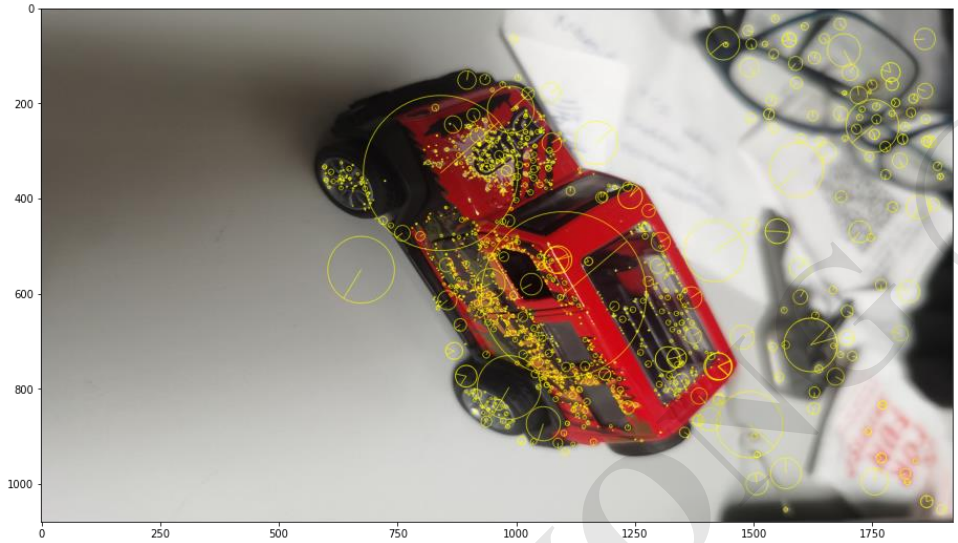

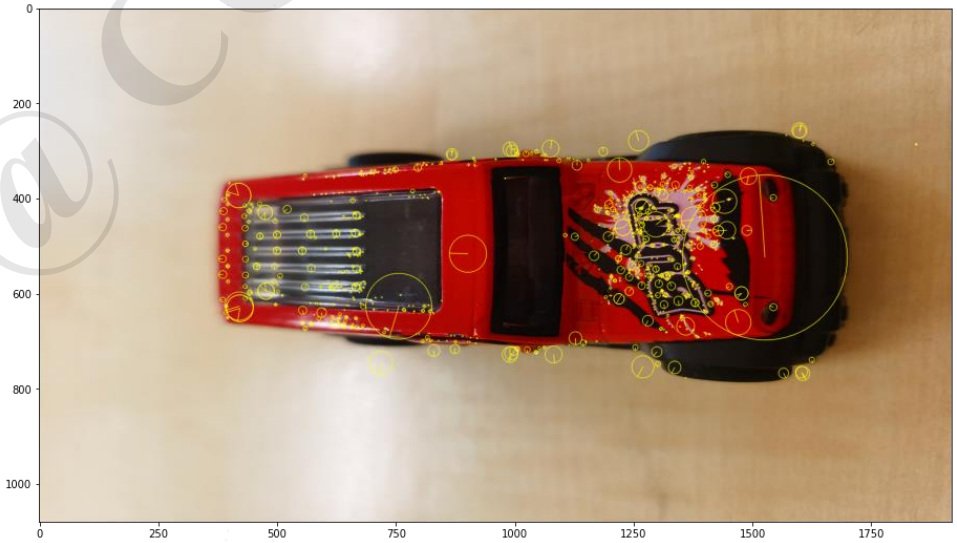

Bag of Visual Words (BoVW)

BoW vs. BoVW

Bag of Visual Words in a Nutshell - The art of choosing important features - Bethea Davida

Visual words and bags

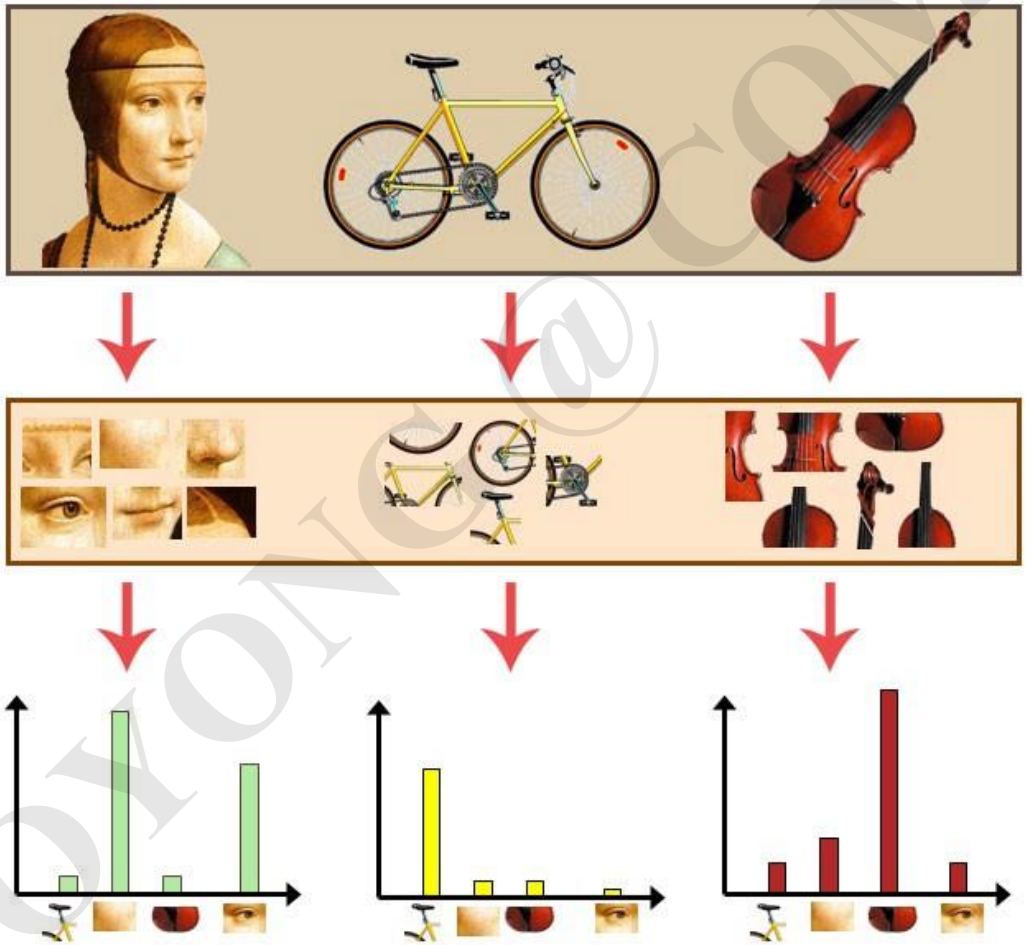

Step 1: Visual Descriptors Extraction

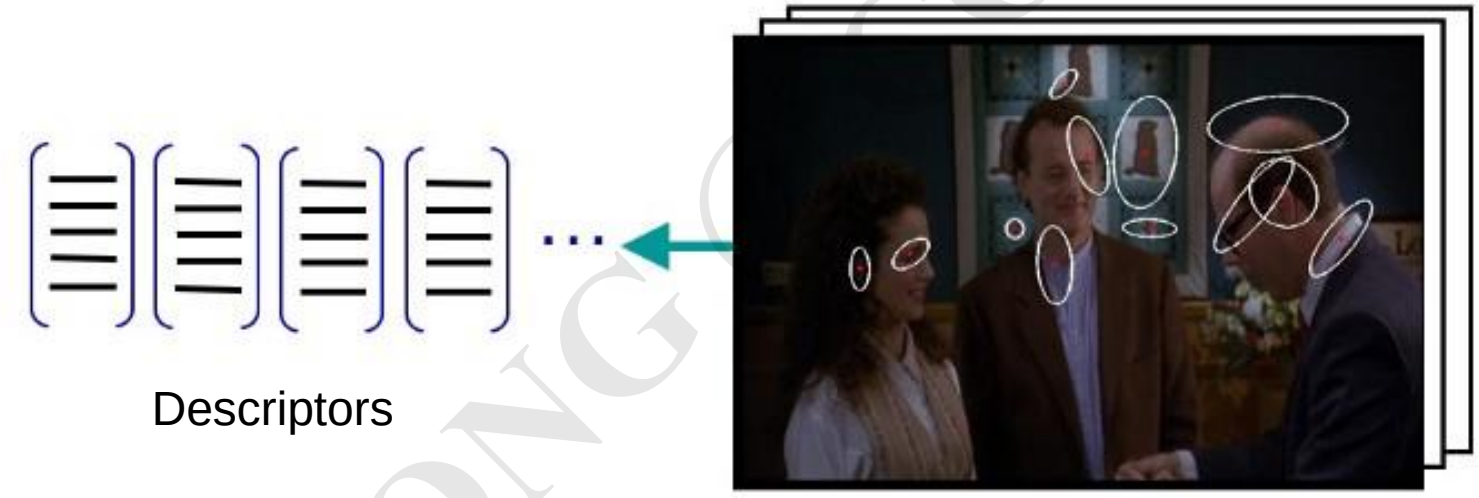

Step 2: Dictionary (CodeBook)

Group the descriptors (from all images) using clustering

Pick one from each cluster as the representative and put them together to construct a dictionary (codebook)

Descriptors in a dictionary are then used as visual words

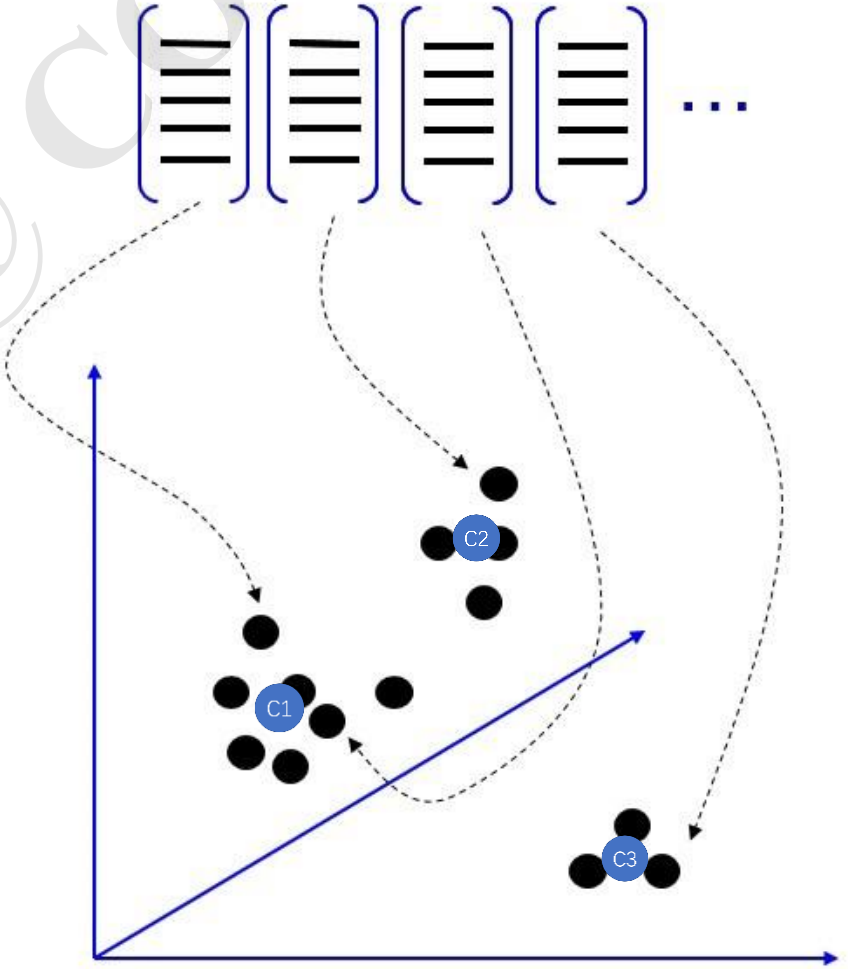

Step 3: Count words in a bag

For each visual word in the dictionary, calculate its frequency in the bag and construct a histogram as the feature vector

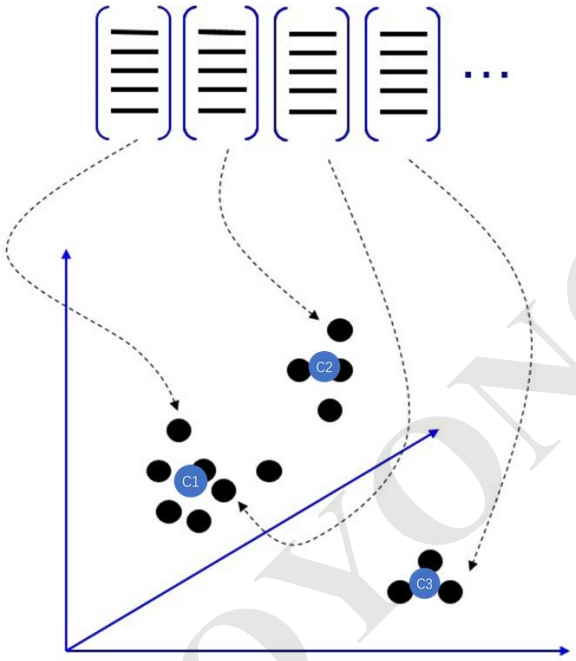

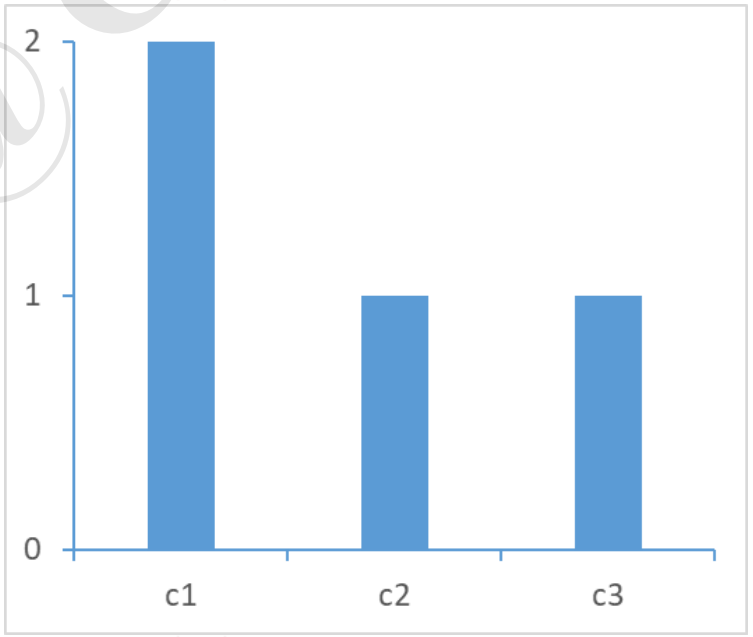

Visual words and bags

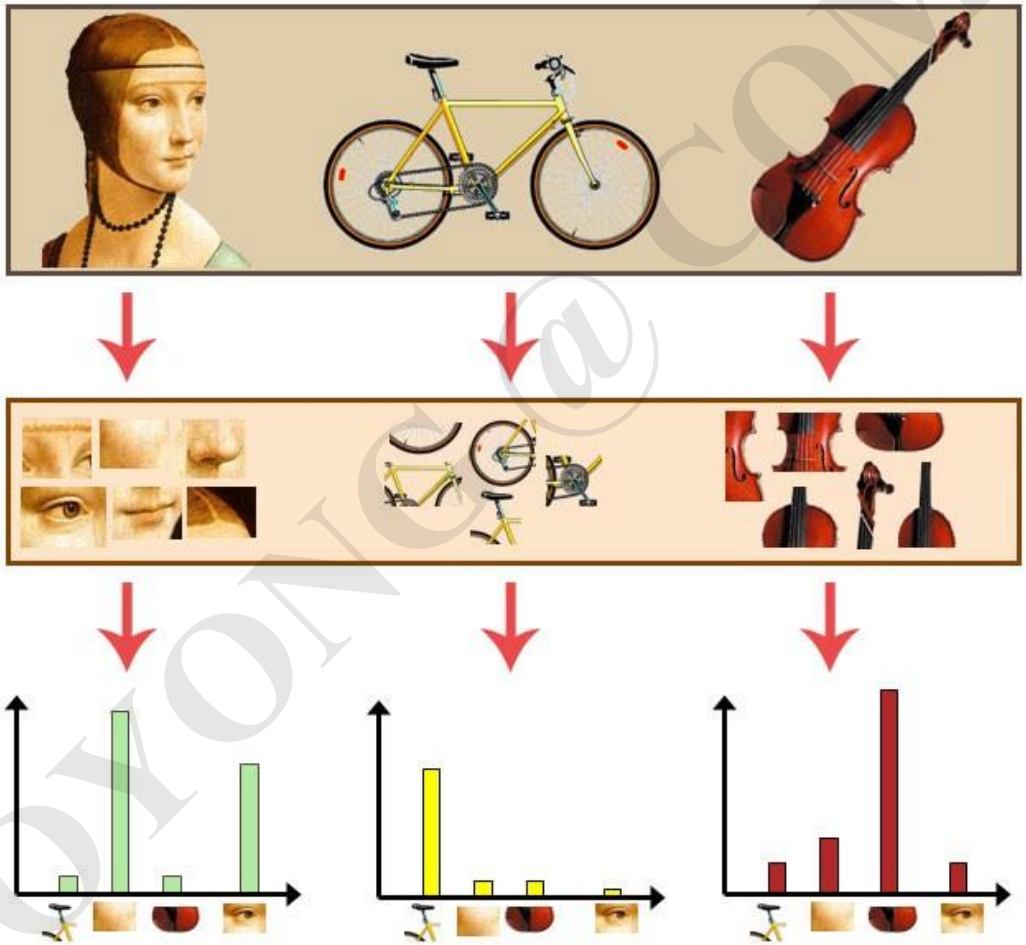

Examples from IMHere

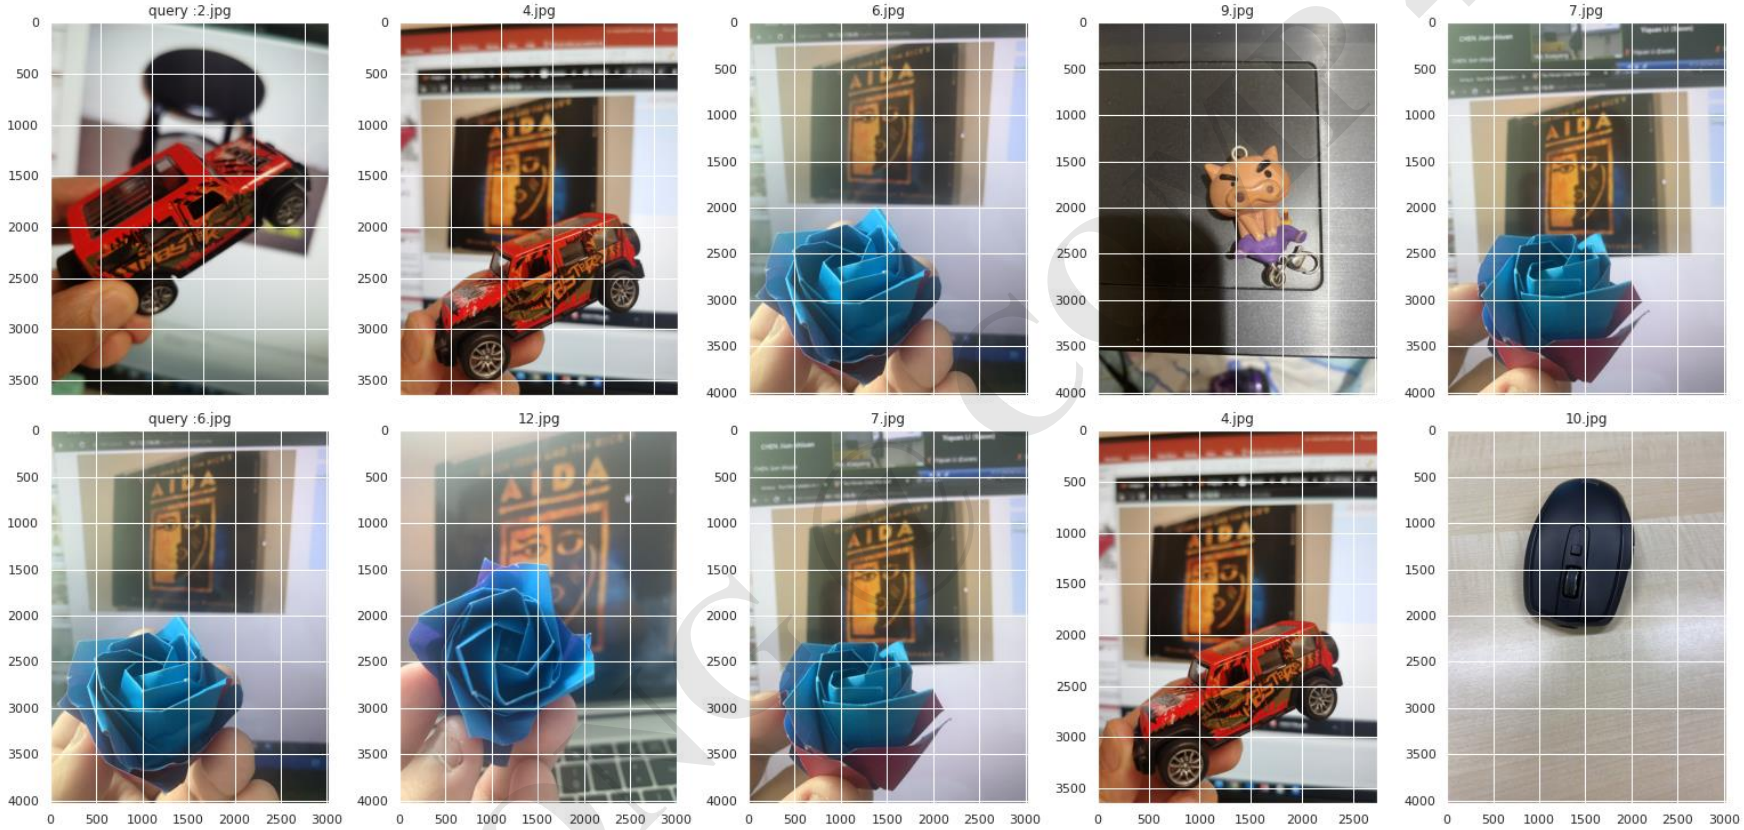

Searching by BoVW

Examples from IMHere

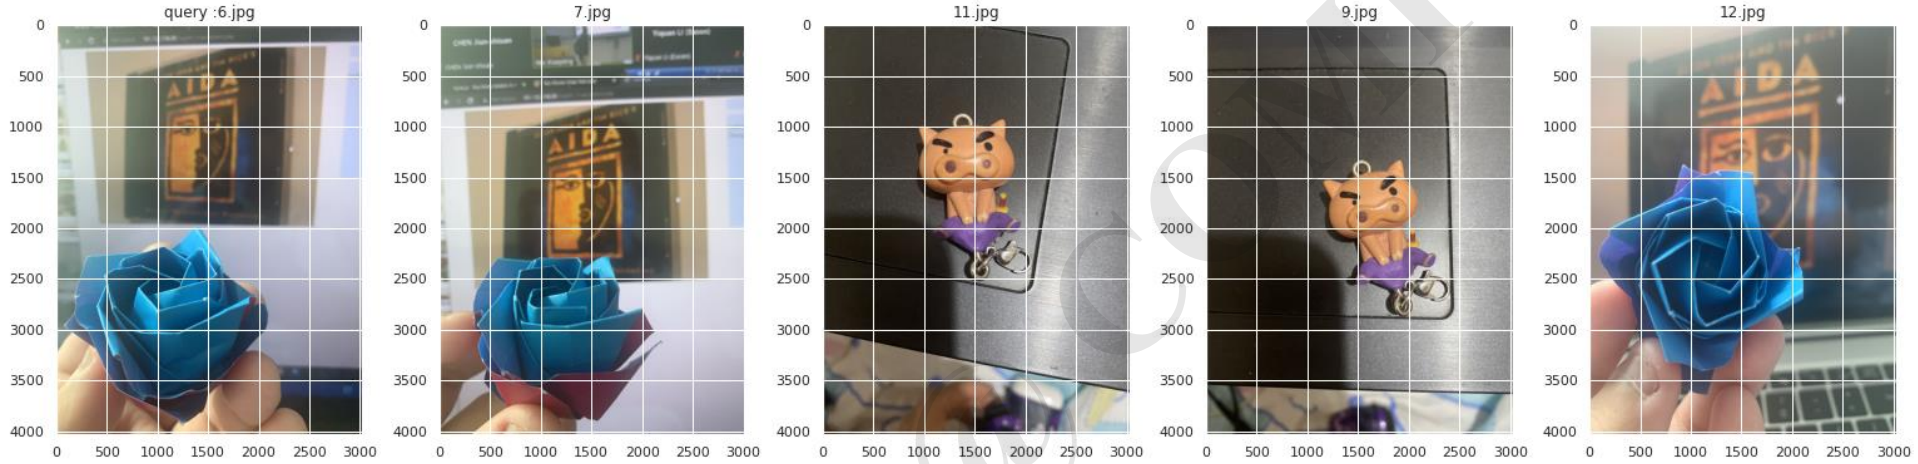

Searching by Histograms

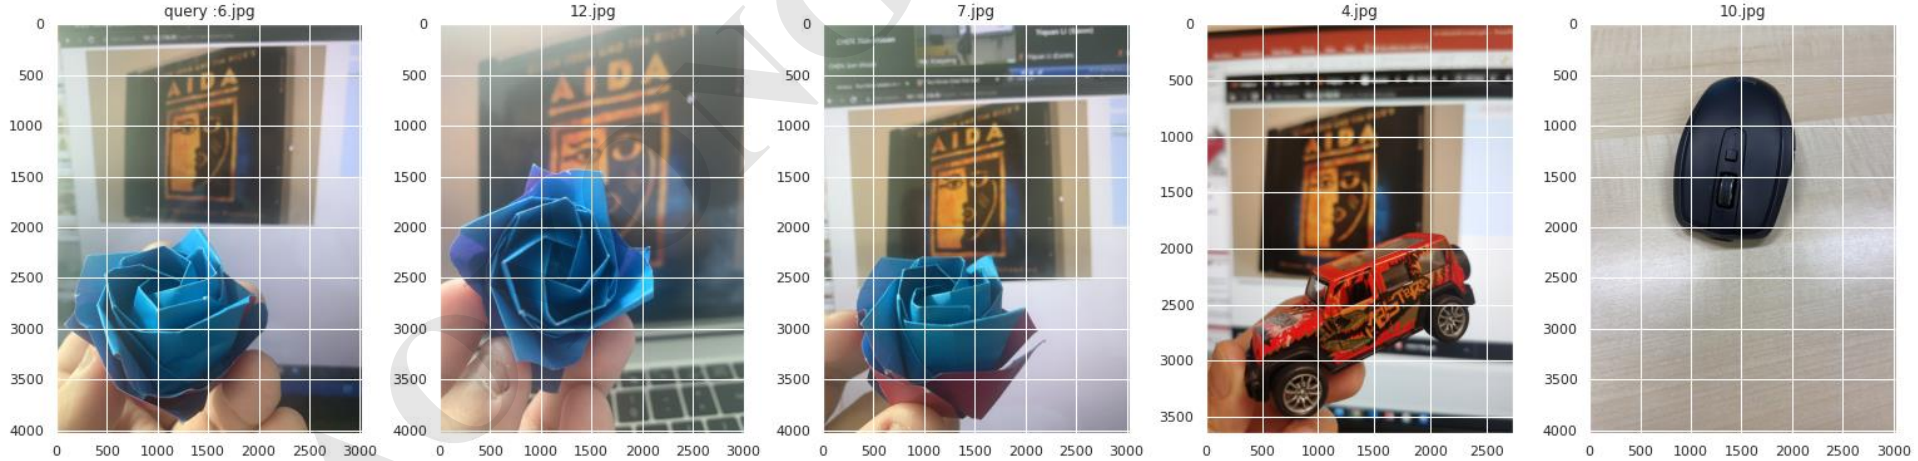

Searching by BoVW

CBIR is not the only solution

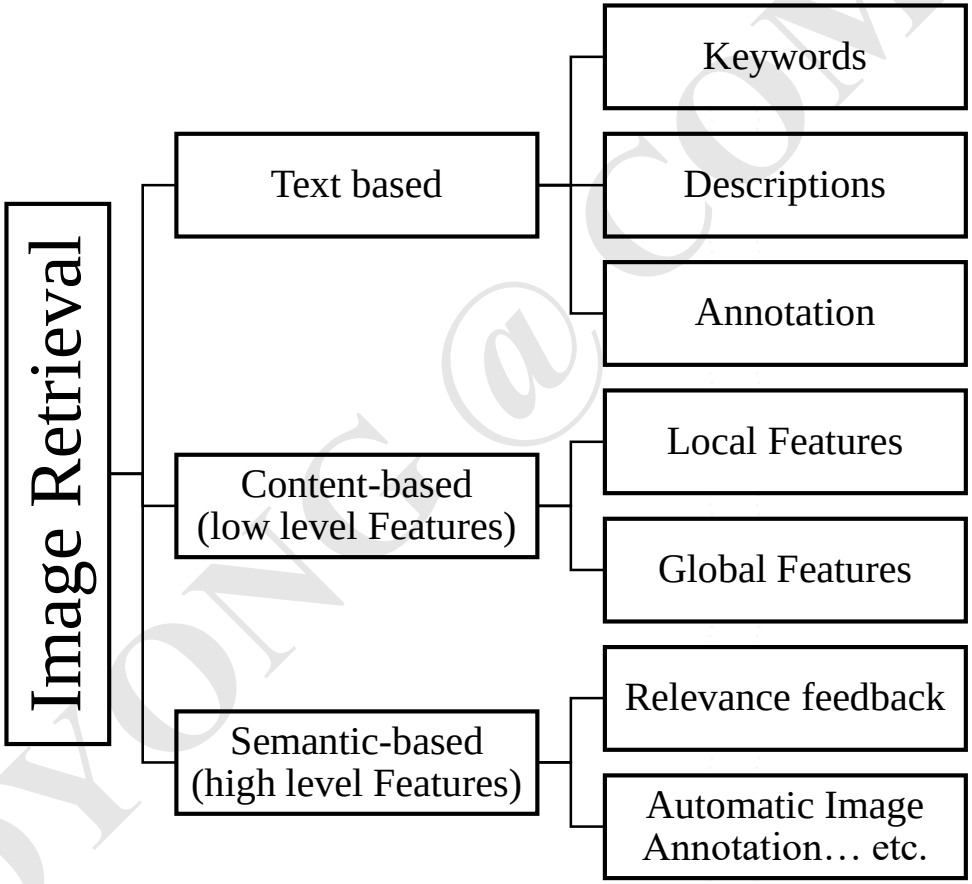

Alkhawlani M , Elmogy M , Bakry H E . Text-based, Content-based, and Semantic-based Image Retrievals: A Survey. 2015.

The New Toy

> The question to answer: is BoVW able to find the nearest keyframes?

0 20 … t t+20

Keyframes selected using a fixed interval of 20 frames

The New Toy

> The question to answer: is BoVW able to find the nearest keyframes?

0 20 … t t+20

0 20 … t t+20

Keyframe tones modified

The New Toy

> The question to answer: is BoVW able to find the nearest keyframes?





0 20 … t t+20

Will the **nearest keyframes** of a query **frame in the feature space** be the same as its **nearest keyframes in the video?**




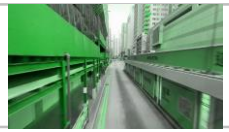


0 20 … t t+20

The New Toy

> The question to answer: is BoVW able to find the nearest keyframes?

0 20 … t t+20

Repeat this for all frames and compose the results as a new video. We should see a smooth tone change if the BoVW worked by locating the rights nearest keyframes

0 20 … t t+20

Code to modify the tone of an image (nparry in BRG)

```
def change_tone(frame,shift):
    # convert the image data to HSV space
    hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
    h,s,v = cv2.split(hsv)
    # modify hue channel by adding shift and modulo 180
    h2 = np.mod(h*0.0 + shift, 180).astype(np.uint8)
    # convert back to RGB space
    frame_new = cv2.cvtColor(cv2.merge([h2,s,v]), cv2.COLOR_HSV2BGR)
    return frame_new
```

Please take it as a challenge and we will release the sample code later. In fact, you will have all the necessary code by completing the tasks of our next tutorial.

Thank you!